



# Reservoir Computing in R: a Tutorial for Using reservoirnet to Predict Complex Time-Series

ISSN 2824-7795

Thomas Ferté Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria

Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux

Kalidou Ba Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria

Dan Dutarte Inria BSO, Inria

Pierrick Legrand Inria BSO, F-33000 Bordeaux, France, Inria

IMB, Institut de Mathématiques de Bordeaux, UMR CNRS 5251, IMB

Vianney Jouhet Inserm Bordeaux Population Health Research Center UMR 1219, team AHeAD, F-33000 Bordeaux, Inserm

Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux

Romain Griffier Inserm Bordeaux Population Health Research Center UMR 1219, team AHeAD, F-33000 Bordeaux, Inserm

Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux

Rodolphe Thiébaut Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria

Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux

Xavier Hinault Inria BSO, F-33000 Bordeaux, France, Inria

Univ. Bordeaux, CNRS, IMN, UMR 5293, Bordeaux, France, CNRS

LaBRI, Univ. Bordeaux, Bordeaux INP, CNRS UMR 5800., LaBRI

Boris Hejblum <sup>1</sup> Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria

Date published: 2024-12-10 Last modified: 2024-12-10

## Abstract

Reservoir Computing (RC) is a machine learning method based on neural networks that efficiently process information generated by dynamical systems. It has been successful in solving various tasks including time series forecasting, language processing or voice processing. RC is implemented in Python and Julia but not in R. This article introduces `reservoirnet`, an R package providing access to the Python API `ReservoirPy`, allowing R users to harness the power of reservoir computing. This article provides an introduction to the fundamentals of RC and showcases its real-world applicability through three distinct sections. First, we cover the foundational concepts of RC, setting the stage for understanding its capabilities. Next, we delve into the practical usage of `reservoirnet` through two illustrative examples. These examples demonstrate how it can be applied to real-world problems, specifically, regression of COVID-19 hospitalizations and classification of Japanese vowels. Finally, we present a comprehensive analysis of a real-world application of `reservoirnet`, where it was used to forecast COVID-19 hospitalizations at Bordeaux University Hospital using public data and electronic health records.

<sup>1</sup>Corresponding author:

*Keywords:* Reservoir Computing, Covid-19, Electronic Health Records, Time series

## 1 **Contents**

2	<b>1 Introduction</b>	3
3	<b>2 RC presentation</b>	4
4	<b>3 Usage workflow</b>	5
5	3.1 Installation . . . . .	5
6	3.2 Package workflow overview . . . . .	7
7	3.3 Basic regression use-case . . . . .	7
8	3.3.1 Covid-19 data . . . . .	7
9	3.3.2 First reservoir . . . . .	7
10	3.3.3 Forecast . . . . .	8
11	3.4 Classification . . . . .	11
12	3.4.1 The Japanese vowel dataset . . . . .	11
13	3.4.2 Classification (sequence-to-vector model) . . . . .	11
14	3.4.3 Transduction (sequence-to-sequence model) . . . . .	12
15	<b>4 Avanced case-study: Covid-19 hospitalizations forecast</b>	13
16	4.1 Introduction . . . . .	13
17	4.2 Methods . . . . .	15
18	4.2.1 Data . . . . .	15
19	4.2.2 Evaluation framework . . . . .	16
20	4.2.3 Models . . . . .	16
21	4.2.4 Hyperparameter optimisation using random search . . . . .	17
22	4.3 Results . . . . .	17
23	4.3.1 Hyperparameter selection . . . . .	17
24	4.3.2 Forecast performance . . . . .	19
25	4.3.3 Number of model to aggregate . . . . .	19
26	4.3.4 Input feature importance . . . . .	19
27	4.4 Discussion . . . . .	24
28	<b>5 Discussion and conclusion</b>	24
29	<b>References</b>	27
30	<b>Session information</b>	30

## 31 1 Introduction

32 Reservoir Computing (RC) is a prominent machine learning method, proposed by Jaeger (2001), Maass,  
33 Natschläger, and Markram (2002) and Lukoševičius and Jaeger (2009) that has gained significant  
34 attention in recent years for its ability to efficiently process information generated by dynamical  
35 systems. This innovative approach leverages the dynamics of a high-dimensional “reservoir” (we  
36 define it below) to perform complex computations and solve various tasks based on the response  
37 of this dynamical system to input signals. RC has demonstrated its efficacy in tackling various  
38 challenges, encompassing pattern classification and time series forecasting in applications ranging  
39 from electrocardiogram analysis to bird calls (Trouvain and Hinaut 2021), language processing  
40 (Hinaut and Dominey 2013), power plants, internet traffic, stock prices, and beyond (Lukoševičius  
41 and Jaeger 2009; Tanaka et al. 2019).

42 Originally, the RC paradigm was implemented in artificial firing-rate neurons (“Echo State Networks”,  
43 Jaeger (2001)) and spiking neurons (“Liquid State Machine”, Maass, Natschläger, and Markram (2002))  
44 as a recurrent neural network (RNN) where the internal recurrent connections, denoted as the  
45 reservoir, are randomly generated and only the output layer (named “read-out”) is trained. The  
46 reservoir projects temporal input signals onto a high-dimensional feature space, facilitating the  
47 learning of non-linear and temporal interactions. Thus, this recurrent layer contains high-dimensional  
48 non-linear recombination of the inputs and past states: it is a “reservoir of computations” from  
49 which useful information can be linearly extracted (or “read-out”) to provide the desired outputs.  
50 This offers the advantage of decreasing the computing time compared to conventional RNNs while  
51 consistently maintaining performance (Vlachas et al. 2020). Besides, this RC paradigm fostered  
52 increasing interest thanks to its ability to be implemented on classical computers, as the hidden  
53 recurrent layer can be kept untrained. A wide range of physical media can be also used to replace  
54 it and Tanaka et al. (2019) recently reviewed this prolific field: from FPGA hardware (Penkovsky,  
55 Larger, and Brunner 2018), to spin waves using magnetic properties (Nakane, Tanaka, and Hirose  
56 2018), skrymions (Prychynenko et al. 2018) or optical implementations (Rafayelyan et al. 2020).  
57 This provides interesting and potentially more efficient alternative to traditional machine learning  
58 computing.

59 RC leverages various hyperparameters to introduce prior knowledge about the relationship between  
60 input variables and output targets. But because the connections within the reservoir are randomly  
61 initialized, the same set of hyperparameters may exhibit diverse behaviors across different instances  
62 of the reservoir connections. This unpredictability makes it challenging to anticipate the performance  
63 of a particular hyperparameter setting, as identical settings may produce varying outcomes when  
64 applied to distinct instances of the reservoir. Moreover, selecting the most suitable hyperparameters  
65 often requires researchers to experiment with multiple combinations on a training dataset and  
66 evaluate their performance on a separate test set<sup>2</sup>. Although this approach can be resource-intensive  
67 and time-consuming, it is a compromise that is acceptable considering the rapid simulation capabilities  
68 offered by RC. Furthermore, there is a current absence of implementation in R, rendering the method  
69 challenging for users unfamiliar with Python (Trouvain and Hinaut 2022) or Julia (Martinuzzi et al.  
70 2022).

71 Here, we offer comprehensive guidance to assist new users in maximizing the benefits of RC. Initially,  
72 a broad introduction to reservoir computing is presented in Section 2, followed in Section 3 by a  
73 tutorial on its application using `reservoirnet`, an R package built upon the `ReservoirPy` Python  
74 module (Trouvain, Rougier, and Hinaut 2022; Trouvain and Hinaut 2022; Trouvain et al. 2020).  
75 Section 3 then introduces the workflow usage on `reservoirnet` for RC with two basic use-cases,

<sup>2</sup>In this article, we employ the term “train set” to refer to the combined dataset consisting of both the training and validation sets, which are cycled through in a cross-validation manner.

and finally, in Section 4 we investigate the various challenges associated with an advanced case-study leveraging RC for forecasting COVID-19 hospitalizations. This case-study exploration includes detailed guidance on the modeling strategy, the selection of hyperparameters, and the implementation process.

## 2 RC presentation

RC is a machine learning paradigm which is most often implemented as Echo State Networks (ESNs), i.e. the firing-rate neuron version (Jaeger 2001). An ESN is described by three matrices of connectivity: an input layer  $W_{in}$ , a recurrent layer  $W$  and an output layer  $W_{out}$ . At each time step, the input vector  $u_t$  is projected into the reservoir which is also combined with reservoir past state  $x(t - 1)$  through the recurrent connections. The output  $y(t)$  is linearly read-out from the reservoir. Input  $W_{in}$  and recurrent  $W$  matrices are kept random; only the output matrix  $W_{out}$  is trained in an offline or online method. Often a ridge regression (i.e. a regularized linear regression) is used to obtain the desired outputs  $y(t)$  from the reservoir states  $x(t)$ . Figure 1 depicts the architecture. For simplicity, we will use the term “reservoir computing” for “Echo State Network” in the remainder of the paper.

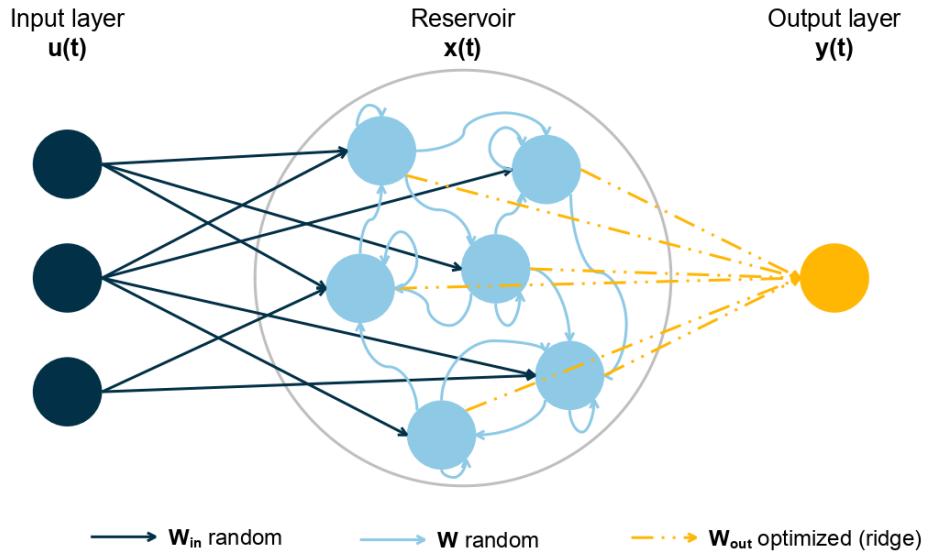


Figure 1: Reservoir computing is composed of an input layer, a reservoir and an output layer. Connection between input layer and reservoir and inside reservoir are random. Only the output layer is optimized based on a ridge regression. Adapted from Trouvain et al. (2020)

The input layer  $u(t)$  is an  $M$ -dimensional vector, where  $M$  is the number of input time series, which corresponds to the values of the input time series at time  $t = 1, \dots, T$ . The reservoir layer  $x(t)$  is an  $N_{res}$ -dimensional vector where  $N_{res}$  is the number of nodes in the reservoir. The value  $x(t)$  is defined as follow:

$$x(t + 1) = (1 - \alpha)x(t) + \alpha \tanh(Wx(t) + W_{in}u(t + 1)) \quad (1)$$

The leaking rate  $\alpha \in [0, 1]$  defines the update rate of the nodes. The closer  $\alpha$  is to 1, the more the reservoir is sensitive to new inputs  $u(t)$ . Therefore, the reservoir state at time  $t + 1$  denoted  $x(t + 1)$  depends on the reservoir state at the previous time  $x(t)$  and the new inputs  $u(t + 1)$ . Both  $W_{in}$  and  $W$  are random matrices of size  $N_{res} \times M$  and  $N_{res} \times N_{res}$  respectively.

98  $W_{in}$  is a matrix (usually sparse) generated using a Bernoulli (bimodal) distribution where each value  
 99 can be either  $-I_{scale}(m)$  or  $I_{scale}(m)$  with an equal probability where  $m = 1, \dots, M$  corresponds to a  
 100 given feature in the input layer. The input scaling, denoted  $I_{scale}$ , is a hyperparameter coefficient  
 101 common to all features from the input layer or specific to each feature  $m$ . In that case, the more  
 102 important the feature is, the greater should be its input scaling.  $W$  is a matrix (usually sparse) where  
 103 values are generated from a Gaussian distribution  $\mathcal{N}(0, 1)$ . Then, the  $W$  matrix is scaled according to  
 104 the defined spectral radius, a hyperparameter defining the highest eigen value of  $W$ .

105 The final layer is a linear regression with ridge penalization where the explanatory features are the  
 106 reservoir state and the variable to be explained is the outcome to predict such that:

$$W_{out} = YX^T(XX^T + \lambda I)^{-1}$$

107 Where  $x(t)$  and  $y(t)$  are accumulated in  $X$  and  $Y$  respectively such that:

$$X = \begin{bmatrix} x(1) \\ x(2) \\ \vdots \\ x(T) \end{bmatrix} \text{ and } Y = \begin{bmatrix} y(1) \\ y(2) \\ \vdots \\ y(T) \end{bmatrix}$$

108 The parameter  $\lambda$  is the ridge penalization which aims to prevent overfitting. Additionally, one can also  
 109 connect the input layer to the output layer to the reservoir nodes. In that case,  $X$  is the accumulation  
 110 of both such that :

$$X = \begin{bmatrix} x(1), u(1) \\ x(2), u(2) \\ \vdots \\ x(T), u(T) \end{bmatrix} \text{ and } Y = \begin{bmatrix} y(1) \\ y(2) \\ \vdots \\ y(T) \end{bmatrix}$$

111 Overall, there are four main hyperparameters to be chosen by the user: i) the leaking rate which  
 112 defines the memory of the RC, ii) the input scaling which defines the relative importance of the  
 113 features, iii) the spectral radius which defines the connections of the neurons inside the reservoir  
 114 which in turn defines the degree of non-linear combination of features, and iv) the ridge penalization  
 115 which controls the degree of overfitting. The choice of hyperparameters often requires the user to  
 116 evaluate the performance of different combinations of hyperparameters on a validation set before  
 117 selecting the optimal combination to forecast on the test set.

### 118 3 Usage workflow

119 In this section, we will cover the basics of `reservoirnet` use including installation, classification and  
 120 regression. A more in depth description is provided in Section 4 with the covid-19 forecast use case.

#### 121 3.1 Installation

122 `reservoirnet` is an R package making the Python module `ReservoirPy` easily callable from R using  
 123 `reticulate` R package Ushey, Allaire, and Tang (2024). It is available on CRAN (see <https://cran.r-project.org/package=reservoirnet>) and can be installed using:

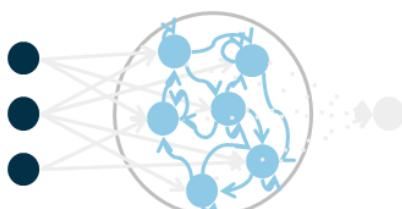
125 Reservoir Computing (RC) is well suited to both regression and classification tasks. We will introduce  
 126 a simple example for both task.

**Input layer :X**



**Instantiate reservoir :**

```
reservoir <- createNode(nodeType = "Reservoir")
```



**Instantiate output layer :**

```
readout <- createNode(nodeType = "Ridge")
```



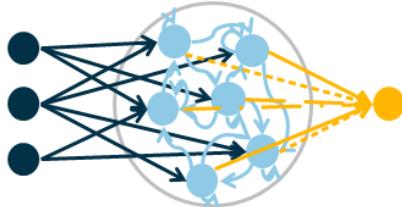
**Build model :**

```
model <- reservoir %>>% readout
```



**Fit model :**

```
fit <- reservoirR_fit(node = model,
  X = X,
  Y = Y)
```



**Forecast :**

```
predict_seq(node = fit$fit, X = X)
```

Figure 2: Workflow of reservoirnet.

127 **3.2 Package workflow overview**

128 The workflow of `reservoirnet` is described in Figure 2. A reservoir model is created by the association  
129 of an input layer (a matrix), a reservoir, and an output layer. Both the reservoir and the output layer  
130 are created using the function `reservoirnet::createNode()` by specifying the node type (i.e., either  
131 `Reservoir` or `Ridge`).

132 This function accepts several arguments to specify the hyperparameters of the reservoir and will be  
133 detailed in future sections. After the reservoir and output layer are created, they can be connected  
134 using the `%>>%` operator, a specific pipe operator dedicated to `reservoirnet`. The model can then be  
135 fitted using `reservoirR_fit()` and used to make predictions on a new dataset using `predict_seq()`.

136 **3.3 Basic regression use-case**

137 **3.3.1 Covid-19 data**

138 In this first use-case, we will introduce the fundamental usage of the `reservoirnet` package. This  
139 demonstration will be conducted using the COVID-19 dataset that is included within the package.  
140 These data encompass hospitalization, positive RT-PCR (Reverse Transcription Polymerase Chain  
141 Reaction) results, and overall RT-PCR data sourced from Santé Publique France, which are publicly  
142 available on `data.gouv.fr` (for further details, refer to `help(dfCovid)`). Our primary objective is to  
143 predict the number of hospitalized patients 14 days into the future. To accomplish this, we will  
144 initially train our model on data preceding the date of January 1, 2022, and subsequently apply it to  
145 forecast values using the subsequent dataset.

146 We can proceed by loading useful packages - namely `ggplot2` Wickham, Navarro, and Pedersen  
147 (2018) and `dplyr` Wickham et al. (2023), data and define the task:

148 Due to the substantial fluctuations observed in both RT-PCR metrics, our initial step involves applying  
149 a moving average computation over the most recent 7-day periods for these features. Additionally,  
150 we augment the dataset by introducing an `outcome` column and an `outcomeDate` column, which  
151 will serve as valuable inputs for model training. Moreover, we calculate the `outcome_deriv` as the  
152 difference between the `outcome` and the number of hospitalized patients (`hosp`), representing the  
153 variation in hospitalization in relation to the current count of hospitalized individuals. The resulting  
154 smoothed data is visualized in Figure 3.

155 **3.3.2 First reservoir**

156 Setting a reservoir is done with the `createNode()` function. The important hyperparameters are the  
157 following :

- 158 • Number of nodes (`units`) : it corresponds to the number of nodes inside the reservoir. Usually,  
159 the more the better, but more nodes increases the computation time.
- 160 • Leaking rate (`lr`) : the leaking rate corresponds to the balance between the new inputs and the  
161 previous state. A leaking rate of 1 only consider information from new inputs.
- 162 • Spectral radius (`sr`): the spectral radius is the maximum absolute eigenvalue of the reservoir  
163 connectivity matrix. A small spectral radius induces stable dynamics inside the reservoir, a  
164 high spectral radius induces a chaotic regime inside the reservoir.
- 165 • Input scaling (`input_scaling`): the input scaling is a gain applied to the input features of the  
166 reservoir.
- 167 • Warmup (`warmup`) : it corresponds to the number of time step during which the data are  
168 propagating into the reservoir but not used to fit the output layer. This hyperparameter is set  
169 in the `reservoirR_fit()` function.

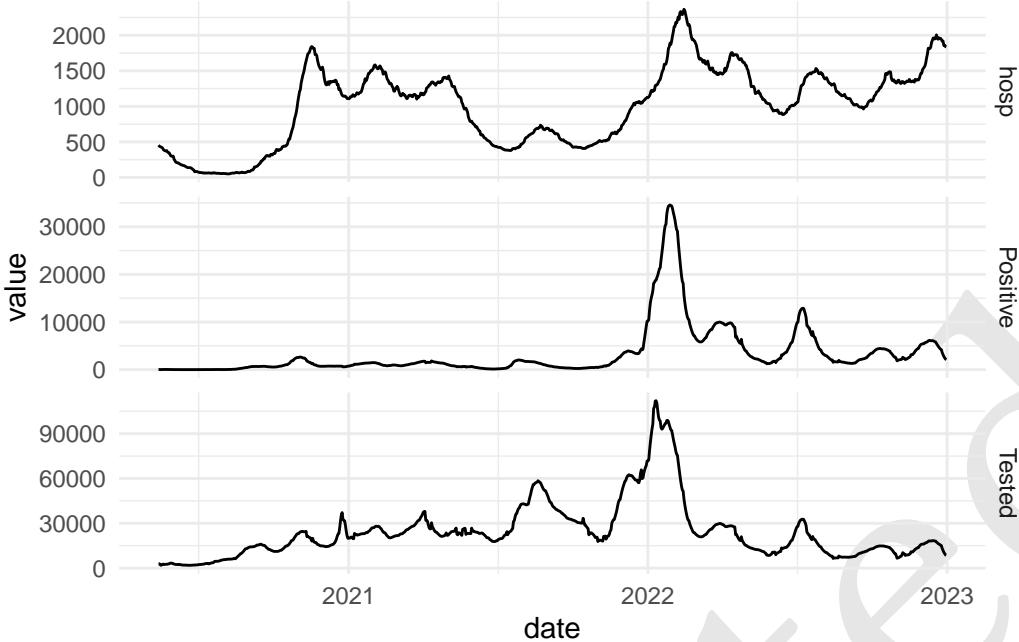


Figure 3: Hospitalizations, number of positive RT-PCR and number ofx RT-PCR of Bordeaux University Hospital.

170 In addition, we can set the seed (seed). Because the reservoir connections are set at random, setting  
171 the seed is a good approach to ensure reproducibility.

172 For this part of the tutorial, we will set the hyperparameter at a given value. Hyperparameter  
173 optimization will be detailed at Section 4.

174 Then we can feed the data to the reservoir and see the activation state of the reservoir  $x(t)$ . To do so,  
175 we first prepare the data and transform it to a matrix.

176 Then we run the `predict_seq()` function. It takes as input a node (i.e a reservoir or a reservoir  
177 associated with an output layer) and the feature matrix.

178 Now we can visualize node activation using the `plot()` function presented at Figure 4 .

179 Numerous nodes within the system exhibit a consistent equilibrium state. The challenge arises when  
180 the output layer attempts to extract knowledge from these nodes, as they do not convey meaningful  
181 information. This issue can be attributed to the disparate scales of the features. To address this  
182 concern, a practical approach involves normalizing the features by dividing each of them by their  
183 respective maximum values, thereby scaling them within the range of  $-1$  to  $1$  by dividing by the  
184 maximum of the absolute value. Of note, here the features will be scaled between  $0$  and  $1$  because all  
185 features are positive.

186 We then feed them to the reservoir and plot the node activation again. Compared to Figure 4, the  
187 obtained node activation at Figure 5 shows interesting trend outputs as no node seems saturated.

### 188 3.3.3 Forecast

189 In order to train the reservoir, we should train the last layer which linearly combines the neuron's  
190 output.

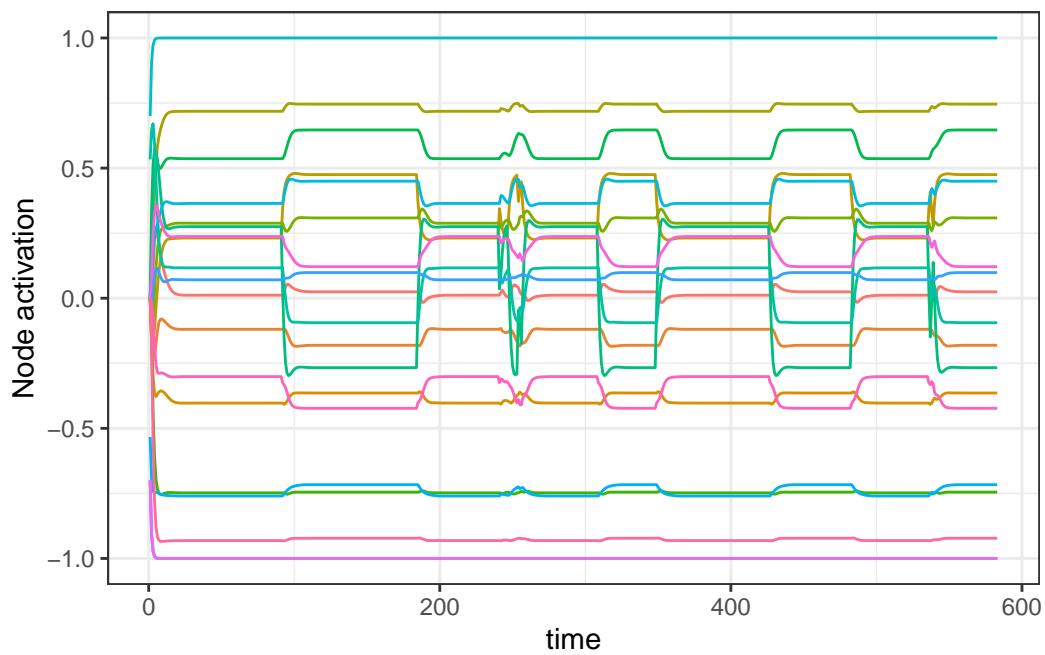


Figure 4: 20 random nodes activation over time.

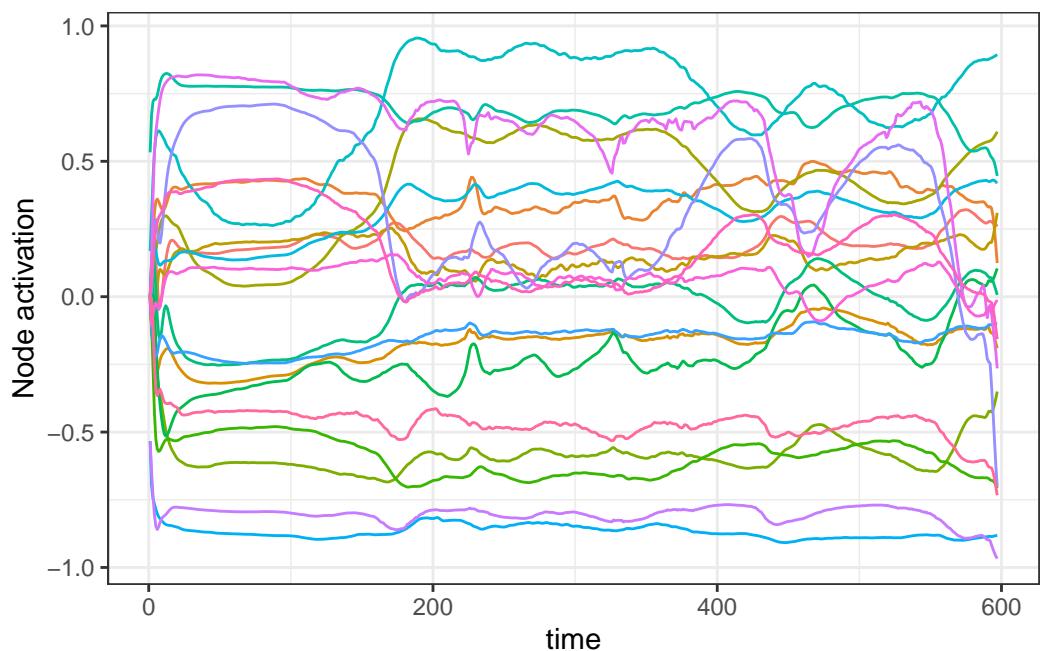


Figure 5: 20 random node activation over time. Scaled features.

191 **3.3.3.1 Set the ESN**

192 Initially, we establish the output layer with the `createNode()` function, incorporating a ridge penalty  
193 set at `1e3`. It's important to note that this hyperparameter can be subject to optimization, a topic  
194 that will be explored in Section 4. This parameter plays a pivotal role in fine-tuning the model's  
195 conformity to the data. When set excessively high, the risk of underfitting arises, whereas setting  
196 it too low can lead to overfitting. We connect the output layer to the reservoir, with the `link()`  
197 function, making the model ready to be trained.

198 **3.3.3.2 Set the data**

199 First we separate the train set on which we will learn the ridge coefficients and the test set on which  
200 we will make the forecast. We define the train set to be all the data before 2022-01-01 and the test  
201 data to be all the data to have forecast both on train and test sets.

202 We standardize with the same formula as seen before. We learn the standardization on the training  
203 set and apply it on the test set. Then we convert the dataframe to matrix.

204 **3.3.3.3 Train the model and predict**

205 We then feed the reservoir with the train set using the `reservoirR_fit()` function. To do so, we set  
206 a `warmup` of 30 days during which the data are propagating into the reservoir but not used to fit the  
207 output layer.

208 Now that the ridge layer is trained, we can forecast using the `predict_seq()` function. We set the  
209 parameter `reset` to `TRUE` in order to clean the reservoir from the data used by the training set.

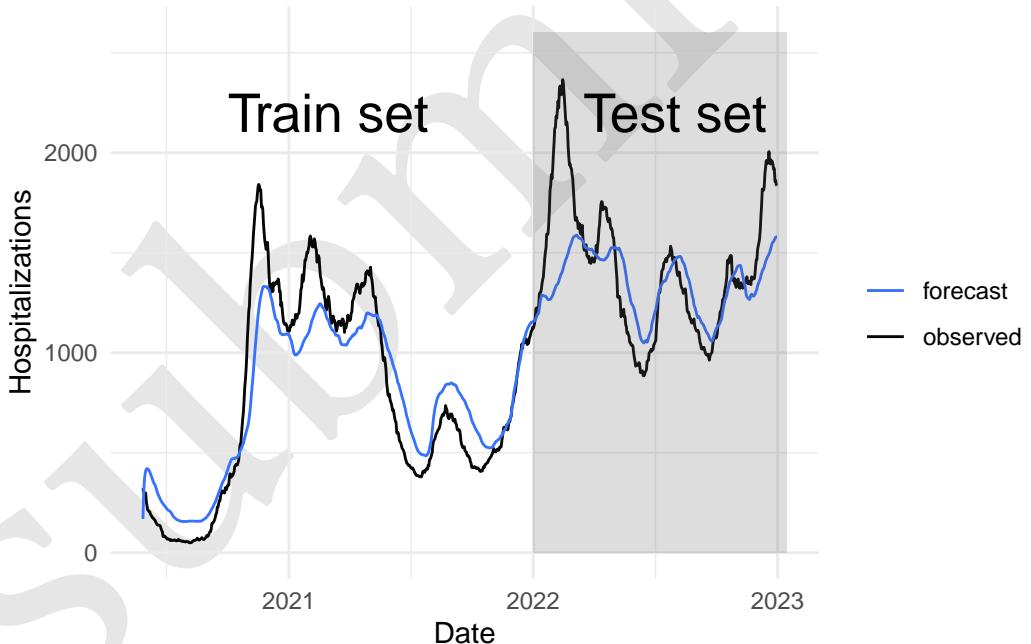


Figure 6: Forecast

210 We observe that the model forecast at Figure 6 is not fully accurate, both on the test set and the  
211 train set. In that case, one option could be to reduce ridge penalization to fit more closely the data,  
212 the optimization of ridge hyperparameter will be discussed at Section 4. Another possibility is to  
213 ease the learning of the algorithm by forecasting the variation of the hospitalization instead of

214 the number of hospitalized patients. For that step, we will learn on the `outcome_deriv` contained  
 215 in `yTrain_variation` data which is defined outcome as `outcome_deriv = outcome - hosp`. As  
 216 depicted at Figure 7, this strategy improved the model forecast.

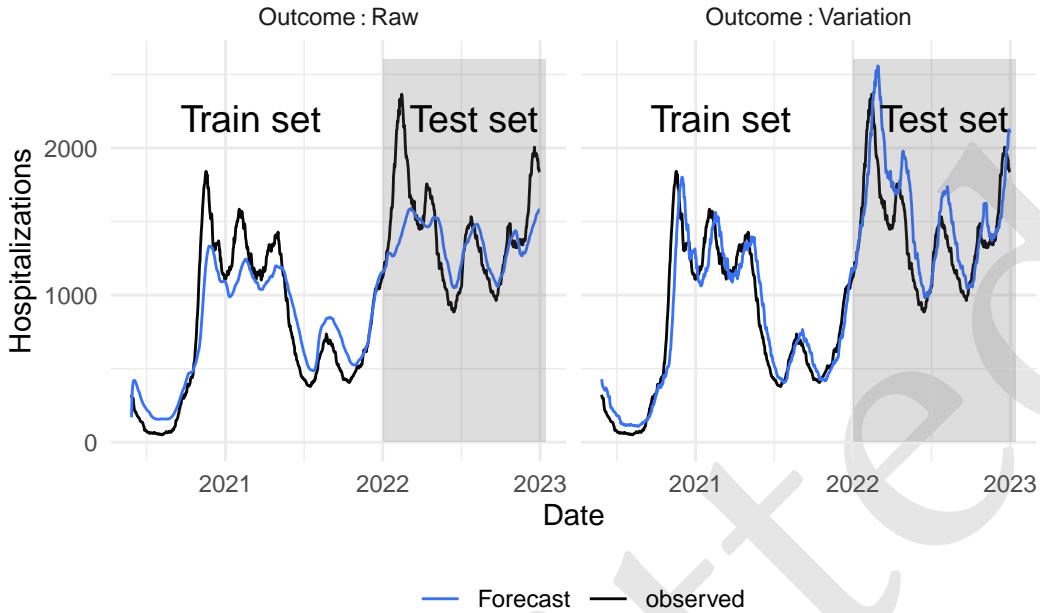


Figure 7: Covid-19 hospitalizations forecast. The model is either trained to forecast the number of hospitalizations (denoted Raw) or the variation of the hospitalizations compared to current level of hospitalisation (denoted Variation)

### 217 3.4 Classification

#### 218 3.4.1 The Japanese vowel dataset

219 This example is largely inspired from the [classification tutorial of reservoirpy](#). To illustrate the  
 220 classification task, we will use the Japanese vowel dataset (Kudo, Toyama, and Shimbo (1999)). The  
 221 data can be loaded from `reservoirnet` as follow :

222 The dataset comprises 640 vocalizations of the Japanese vowel æ, contributed by nine distinct  
 223 speakers. Each vocalization represents a time series spanning between 7 and 29 time steps, encoded  
 224 as a 12-dimensional vector denoting the Linear Prediction Coefficients (LPC). A visual representation  
 225 of six distinct utterances from the test set, originating from three different speakers, is depicted in  
 226 Figure 8.

227 The primary objective involves the attribution of each utterance to its respective speaker, this is  
 228 denoted as classification or sequence-to-vector encoding. The secondary objective involves the  
 229 attribution of each time step of each utterance to its speaker, this is denoted as transduction or  
 230 sequence-to-sequence encoding.

#### 231 3.4.2 Classification (sequence-to-vector model)

232 The first approach is the sequence-to-vector encoding. For this task we aim to predict the speaker of  
 233 the whole utterance (i.e the label is assigned to the whole sequence). We first start by creating the  
 234 reservoir and the output layer using `createNode()` function.

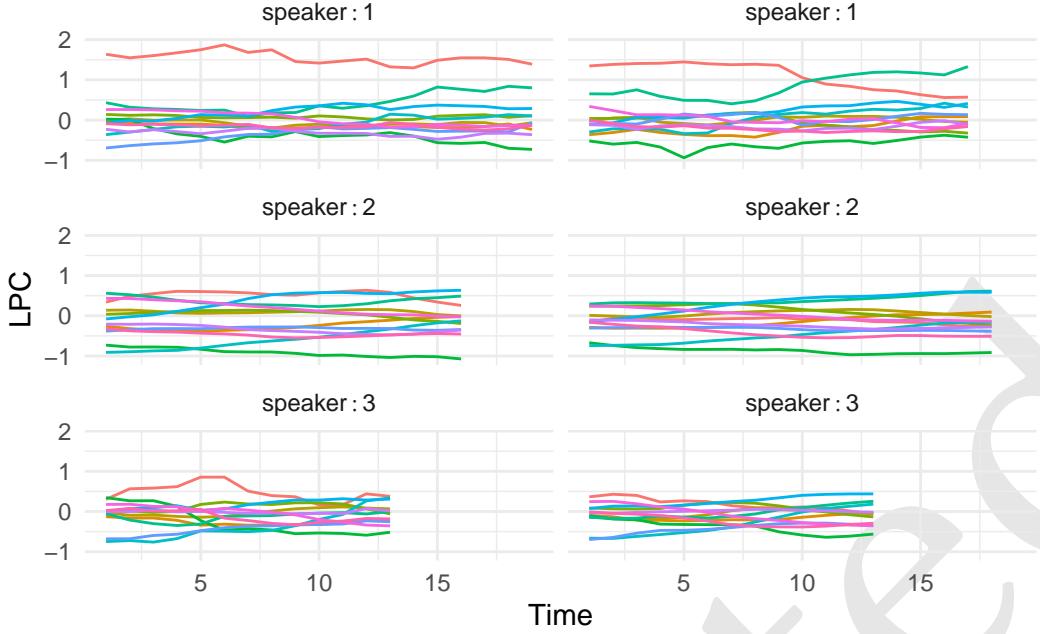


Figure 8: Vowel dataset, sample with 3 speakers and 2 utterance each.

235 To perform this task, we need to modify the training and testing processes. Leveraging the inherent  
 236 inertia of the reservoir, information from preceding time steps is preserved, effectively endowing the  
 237 RC with a form of memory. Consequently, the final state vector encapsulates insights gathered from  
 238 all antecedent states. In the context of the sequence-to-vector encoding task, only this ultimate state  
 239 is employed. This process is executed as follows:

240 Then we can train the readout based on this last state vector. In that case, `Y_train` contains a single  
 241 label for each utterance.

242 The prediction is also modified using only the final state :

243 Figure 9 shows the prediction for the 6 utterances depicted at Figure 8 where the model correctly  
 244 identifies the speaker.

245 Then, we can also compute the overall accuracy :

246 [1] "Accuracy: 92.703%"

### 247 3.4.3 Transduction (sequence-to-sequence model)

248 For this task, the goal is to predict the speaker for each time step of each utterance. The first  
 249 step is to get the data where the label is repeated for each time step. This is easily done with the  
 250 `repeat_targets` argument as follow :

251 Then we can train a simple Echo State Network to solve this task. For this example we will connect  
 252 both the input layer and the reservoir layer to the readout layer which is performed by the `%>>%`  
 253 operator :

254 We can then fit the model and predict the labels for the test data. The `reset` parameter is set to TRUE  
 255 to remove information from the reservoir from the training process.

256 From the `Y_pred` and `Y_test` we represent at Figure 10 the predictions for the same patients as in  
 257 Figure 8.

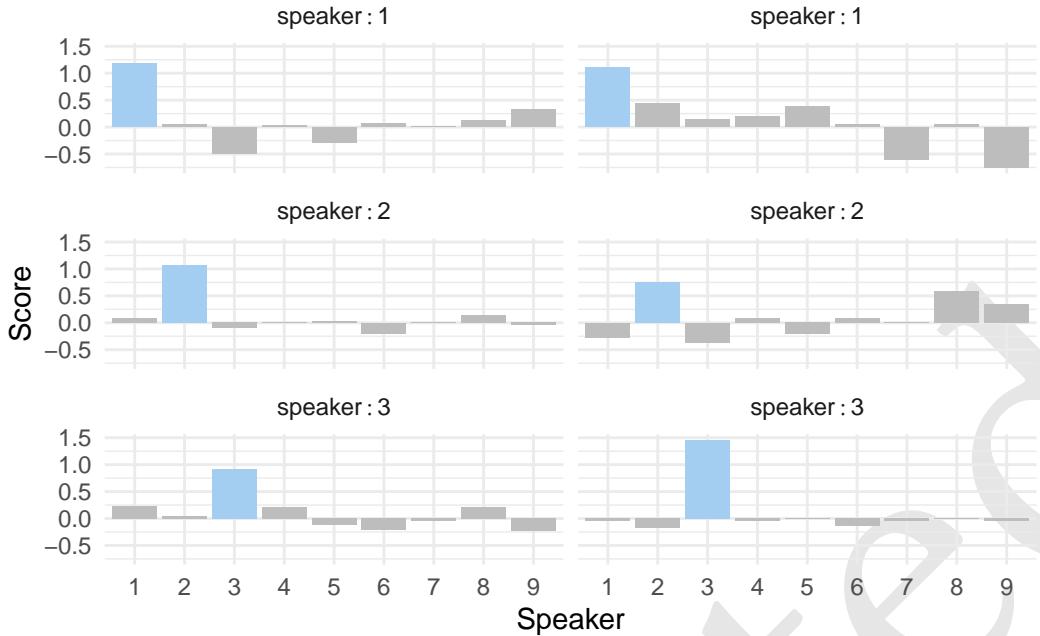


Figure 9: Prediction in a sequence-to-sequence approach 6 samples with 3 speakers and 2 utterance each. The speaker to predict is depicted in blue. For each of the 6 utterance, the model correctly identifies the speaker.

258 For those 6 utterances, the model correctly identify the speaker for most of the time steps. We can  
 259 then evaluate the overall accuracy of the model :

260 [1] "Accuracy: 92.456%"

## 261 **4 Avanced case-study: Covid-19 hospitalizations forecast**

### 262 **4.1 Introduction**

263 Since late 2020, millions of cases of SARS-CoV-2 infection have been documented across the globe  
 264 (World Health Organisation 2020; COVID-19 Cumulative Infection Collaborators 2022; Carrat et al.  
 265 2021). This ongoing pandemic has exerted significant strain on healthcare systems, resulting in a surge  
 266 in hospitalizations. This surge, in turn, necessitated modifications to the healthcare infrastructure and  
 267 gave rise to population-wide lockdown measures aimed at preventing the saturation of healthcare  
 268 facilities (Simões et al. 2021; Hübner et al. 2020; Kim et al. 2020). The capacity to predict the  
 269 trajectory of the epidemic on a regional scale is of paramount importance for effective healthcare  
 270 system management.

271 Numerous COVID-19 forecasting algorithms have been proposed using different methods (e.g en-  
 272 semble, deep learning, compartmental), yet none has proven entirely satisfactory (Cramer et al. 2022;  
 273 Rahimi, Chen, and Gandomi 2021). In France, short-term forecasts with different methods have  
 274 been evaluated with similar results (Paireau et al. 2022; Carvalho et al. 2021; Mohimont et al. 2021;  
 275 Pottier 2021). In this context a machine learning algorithm based on linear regression with elastic-net  
 276 penalization, leveraging both Electronic Health Records (EHRs) and public data, was implemented at  
 277 Bordeaux University Hospital (Ferté et al. 2022). This model, which aimed at forecasting the number  
 278 of hospitalized patients at 14 days, showed good performance but struggled to accurately anticipate  
 279 dynamic shifts of the epidemic.

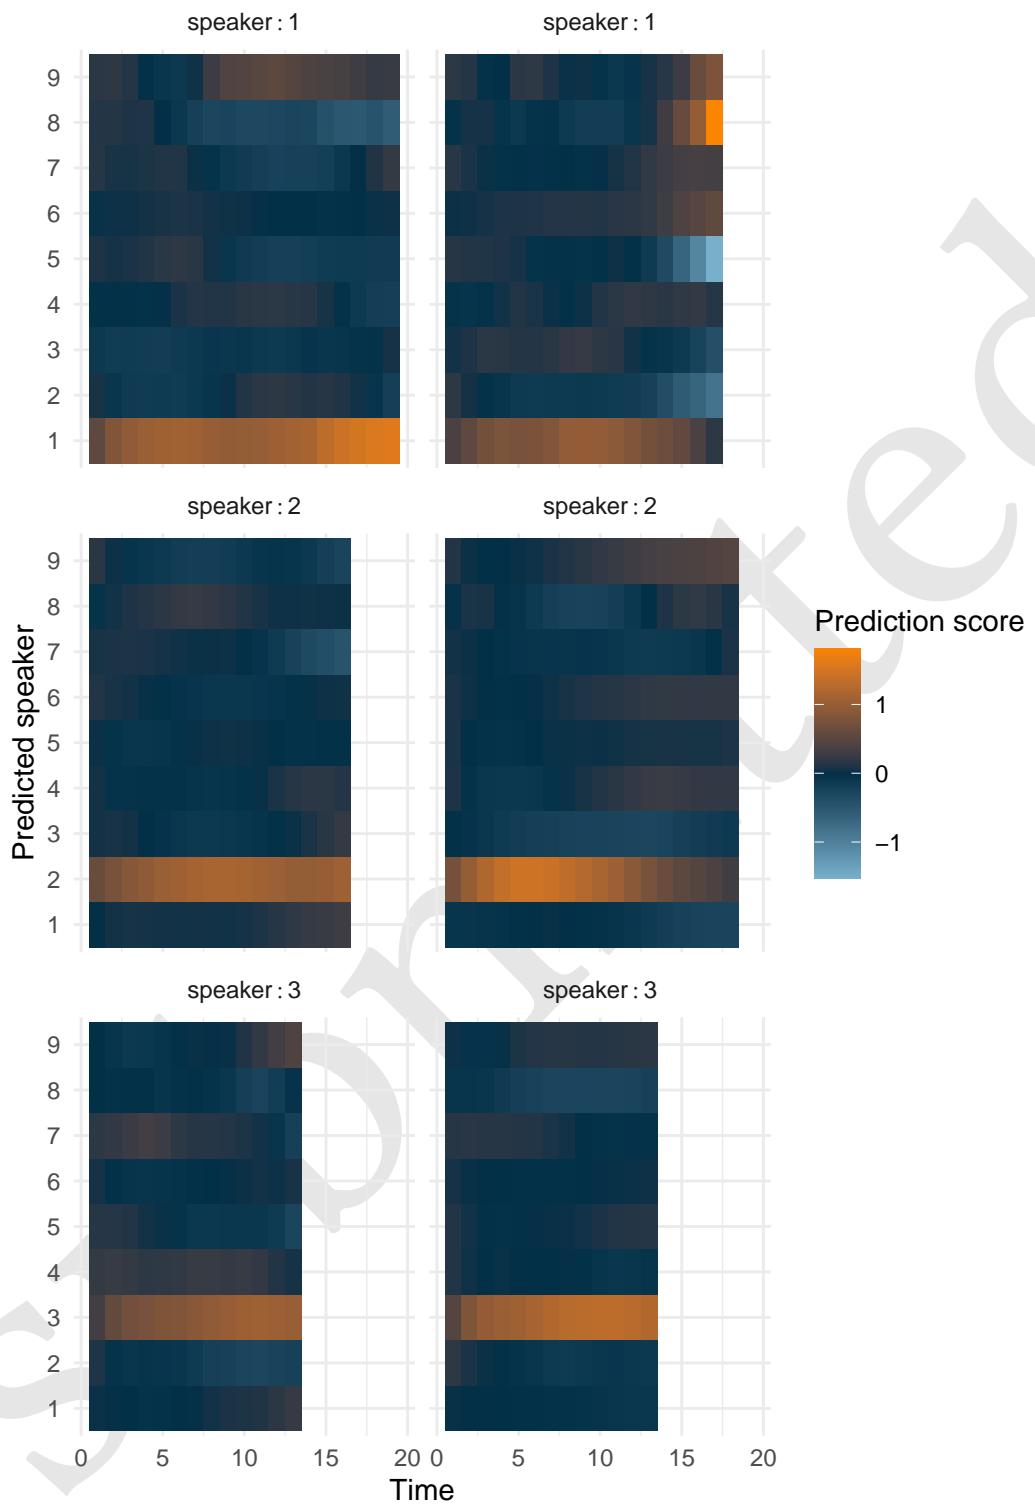


Figure 10: Prediction in a sequence-to-sequence approach 6 samples with 3 speakers and 2 utterance each. The higher the score of the speaker, the lighter the color.

RC has been used in the context of covid-19 epidemic forecast (Kmet and Kmetova 2019; Liu et al. 2023; Ray, Chakraborty, and Ghosh 2021; Zhang et al. 2023; Ghosh et al. 2021). Among them, Ghosh et al. (2021), Liu et al. (2023) and Ray, Chakraborty, and Ghosh (2021) used it to forecast epidemic, Zhang et al. (2023) performed sentiment analysis and Kmet and Kmetova (2019) used it to solve optimal control related to vaccine. The evaluation of RC for epidemic forecast showed promising results in all approaches, being competitive with Long-Short Term Memory (LSTM) and Feed-Forward Neural Network (FFNN) in Ray, Chakraborty, and Ghosh (2021). However, the test period was short for Ghosh et al. (2021) (21 and 14 days) and Ray, Chakraborty, and Ghosh (2021) (86 days) making it difficult to evaluate the behavior of the methods during epidemic dynamic shift. This was not the case for Liu et al. (2023) (6 months) but they implemented daily ahead forecast which would be difficult to use to manage a hospital. Finally, all three implementations used only one time series as input whereas it has been shown that using different data sources could improve forecast Ferté et al. (2022). Therefore, it is still difficult to assess the usefulness of RC over a large period and using many time series as inputs.

RC can be viewed as an extension of penalized linear regression, where inputs undergo processing by a reservoir, introducing the capacity for memory and non-linear combinations. Given the effectiveness of penalized linear regression in COVID-19 forecasting, as highlighted in Ferté et al. (2022), and the promising results exhibited by RC in epidemic forecasting, as demonstrated in studies such as Ghosh et al. (2021), Liu et al. (2023), and Ray, Chakraborty, and Ghosh (2021), we have opted to employ RC for the prediction of hospitalizations at 14 days at the University Hospital of Bordeaux.

The primary aim of this study is to assess the performance of RC in this forecasting task. Secondary objectives include (i) comparing the performance of RC with that of elastic-net penalized regression (identified as the optimal model in Ferté et al. (2022)) and (ii) evaluating variations in RC performance based on different architectural choices, such as the connection between the input layer and the output layer, and the use of one input scaling per feature versus a common input scaling.

## 4.2 Methods

### 4.2.1 Data

The study utilized aggregated data spanning from May 16, 2020, to January 17, 2022, regarding the COVID-19 epidemic in France, drawing from various sources to enhance forecasting accuracy. These sources encompassed epidemiological statistics from Santé Publique France, weather data from the National Oceanic and Atmospheric Administration (NOAA), both providing department-level data (Smith, Lott, and Vose 2011; Etalab 2020) and Electronic Health Record (EHR) data from the Bordeaux Hospital providing hospital-level data. All data were daily updated. Santé Publique France data included information on hospitalizations, RT-PCR tests, positive RT-PCR results, variant prevalence, and vaccination data, categorized by age groups. NOAA data contributed temperature, wind speed, humidity, and dew point data, allowing for the computation of the COVID-19 Climate Transmissibility Predict Index (Roumagnac et al. 2021). EHRs data included hospitalizations, ICU admissions, ambulance service records, and emergency unit notes, with relevant COVID-19-related concepts extracted from the notes. Data are discussed more in depth in Ferté et al. (2022).

First derivative over the last 7 days were computed to enrich model information. To take into account measurement error and daily noise variation, data were smoothed using a local polynomial regression with a span of 21 days. As previously described, input features were scaled between -1 and 1 by dividing the observed value by the maximum of the absolute value of the given input feature.

All data are publicly available. Weather data can be obtained from Smith, Lott, and Vose (2011) using R package `worldmet` (Carslaw 2023). Vaccine data can be downloaded from Etalab (2020). EHRs data can be downloaded on dryad (Ferté et al. 2023). For privacy issues, publicly available EHRs data

326 below 10 patients were obfuscated to 0. For convenience, all data were downloaded, merged and  
327 provided as replication material.

### 328 4.2.2 Evaluation framework

329 The task was to forecast 14 days ahead the number of hospitalized patients. As seen at Section 3.3,  
330 we will train the model to predict the variation of hospitalization, denoted as  $hosp$ , defined as  
331  $outcome_{t+14} = hosp_{t+14} - hosp_t$ . Metrics computation and visualizations will be performed on the  
332 predicted number of hospitalizations denoted as  $\widehat{hosp}_{t+14} = \widehat{outcome}_{t+14} + hosp_t$ .

333 The dataset was separated into two periods. First period from May 16, 2020 to March 1, 2021 served  
334 to identify relevant hyperparameters. Remaining data was used to evaluate the model performance.

335 The performance of the model was evaluated according to several metrics:

- 336 • the mean absolute error :  $MAE = mean(|\widehat{hosp}_{t+14} - hosp_{t+14}|)$ .
- 337 • the median relative error :  $MRE = median(|\frac{\widehat{hosp}_{t+14} - hosp_{t+14}}{hosp_{t+14}}|)$ .
- 338 • the mean absolute error to baseline :  $MAEB = mean(|\widehat{hosp}_{t+14} - hosp_{t+14}| - |hosp_t - hosp_{t+14}|)$ .
- 339 • the median relative error to baseline :  $MREB = median(|\frac{\widehat{hosp}_{t+14} - hosp_{t+14}}{hosp_t - hosp_{t+14}}|)$

340 Median was chosen over mean for  $MRE$  and  $MREB$  because those metrics tend to have extremely  
341 high values when the denominator is close to 0 (i.e when the number of hospitalized patients is close  
342 to 0 or the number of patients hospitalized at 14 days is close to the current number of hospitalized  
343 patients respectively).  $MAEB$  and  $MREB$  compare model performance to a baseline model which  
344 predicts the current number of hospitalized patients at 14 days. Those metrics help to determine the  
345 information added by the model and is a good baseline as covid-19 forecast model do not always  
346 outperform this basic forecast (Cramer et al. (2022)).

347 Because the outcome is obfuscated below 10 hospitalizations for privacy reason, we set both the  
348 outcome and the forecast to 10 when the observed value was 0 or the forecasted value was below 10  
349 when evaluating the model performance.

### 350 4.2.3 Models

351 We compared RC to elastic-net penalized regression (denoted as *Enet*). Furthermore we evaluated  
352 RC based on several architectures. First we compared RC with a single input scaling common to  
353 all features and a RC with on specific input scaling per feature. Second we compared RC where the  
354 input layer is connected to the output layer in addition to the connection between reservoir and  
355 output layer. Therefore, five models were evaluated :

- 356 • Elastic-net penalized regression denoted *Enet*
- 357 • RC with a single input scaling and no connection between input and ouput layers denoted  
*Common IS R %»% O*
- 358 • RC with a single input scaling and connection between input and ouput layers denoted *Common*  
*IS I+R %»% O*
- 359 • RC with multiple input scaling and no connection between input and ouput layers denoted  
*Multiple IS R %»% O*
- 360 • RC with multiple input scaling and connection between input and ouput layers denoted *Multiple*  
*IS I+R %»% O*

365 Because of the randomness of the reservoir, we took the median forecast of 10 reservoir on the train  
366 set to evaluate the performance of a given hyperparameter set. On the test set we aggregated the  
367 forecast of 40 reservoirs, each of them having one of the 40 best hyperparameter sets found on the

368 train set. In addition, because covid-19 hospitalization is a non-stationary process, models were  
369 re-trained everyday using all previous days. To ease computation burden, only one day over two  
370 was used to find hyperparameters on the training set.

#### 371 **4.2.4 Hyperparameter optimisation using random search**

372 RC relies mainly on 4 hyperparameters including the leaking rate (i.e “memory” parameter), spectral  
373 radius (i.e “chaoticity” parameter), input scaling (i.e “feature gain” parameter) and ridge (i.e penaliza-  
374 tion parameter). Input scaling can be either, common to all features or specific to each feature which  
375 increases the number of hyperparameter by the number of features.

376 Following the notation from `glmnet` package (Friedman, Hastie, and Tibshirani 2010), elastic-net  
377 penalized linear regression relies on two hyperparameters, lambda (i.e the penalization parameter)  
378 and alpha (i.e the compromise between lasso and ridge penalty)

379 Hyperparameter were selected in the training set (i.e before March 1, 2021) using a wrapper approach  
380 and a random search sampler using 2000 samples for each model. The sampling distribution were  
381 defined as follow :

- 382 • (RC) ridge and (Enet) lambda : log-uniform law defined between 1e-10 and 1e5
- 383 • (RC) input scaling and spectral radius : log-uniform law defined between 1e-5 and 1e5
- 384 • (RC) leaking rate : log-uniform law defined between 1e-3 and 1
- 385 • (Enet) alpha : uniform defined between 0 and 1

386 We provided large search space for all hyper-parameters. Search space was slightly reduced for  
387 leaking rate based on previous results and because a leaking rate of 1e-3 already imply that new  
388 inputs make the reservoir change really slowly which is not inline with the dynamic of covid-19 but  
389 would be appropriate for an application where the phenomena to forecast has a slow dynamic.

390 Finally, we provided an additional Enet model similar to the one in Ferté et al. (2022) where alpha  
391 was set to 0.5 and lambda was re-evaluated everyday in the test set based on previous data using the  
392 cross-validation procedure provided by `glmnet`.

### 393 **4.3 Results**

394 The goal of this task is to predict 14 days ahead the hospitalization. Figure 11 shows both the training  
395 set (i.e before 2021-03-01) and the test set where the blue curve correspond to the input features (first  
396 derivatives are not shown) and the orange curves correspond to the outcome the model is trained  
397 on (i.e the hospitalization variation) and the hospitalizations at 14 days on which the performance  
398 metrics are computed. The figures outline that the relation between the input features and the  
399 outcome evolve over time and that the time series is not stationary. For instance IPTCC (*Index*  
400 *PREDICT de Transmissivité Climatique de la COVID-19*) seems correlated to the outcome except that  
401 it completely miss the summer 2021 increase.

#### 402 **4.3.1 Hyperparameter selection**

403 Figure 12 shows the hyperparameter optimisation using random search for the different RC architec-  
404 tures. We observe that model with multiple input scaling achieved better performance on the train  
405 set compared to model with single input scaling which is expected as they can adapt more closely to  
406 the data thanks to specific input scaling for each feature.

407 As expected, we observe that the optimal leaking rate is above 1e-2 for all RC which is coherent with  
408 the short term dynamic of covid-19 epidemic. Trend for other hyperparameter are less clear even

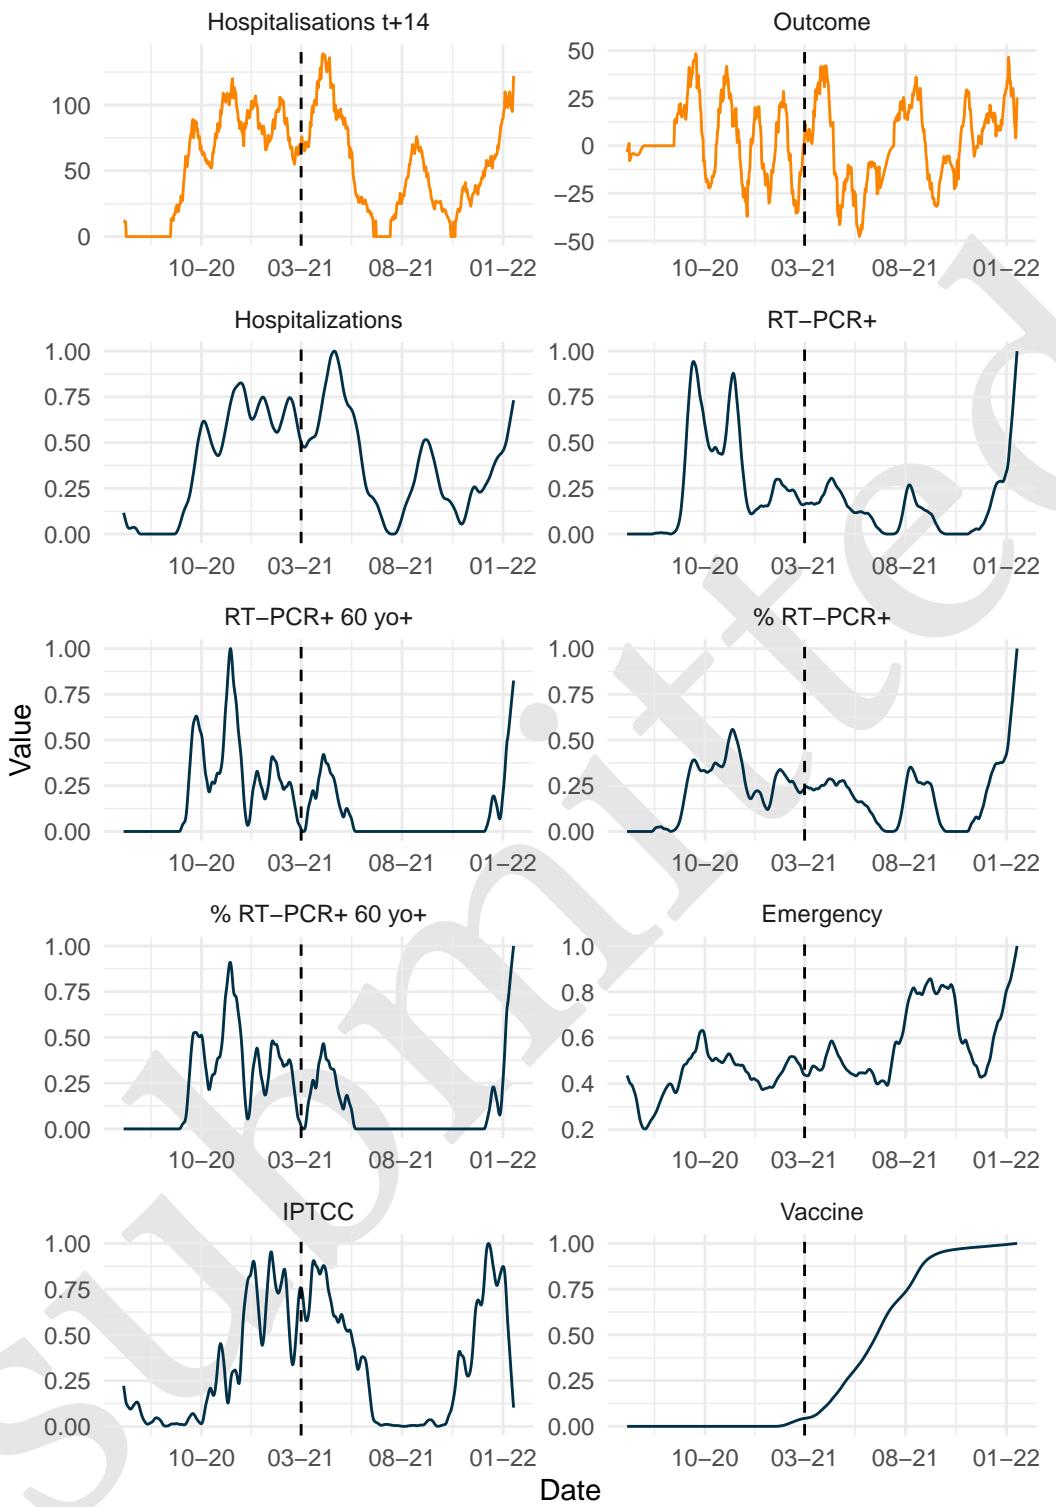


Figure 11: Covid-19 epidemic at BUH. Outcome of interest is presented in orange. Model is trained to forecast Outcome curve which corresponds to the difference between Hospitalisations at 14 days and current hospitalisations. Other features are scaled (divide by the maximum of the feature) represented in darkblue.

Table 1: Model performance with several reservoir configuration. For each setting, 40 reservoirs are computed and the forecast is the median of the 40 forecasts. Results show the performance metrics : MAE = Mean Absolute Error, MRE = Median Relative Error, MAEB = Mean Absolute Error to Baseline, MREB = Median Relative Error to Baseline.

Model Performance

Model	MAE	MRE	MAEB	MREB
Common IS R %»% O	15.23	0.26	-3.50	0.85
Common IS I + R %»% O	14.84	0.26	-3.89	0.83
Multiple IS R %»% O	15.38	0.28	-3.35	0.82
Multiple IS I + R %»% O	15.25	0.28	-3.49	0.83
Elastic-net	16.40	0.29	-2.34	0.93

409 though best hyperparameter sets were close for RC with common input scaling and for RC with  
410 multiple input scaling.

411 Figure 13 shows the hyperparameter search for RC with multiple input scaling and connected input  
412 layer. We observe that the random search tend to favor high importance given to derivative of  
413 positive RT-PCR (including the elderly) and the derivative of IPTCC. The rest of the feature do not  
414 show clear pattern.

#### 415 4.3.2 Forecast performance

416 Table 1 shows the performance on the test set. Best model according to all metrics was RC with  
417 common input scaling and connection between input and output layers. Having one input scaling per  
418 feature did not improve the model which might be due to low generalisability of the hyperparameter  
419 of the training set to the test set due to non-stationarity. Additionally, connecting input layer to  
420 output layer improved the model forecast. All RC models performed better than the elastic-net  
421 model.

422 Figure 14 shows the forecast of the different models. We note that models struggle to accurately  
423 forecast slope shifts. For instance, summer 2021 initial increase is partially predicted by all models  
424 but its decrease is not well predicted. Winter 2021 increase is anticipated by all models but they tend  
425 to overestimate it because of the rise of vaccine effect.

#### 426 4.3.3 Number of model to aggregate

427 Figure 15 show the individual forecast for the 40 best sets of hyperparameters of each RC architecture.  
428 Due to the internal random connection of the reservoir, we observe forecast stochasticity and relying  
429 on only one forecast is unreliable. We explored the number of model needed at Figure 16 which  
430 shows that after 10 models, forecast is stable and even 5 models for the simpler model with common  
431 input scaling which rely on less hyperparamters.

#### 432 4.3.4 Input feature importance

433 We compared the coefficients of the output layer estimated for the input layer and the reservoir  
434 nodes. Additionally, we compared the coefficient given to the input layer by the output layer in the  
435 reservoir and the coefficient estimated by the elastic-net model.

436 Figure 17 illustrates the ranking of input layer compared to all connections to the output layer,  
437 including the 500 reservoir nodes and the 16 features of the input layer (excluding bias). The figure

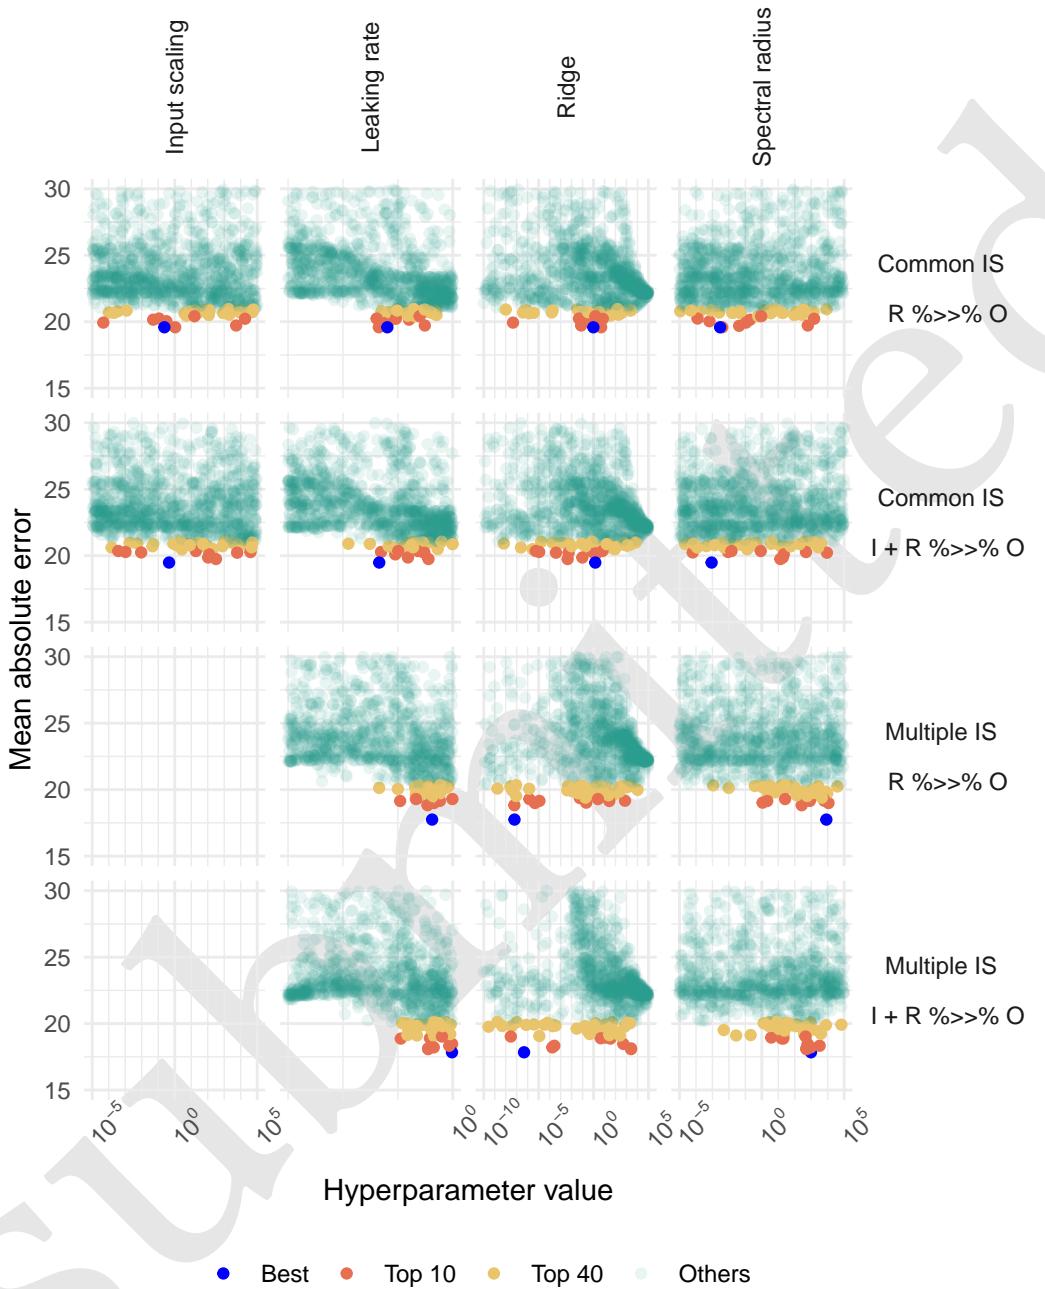


Figure 12: Hyperparameter evaluation on training set by random search. Hp sets with MAE above 30 were removed for clarity of visualisation.

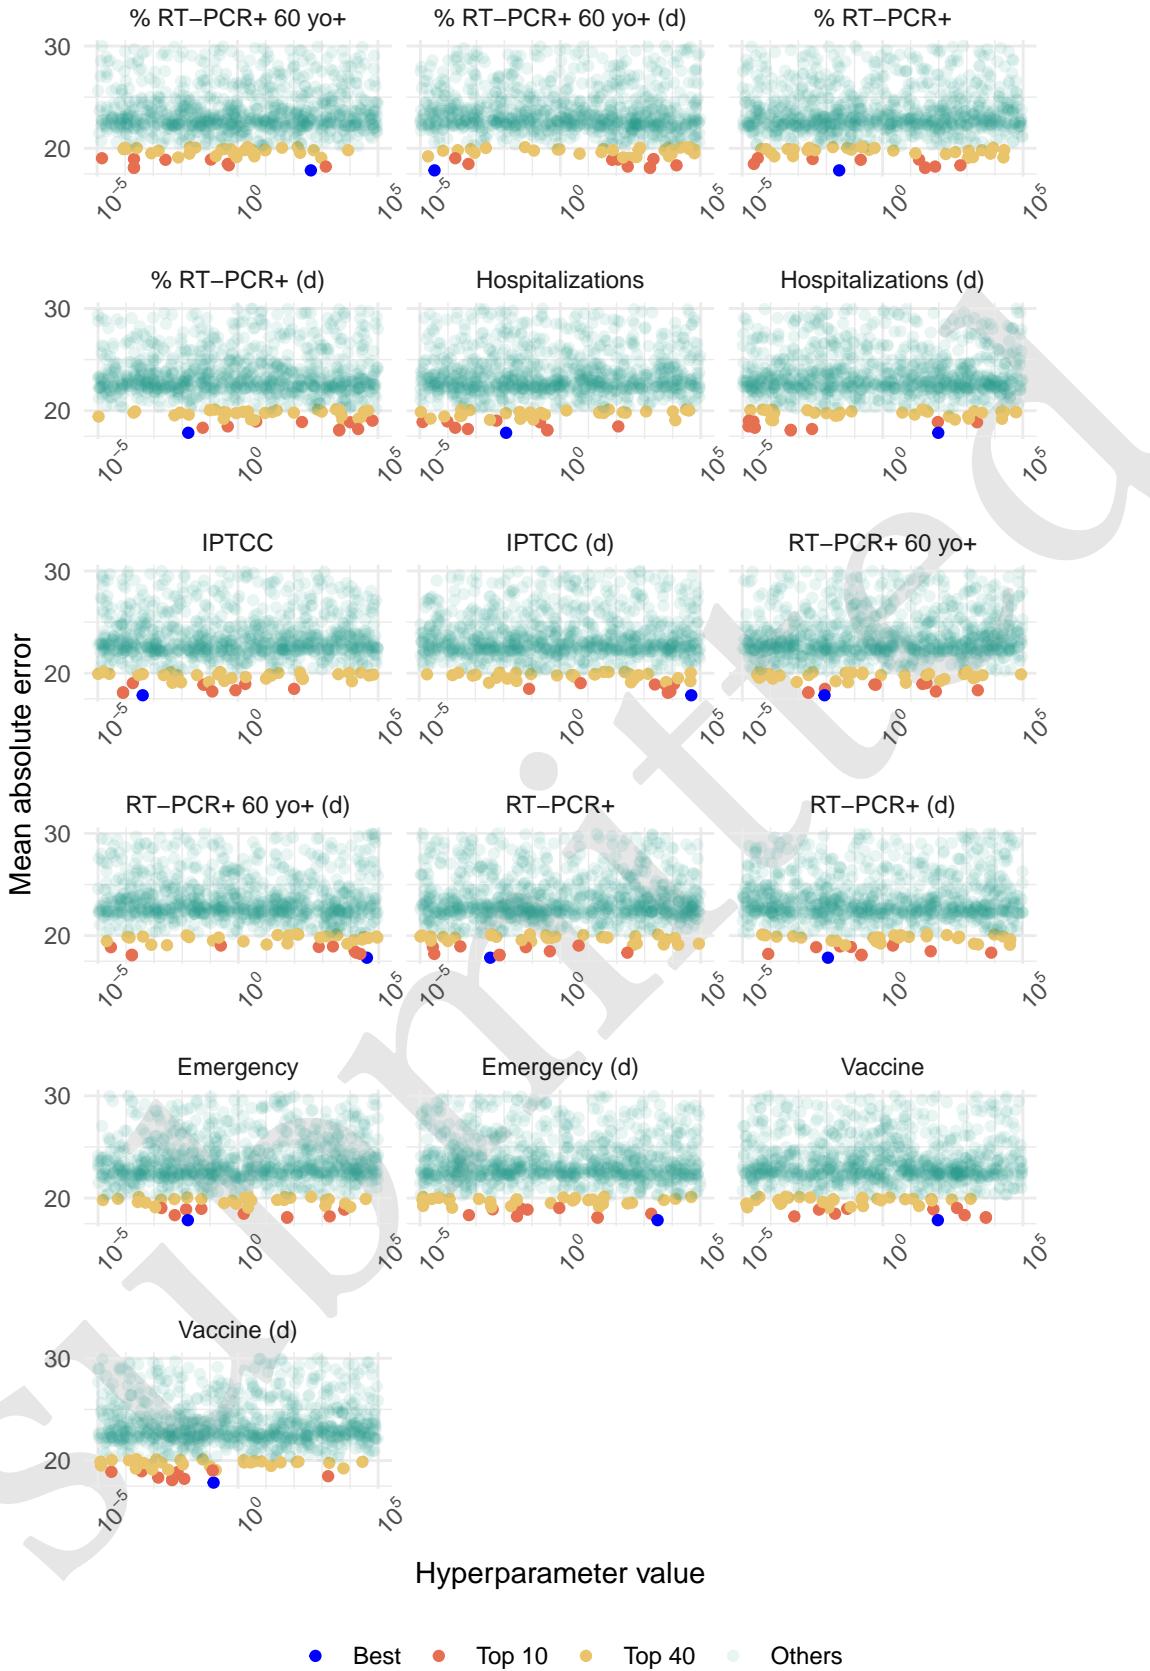


Figure 13: Hyperparameter evaluation on training set by random search of the model with multiple input scaling and no connection between input layer and output layer. Hp sets with MAE above 30 were removed for clarity of visualisation.

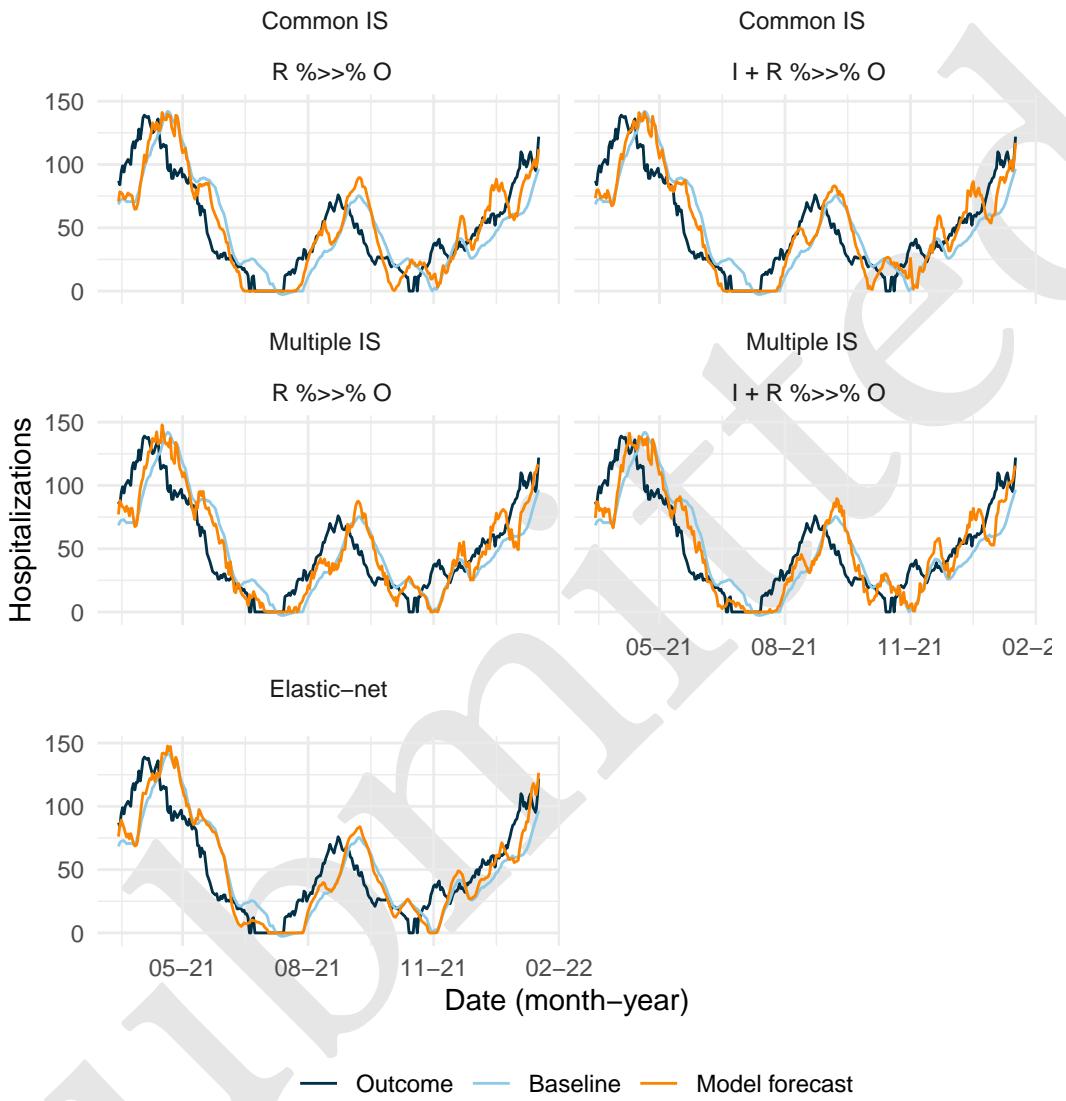


Figure 14: Reservoir computing forecast depending on the setting with and without monthly update. Red line is the median forecast of 40 reservoirs. Grey lines are individual forecast of each of the 40 reservoirs.

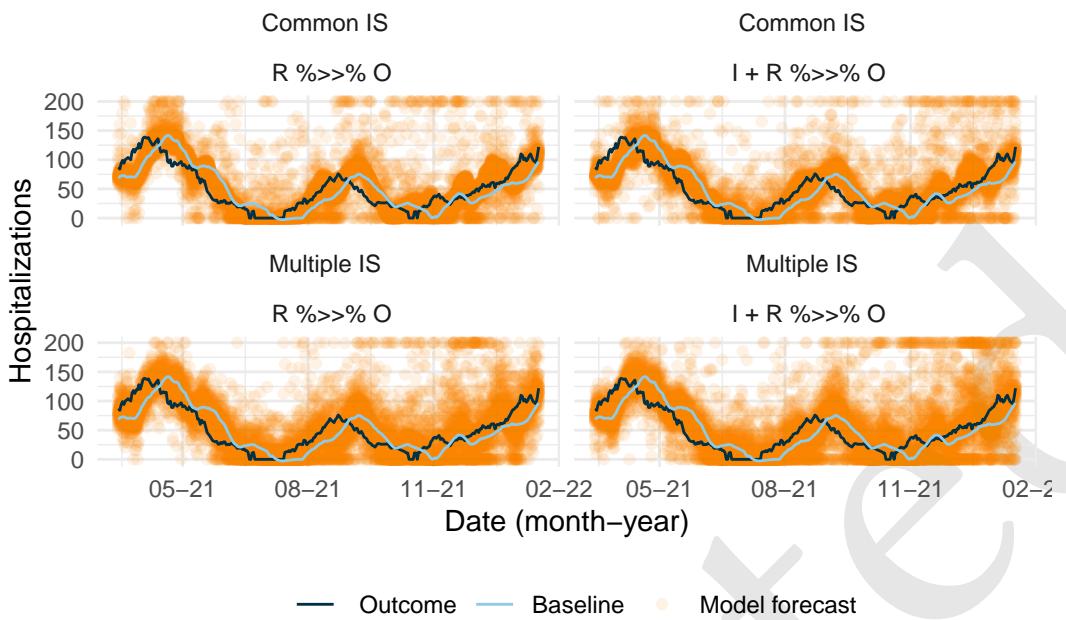


Figure 15: Individual forecast the 40 best hyperparameter sets for the different RC configuration. Forecast value above 200 were set to 200 for clarity.

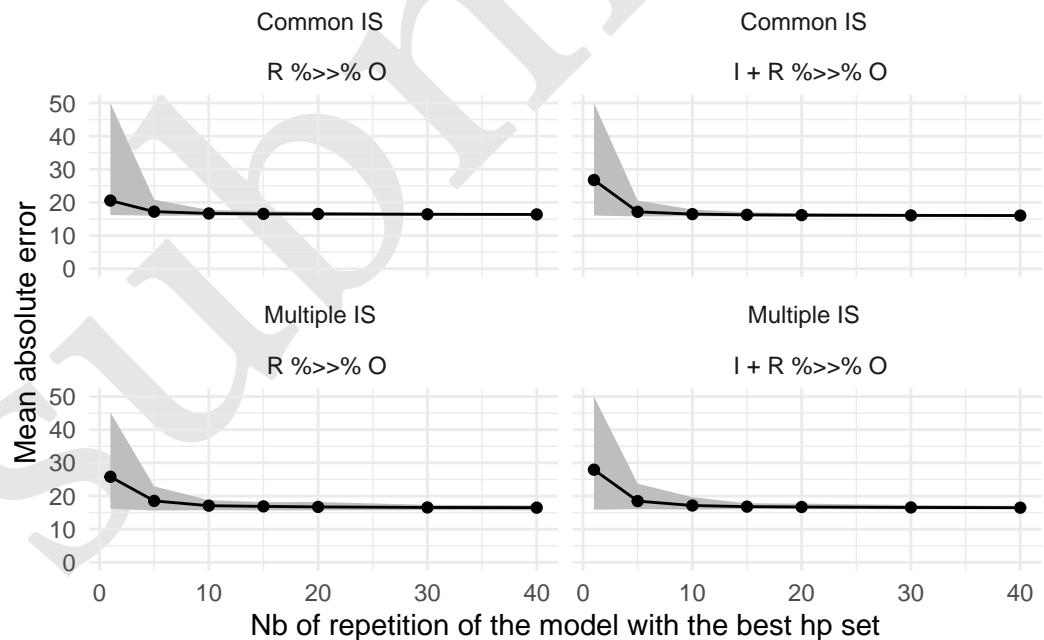


Figure 16: Mean absolute error depending on the number of aggregated reservoir.

438 shows that the model with common input scaling tends to assign less weight to input layer compared  
439 to the model with multiple input scaling. This suggests that the reservoir with common input scaling  
440 provides more information than the reservoir with multiple input scaling, which aligns with its better  
441 performance, as shown in Table Table 1.

442 Furthermore, Figure 18 compares the coefficients assigned to input features by the elastic-net model  
443 and the RC models. While the coefficients are generally consistent across RC models, there are  
444 some notable differences with elastic-net. Specifically, certain features deemed important by the  
445 elastic-net model, such as the derivative of RT-PCR, and the derivative of Vaccine, are less important  
446 for the reservoir computing model. This may indicate that these features predictive ability is better  
447 conveyed by their relationship with other features, which is captured by the reservoir computing  
448 model but might not be by the elastic-net model. Conversely, emergency, IPTCC, proportion of  
449 positive RT-PCR, and hospitalizations are more important for the reservoir computing model than  
450 for the elastic-net model.

#### 451 **4.4 Discussion**

452 In this specific application, we have demonstrated that RC exhibits commendable performance in  
453 comparison to Elastic-net, which serves as the reference model. Furthermore, we highlight the  
454 inherent challenges in forecasting within this context, primarily stemming from the non-stationarity  
455 of the time series.

456 All computations in this study were conducted using the `reservoirnet` package, and the entire  
457 codebase is accessible on Zenodo (Ferté et al. 2024). This R package demonstrates its efficacy in  
458 implementing various reservoir architectures, including connection between the input layer and the  
459 output layer, as well as the utilization of several input scaling, all within the context of a real-world  
460 use case.

461 Given the substantial number of hyperparameters involved, we acknowledge that random search  
462 may not be the most efficient optimization algorithm. We have retained this approach for the sake  
463 of simplicity in this tutorial paper; however, meta-heuristic approaches, particularly those utilizing  
464 evolutionary algorithms, may prove more efficient, especially when employing multiple input scaling  
465 (Bala et al. 2018).

466 This study represents a novel contribution to epidemic forecasting utilizing RC. Notably, previous  
467 literature predominantly focused on simpler problems characterized by fewer input features or  
468 shorter evaluation periods (Liu et al. 2023; Ray, Chakraborty, and Ghosh 2021; Ghosh et al. 2021).  
469 Our findings underscore the potential of this approach for future epidemics, suggesting its potential  
470 to surpass more traditional epidemiological tools while maintaining a lightweight model structure  
471 compared to other RNNs.

472 It is worth noting that all models, including the Ferté et al. (2022) models, encounter challenges  
473 in accurately anticipating slope shifts, indicating the need for further investigation. Specifically,  
474 additional work is warranted to extend the application of RC to high-dimensional settings, building  
475 upon the insights gained from models based on a more extensive set of features.

### 476 **5 Discussion and conclusion**

477 In this paper, we introduce the R package `reservoirnet`, which serves as a versatile tool for imple-  
478 menting reservoir computing based on ReservoirPy’s Python library. It offers flexibility in defining  
479 the reservoir architecture, including options for specifying connections between the input layer and  
480 the output layer, as well as variations in input scaling as demonstrated on a real-world use case.

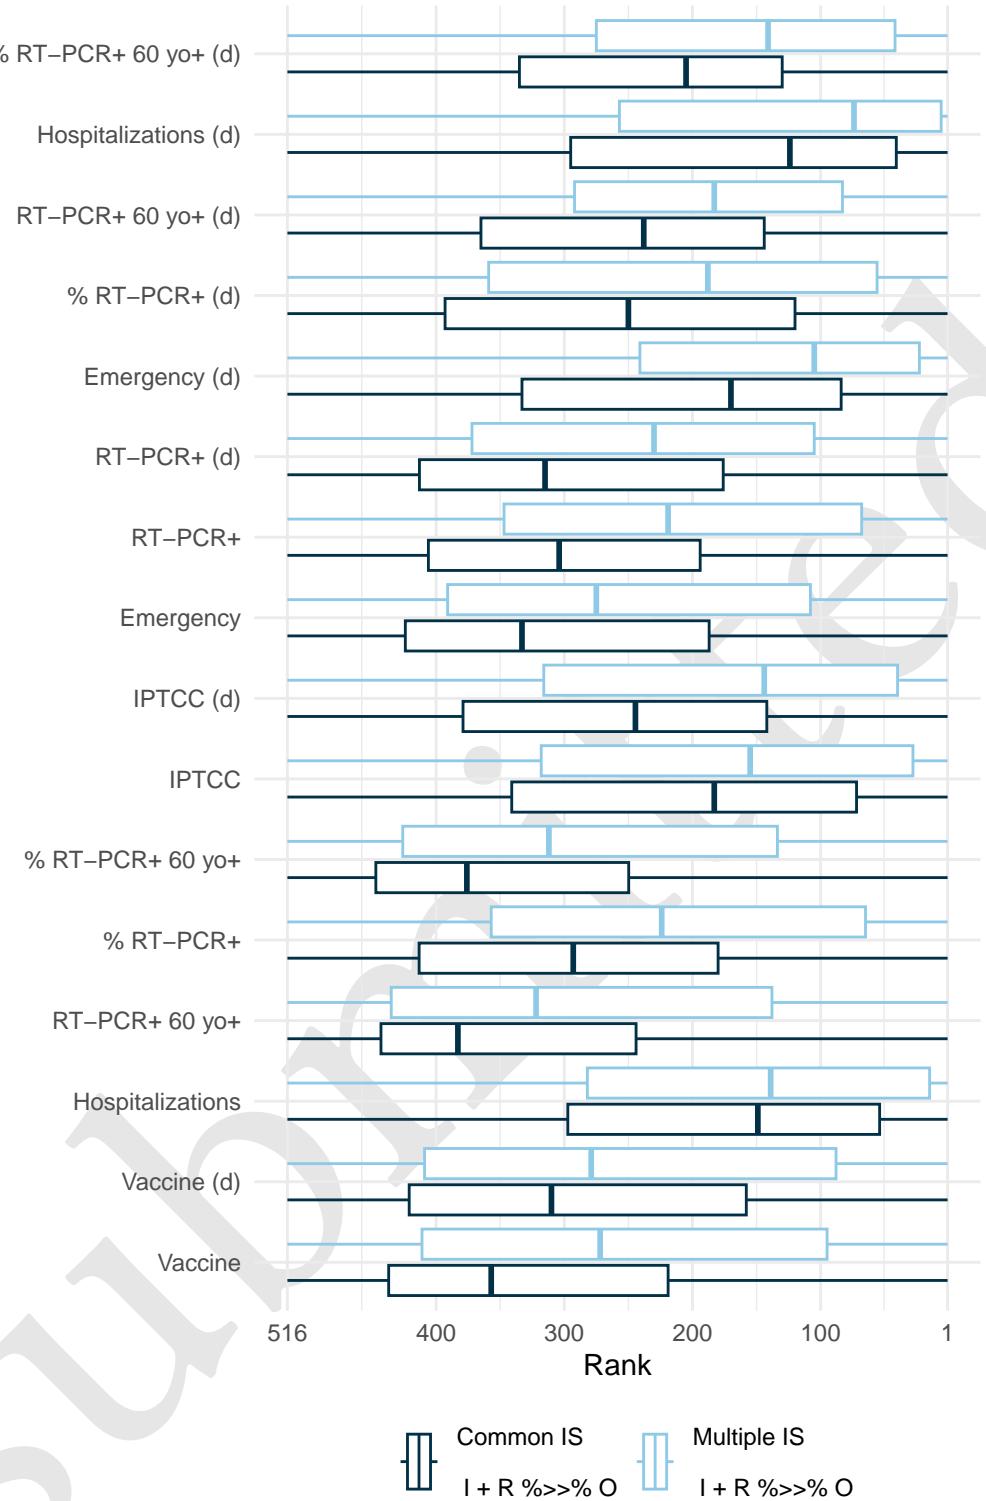


Figure 17: Mean feature importance of the 40 best hyperparameter sets by model, focus on the connection between the input and output layers. Models with direct connection between input and output layer are included. The rank is obtained by comparing the feature input layer and all other connection coefficients (both input and reservoir corresponding coefficients) attributed by the output layer at each date for each hyperparameter set. The higher the output layer's coefficient for the input layer, the closer its rank will be to 1 and the more important the feature is.

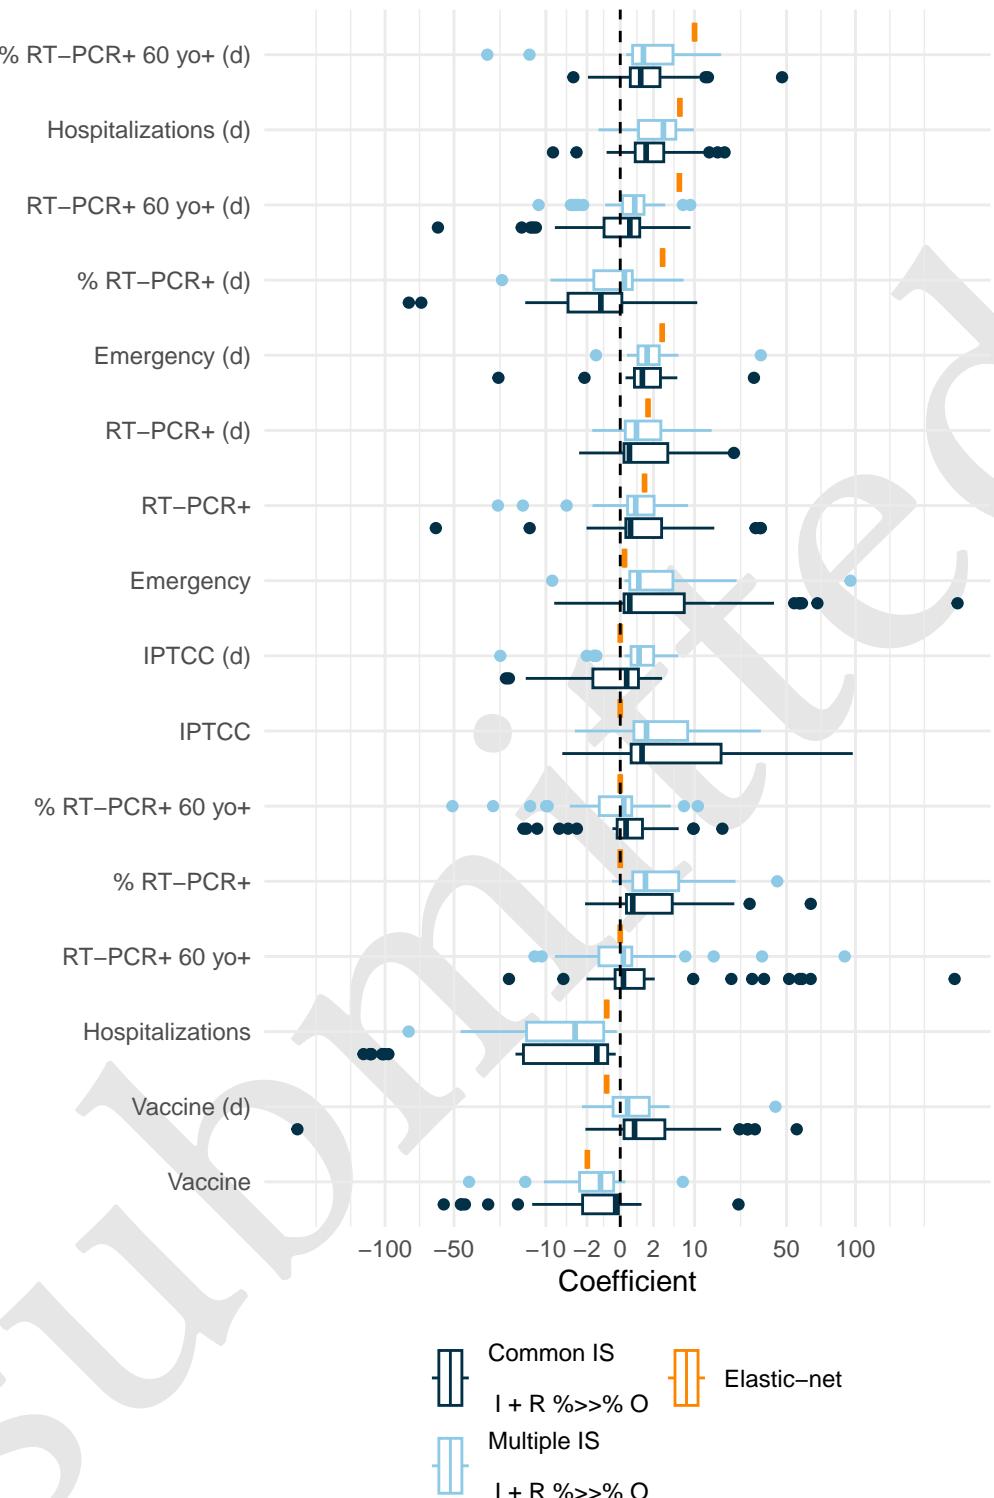


Figure 18: Mean feature coefficient of the 40 best hyperparameter sets by model and the elastic-net model. Only models with direct connection between input and output layer are included. The coefficients were calculated as the average value across all dates for each feature, model and hyperparameter set.

481 We provided a comprehensive overview of the basic usage of the `reservoirnet` package through  
482 illustrative examples in regression and classification tasks. This introductory section serves as  
483 a foundation for R users, offering step-by-step guidance on constructing and training reservoir  
484 computing models using the package. By demonstrating the application of RC in both regression  
485 and classification scenarios, we aim to equip users with the essential knowledge and skills needed to  
486 harness the capabilities of reservoir computing for diverse tasks.

487 Drawing on the robust foundation of the `ReservoirPy` structure, a well-maintained Python library,  
488 this package inherits its reliability and longevity. We have focused on providing access to the  
489 fundamental features, building upon the strong base provided by `ReservoirPy`. Therefore, this initial  
490 version of `reservoirnet` must evolve in tandem with the growing understanding and adoption of  
491 RC within the R community.

## 492 References

- 493 Bala, Abubakar, Idris Ismail, Rosdiazli Ibrahim, and Sadiq M. Sait. 2018. “Applications of Meta-  
494 heuristics in Reservoir Computing Techniques: A Review.” *IEEE Access* 6: 58012–29. <https://doi.org/10.1109/ACCESS.2018.2873770>.
- 495 Carrat, Fabrice, Julie Figoni, Joseph Henny, Jean-Claude Desenclos, Sofiane Kab, Xavier de Lamballerie,  
496 and Marie Zins. 2021. “Evidence of Early Circulation of SARS-CoV-2 in France: Findings from  
497 the Population-Based ‘CONSTANCES’ Cohort.” *European Journal of Epidemiology*, February, 1–4.  
498 <https://doi.org/10.1007/s10654-020-00716-2>.
- 499 Carslaw, David. 2023. “Worldmet: Import Surface Meteorological Data from NOAA Integrated  
500 Surface Database (ISD).” <https://cran.r-project.org/web/packages/worldmet/index.html>.
- 501 Carvalho, Kathleen, João Paulo Vicente, Mihajlo Jakovljevic, and João Paulo Ramos Teixeira. 2021.  
502 “Analysis and Forecasting Incidence, Intensive Care Unit Admissions, and Projected Mortality  
503 Attributable to COVID-19 in Portugal, the UK, Germany, Italy, and France: Predictions for 4  
504 Weeks Ahead.” *Bioengineering* 8 (6): 84. <https://doi.org/10.3390/bioengineering8060084>.
- 505 COVID-19 Cumulative Infection Collaborators. 2022. “Estimating Global, Regional, and National  
506 Daily and Cumulative Infections with SARS-CoV-2 Through Nov 14, 2021: A Statistical Analysis.”  
507 *Lancet*. [https://doi.org/10.1016/S0140-6736\(22\)00484-6](https://doi.org/10.1016/S0140-6736(22)00484-6).
- 508 Cramer, Estee Y., Evan L. Ray, Velma K. Lopez, Johannes Bracher, Andrea Brennen, and et al.  
509 2022. “Evaluation of Individual and Ensemble Probabilistic Forecasts of COVID-19 Mortality  
510 in the United States.” *Proceedings of the National Academy of Sciences* 119 (15): e2113561119.  
511 <https://doi.org/10.1073/pnas.2113561119>.
- 512 Etalab. 2020. “Les Données Relatives Au COVID-19 En France - Data.gouv.fr.” <https://www.data.gouv.fr/fr/pages/donnees-coronavirus/>.
- 513 Ferté, Thomas, Kalidou Ba, Dan Dutartre, Pierrick Legrand, Vianney Jouhet, Romain Griffier, Rodolphe  
514 Thiébaut, Xavier Hinaut, and Boris P. Hejblum. 2024. “Thomasferte/Jss\_reservoirnet: First  
515 Release.” Zenodo. <https://doi.org/10.5281/ZENODO.11281341>.
- 516 Ferté, Thomas, Vianney Jouhet, Romain Griffier, Boris P. Hejblum, Rodolphe Thiébaut, and Bordeaux  
517 University Hospital Covid-19 Crisis Task Force. 2022. “The Benefit of Augmenting Open Data  
518 with Clinical Data-Warehouse EHR for Forecasting SARS-CoV-2 Hospitalizations in Bordeaux  
519 Area, France.” *JAMIA Open* 5 (4): ooac086. <https://doi.org/10.1093/jamiaopen/ooac086>.
- 520 Ferté, Thomas, Vianney Jouhet, Romain Griffier, Boris Hejblum, Rodolphe Thiébaut, and Bordeaux  
521 University Hospital Covid-19 Crisis Task Force. 2023. “The Benefit of Augmenting Open Data  
522 with Clinical Data-Warehouse EHR for Forecasting SARS-CoV-2 Hospitalizations in Bordeaux  
523 Area, France.” Dryad. <https://doi.org/10.5061/DRYAD.HHMGQNKX>.
- 524 Friedman, Jerome, Trevor Hastie, and Rob Tibshirani. 2010. “Regularization Paths for Generalized  
525 Linear Models via Coordinate Descent.” *Journal of Statistical Software* 33 (1): 1–22.
- 526

- 528 Ghosh, Subrata, Abhishek Senapati, Arindam Mishra, Joydev Chattopadhyay, Syamal K. Dana,  
529 Chittaranjan Hens, and Dibakar Ghosh. 2021. “Reservoir Computing on Epidemic Spreading:  
530 A Case Study on COVID-19 Cases.” *Physical Review E* 104 (1): 014308. <https://doi.org/10.1103/PhysRevE.104.014308>.
- 532 Hinaut, Xavier, and Peter Ford Dominey. 2013. “Real-Time Parallel Processing of Grammatical  
533 Structure in the Fronto-Striatal System: A Recurrent Network Simulation Study Using Reservoir  
534 Computing.” *PLOS ONE* 8 (2): e52946. <https://doi.org/10.1371/journal.pone.0052946>.
- 535 Hübner, Martin, Tobias Zingg, David Martin, Philippe Eckert, and Nicolas Demartines. 2020. “Surgery  
536 for Non-Covid-19 Patients During the Pandemic.” *PLoS ONE* 15 (10): e0241331. <https://doi.org/10.1371/journal.pone.0241331>.
- 538 Jaeger, Herbert. 2001. “The” Echo State” Approach to Analysing and Training Recurrent Neu-  
539 ral Networks-with an Erratum Note’.” *Bonn, Germany: German National Research Center for  
540 Information Technology GMD Technical Report* 148 (January).
- 541 Kim, Gina, Mengru Wang, Hanh Pan, Giana H. Davidson, Alison C. Roxby, Jen Neukirch, Danna Lei,  
542 Elicia Hawken-Dennis, Louise Simpson, and Thuan D. Ong. 2020. “A Health System Response to  
543 COVID-19 in Long-Term Care and Post-Acute Care: A Three-Phase Approach.” *Journal of the  
544 American Geriatrics Society* 68 (6): 1155–61. <https://doi.org/10.1111/jgs.16513>.
- 545 Kmet, Tibor, and Maria Kmetova. 2019. “Bézier Curve Parametrisation and Echo State Network  
546 Methods for Solving Optimal Control Problems of SIR Model.” *Biosystems* 186 (December): 104029.  
547 <https://doi.org/10.1016/j.biosystems.2019.104029>.
- 548 Kudo, Mineichi, Jun Toyama, and Masaru Shimbo. 1999. “Multidimensional Curve Classification  
549 Using Passing-Through Regions.” *Pattern Recognition Letters* 20 (11): 1103–11. [https://doi.org/10.1016/S0167-8655\(99\)00077-X](https://doi.org/10.1016/S0167-8655(99)00077-X).
- 551 Liu, Bocheng, Yiyuan Xie, Weichen Liu, Xiao Jiang, Yichen Ye, Tingting Song, Junxiong Chai, Manying  
552 Feng, and Haodong Yuan. 2023. “Nanophotonic Reservoir Computing for COVID-19 Pandemic  
553 Forecasting.” *Nonlinear Dynamics* 111 (7): 6895–6914. <https://doi.org/10.1007/s11071-022-08190-z>.
- 554 Lukoševičius, Mantas, and Herbert Jaeger. 2009. “Reservoir Computing Approaches to Recurrent  
555 Neural Network Training.” *Computer Science Review* 3 (3): 127–49. <https://doi.org/10.1016/j.cosrev.2009.03.005>.
- 557 Maass, Wolfgang, Thomas Natschläger, and Henry Markram. 2002. “Real-Time Computing Without  
558 Stable States: A New Framework for Neural Computation Based on Perturbations.” *Neural  
559 Computation* 14 (11): 2531–60. <https://doi.org/10.1162/089976602760407955>.
- 560 Martinuzzi, Francesco, Chris Rackauckas, Anas Abdelrehim, Miguel D. Mahecha, and Karin Mora.  
561 2022. “ReservoirComputing.jl: An Efficient and Modular Library for Reservoir Computing  
562 Models.” *Journal of Machine Learning Research* 23 (288): 1–8. <http://jmlr.org/papers/v23/22-0611.html>.
- 564 Mohimont, Lucas, Amine Chemchem, François Alin, Michaël Krajecki, and Luiz Angelo Steffenel.  
565 2021. “Convolutional Neural Networks and Temporal CNNs for COVID-19 Forecasting in France.”  
566 *Applied Intelligence*, April. <https://doi.org/10.1007/s10489-021-02359-6>.
- 567 Nakane, Ryosho, Gouhei Tanaka, and Akira Hirose. 2018. “Reservoir Computing With Spin Waves  
568 Excited in a Garnet Film.” *IEEE Access PP* (January): 1–1. <https://doi.org/10.1109/ACCESS.2018.2794584>.
- 570 Paireau, Juliette, Alessio Andronico, Nathanaël Hozé, Maylis Layen, Pascal Crépey, Alix Roumagnac,  
571 Marc Lavielle, Pierre-Yves Boëlle, and Simon Cauchemez. 2022. “An Ensemble Model Based  
572 on Early Predictors to Forecast COVID-19 Health Care Demand in France.” *Proceedings of the  
573 National Academy of Sciences* 119 (18): e2103302119. <https://doi.org/10.1073/pnas.2103302119>.
- 574 Penkovsky, Bogdan, Laurent Larger, and Daniel Brunner. 2018. “Efficient Design of Hardware-  
575 Enabled Reservoir Computing in FPGAs.” *Journal of Applied Physics* 124 (16): 162101. <https://doi.org/10.1063/1.5039826>.
- 577 Pottier, Loïc. 2021. “Forecast of the Covid19 Epidemic in France.” *medRxiv*. <https://doi.org/10.1101/>

- 578 2021.04.13.21255418.
- 579 Prychynenko, Diana, Matthias Sitte, Kai Litzius, Benjamin Krüger, George Bourianoff, Mathias Kläui,  
580 Jairo Sinova, and Karin Everschor-Sitte. 2018. “Magnetic Skyrmion as a Nonlinear Resistive  
581 Element: A Potential Building Block for Reservoir Computing.” *Physical Review Applied* 9 (1):  
582 014034. <https://doi.org/10.1103/PhysRevApplied.9.014034>.
- 583 Rafayelyan, Mushegh, Jonathan Dong, Yongqi Tan, Florent Krzakala, and Sylvain Gigan. 2020. “Large-  
584 Scale Optical Reservoir Computing for Spatiotemporal Chaotic Systems Prediction.” *Physical  
585 Review X* 10 (4): 041037. <https://doi.org/10.1103/PhysRevX.10.041037>.
- 586 Rahimi, Iman, Fang Chen, and Amir H. Gandomi. 2021. “A Review on COVID-19 Forecasting Models.”  
587 *Neural Computing & Applications*, February, 1–11. <https://doi.org/10.1007/s00521-020-05626-8>.
- 588 Ray, Arnob, Tanujit Chakraborty, and Dibakar Ghosh. 2021. “Optimized Ensemble Deep Learning  
589 Framework for Scalable Forecasting of Dynamics Containing Extreme Events.” *Chaos (Woodbury,  
590 N.Y.)* 31 (11): 111105. <https://doi.org/10.1063/5.0074213>.
- 591 Roumagnac, Alix, Eurico de Carvalho Filho, Raphaël Bertrand, Anne-Kim Banchereau, and Guillaume  
592 Lahache. 2021. “Étude de l'influence Potentielle de l'humidité Et de La Température Dans La  
593 Propagation de La Pandémie COVID-19.” *Médecine de Catastrophe - Urgences Collectives, Douleur  
594 et situations d'exceptionPandémie COVID-19*, 5 (1): 87–102. <https://doi.org/10.1016/j.pxur.2021.01.002>.
- 595 Simões, Jorge, João Paulo Moreira Magalhães, André Biscaia, António da Luz Pereira, Gonçalo  
596 Figueiredo Augusto, and Inês Fronteira. 2021. “Organisation of the State, Model of Health System  
597 and COVID-19 Health Outcomes in Six European Countries, During the First Months of the  
598 COVID-19 Epidemic in 2020.” *The International Journal of Health Planning and Management*, June,  
599 10.1002/hpm.3271. <https://doi.org/10.1002/hpm.3271>.
- 600 Smith, Adam, Neal Lott, and Russ Vose. 2011. “The Integrated Surface Database: Recent Developments  
601 and Partnerships.” *Bulletin of the American Meteorological Society* 92 (6): 704–8. <https://doi.org/10.1175/2011BAMS3015.1>.
- 601 Tanaka, Gouhei, Toshiyuki Yamane, Jean Benoit Héroux, Ryosho Nakane, Naoki Kanazawa, Seiji  
602 Takeda, Hidetoshi Numata, Daiju Nakano, and Akira Hirose. 2019. “Recent Advances in Physical  
603 Reservoir Computing: A Review.” *Neural Networks* 115 (July): 100–123. <https://doi.org/10.1016/j.neunet.2019.03.005>.
- 604 Trouvain, Nathan, and Xavier Hinaut. 2021. “Canary Song Decoder: Transduction and Implicit  
605 Segmentation with ESNs and LTSMs.” In *Artificial Neural Networks and Machine Learning – ICANN  
606 2021*, edited by Igor Farkaš, Paolo Masulli, Sebastian Otte, and Stefan Wermter, 71–82. Lecture  
607 Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-030-86383-8\\_6](https://doi.org/10.1007/978-3-030-86383-8_6).
- 608 ——. 2022. “Reservoirpy: A Simple and Flexible Reservoir Computing Tool in Python.” <https://inria.hal.science/hal-03699931>.
- 609 Trouvain, Nathan, Luca Pedrelli, Thanh Trung Dinh, and Xavier Hinaut. 2020. “ReservoirPy:  
610 An Efficient and User-Friendly Library to Design Echo State Networks.” In *Artificial Neural  
611 Networks and Machine Learning – ICANN 2020*, 494–505. Springer International Publishing.  
612 <https://inria.hal.science/hal-02595026>.
- 613 Trouvain, Nathan, Nicolas Rougier, and Xavier Hinaut. 2022. “Create Efficient and Complex Reservoir  
614 Computing Architectures with ReservoirPy.” In *From Animals to Animats 16*, edited by Lola  
615 Cañamero, Philippe Gaussier, Myra Wilson, Sofiane Boucenna, and Nicolas Cuperlier, 91–102.  
616 Lecture Notes in Computer Science. Cham: Springer International Publishing. [https://doi.org/10.1007/978-3-031-16770-6\\_8](https://doi.org/10.1007/978-3-031-16770-6_8).
- 617 Ushey, Kevin, JJ Allaire, and Yuan Tang. 2024. *Reticulate: Interface to 'Python'*. <https://rstudio.github.io/reticulate/>.
- 618 Vlachas, P. R., J. Pathak, B. R. Hunt, T. P. Sapsis, M. Girvan, E. Ott, and P. Koumoutsakos. 2020.  
619 “Backpropagation Algorithms and Reservoir Computing in Recurrent Neural Networks for the

628 Forecasting of Complex Spatiotemporal Dynamics.” *Neural Networks* 126 (June): 191–217. <https://doi.org/10.1016/j.neunet.2020.02.016>.  
 629 Wickham, Hadley, Romain François, Lionel Henry, Kirill Müller, and Davis Vaughan. 2023. “dplyr: A  
 630 Grammar of Data Manipulation.” <https://CRAN.R-project.org/package=dplyr>.  
 631 Wickham, Hadley, Danielle Navarro, and Thomas Lin Pedersen. 2018. *ggplot2: Elegant Graphics for*  
 632 *Data Analysis (3e)*. 3rd ed. Springer-Verlag New York. <https://ggplot2-book.org/>.  
 633 World Health Organisation. 2020. “WHO Coronavirus (COVID-19) Dashboard.” <https://covid19.who.int>.  
 634 Zhang, Qihuang, Grace Y. Yi, Li-Pang Chen, and Wenqing He. 2023. “Sentiment Analysis and  
 635 Causal Learning of COVID-19 Tweets Prior to the Rollout of Vaccines.” *PLoS One* 18 (2): e0277878.  
 636 <https://doi.org/10.1371/journal.pone.0277878>.

## 639 Session information

```

640 R version 4.4.1 (2024-06-14)
641 Platform: x86_64-pc-linux-gnu
642 Running under: Ubuntu 24.04.1 LTS
643
644 Matrix products: default
645 BLAS: /usr/lib/x86_64-linux-gnublas/libblas.so.3.12.0
646 LAPACK: /usr/lib/x86_64-linux-gnulapack/liblapack.so.3.12.0
647
648 locale:
649 [1] LC_CTYPE=C.UTF-8          LC_NUMERIC=C           LC_TIME=C.UTF-8
650 [4] LC_COLLATE=C.UTF-8        LC_MONETARY=C.UTF-8   LC_MESSAGES=C.UTF-8
651 [7] LC_PAPER=C.UTF-8         LC_NAME=C             LC_ADDRESS=C
652 [10] LC_TELEPHONE=C        LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C
653
654 time zone: Etc/UTC
655 tzcode source: system (glibc)
656
657 attached base packages:
658 [1] stats      graphics   grDevices datasets  utils      methods   base
659
660 other attached packages:
661 [1] reservoirnet_0.2.0 patchwork_1.3.0     ggplot2_3.5.1       dplyr_1.1.4
662
663 loaded via a namespace (and not attached):
664 [1] gt_0.11.1        utf8_1.2.4        generics_0.1.3    tidyverse_1.3.1
665 [5] renv_1.0.11      xml2_1.3.6        rstatix_0.7.2     lattice_0.20-45
666 [9] stringi_1.8.4    digest_0.6.37     magrittr_2.0.3    evaluate_1.0.1
667 [13] grid_4.4.1      timechange_0.3.0  fastmap_1.2.0    rprojroot_2.0.4
668 [17] Matrix_1.7-1    jsonlite_1.8.9    slider_0.3.2     backports_1.5.0
669 [21] brio_1.1.5      Formula_1.2-5    purrr_1.0.2       fansi_1.0.6
670 [25] scales_1.3.0    abind_1.4-8       cli_3.6.3        rlang_1.1.4
671 [29] munsell_0.5.1    withr_3.0.2       yaml_2.3.10      tools_4.4.1
672 [33] ggsignif_0.6.4   colorspace_2.1-1  ggpubr_0.6.0     here_1.0.1
673 [37] broom_1.0.7     reticulate_1.40.0  png_0.1-8       vctrs_0.6.5
674 [41] R6_2.5.1        lifecycle_1.0.4    lubridate_1.9.3   snakecase_0.11.1
675 [45] stringr_1.5.1    car_3.1-3        janitor_2.2.0    warp_0.2.1
  
```

```
676 [49] pkgconfig_2.0.3    pillar_1.9.0      gtable_0.3.6      Rcpp_1.0.13-1  
677 [53] glue_1.8.0        xfun_0.49       tibble_3.2.1      tidyselect_1.2.1  
678 [57] knitr_1.49       farver_2.1.2     htmltools_0.5.8.1 labeling_0.4.3  
679 [61] rmarkdown_2.29    carData_3.0-5    testthat_3.2.1.1 compiler_4.4.1
```

Submitted