



Reservoir Computing in R: a Tutorial for Using reservoirnet to Predict Complex Time-Series

ISSN 2824-7795

- Thomas Ferté Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria
Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux
- Kalidou Ba Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria
Dan Dutarte Inria BSO, Inria
Pierrick Legrand Inria BSO, F-33000 Bordeaux, France, Inria
IMB, Institut de Mathématiques de Bordeaux, UMR CNRS 5251, IMB
- Vianney Jouhet Inserm Bordeaux Population Health Research Center UMR 1219, team AHeAD, F-33000 Bordeaux, Inserm
Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux
- Rodolphe Thiébaut Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria
Bordeaux Hospital University Center, Pôle de santé publique, Service d'information médicale, F-33000 Bordeaux, France, CHU de Bordeaux
Xavier Hinaut Inria BSO, F-33000 Bordeaux, France, Inria
Univ. Bordeaux, CNRS, IMN, UMR 5293, Bordeaux, France, CNRS
LaBRI, Univ. Bordeaux, Bordeaux INP, CNRS UMR 5800., LaBRI
- Boris Hejblum ¹ Inserm Bordeaux Population Health Research Center UMR 1219, Inria BSO, team SISTM, F-33000 Bordeaux, France, Inserm, Inria

Date published: 2025-01-09 Last modified: 2025-01-09

Abstract

Reservoir Computing (RC) is a machine learning method based on neural networks that efficiently process information generated by dynamical systems. It has been successful in solving various tasks including time series forecasting, language processing or voice processing. RC is implemented in Python and Julia but not in R. This article introduces `reservoirnet`, an R package providing access to the Python API `ReservoirPy`, allowing R users to harness the power of reservoir computing. This article provides an introduction to the fundamentals of RC and showcases its real-world applicability through three distinct sections. First, we cover the foundational concepts of RC, setting the stage for understanding its capabilities. Next, we delve into the practical usage of `reservoirnet` through two illustrative examples. These examples demonstrate how it can be applied to real-world problems, specifically, regression of COVID-19 hospitalizations and classification of Japanese vowels. Finally, we present a comprehensive analysis of a real-world application of `reservoirnet`, where it was used to forecast COVID-19 hospitalizations at Bordeaux University Hospital using public data and electronic health records.

Keywords: Reservoir Computing, Covid-19, Electronic Health Records, Time series

¹Corresponding author: boris.hejblum@u-bordeaux.fr

1 Contents

2	1 Introduction	2
3	2 RC presentation	3
4	3 Usage workflow	5
5	3.1 Installation	5
6	3.2 Package workflow overview	5
7	3.3 Basic regression use-case	7
8	3.3.1 Covid-19 data	7
9	3.3.2 First reservoir	8
10	3.3.3 Forecast	11
11	3.4 Classification	16
12	3.4.1 The Japanese vowel dataset	16
13	3.4.2 Classification (sequence-to-vector model)	17
14	3.4.3 Transduction (sequence-to-sequence model)	19
15	4 Avanced case-study: Covid-19 hospitalizations forecast	23
16	4.1 Introduction	23
17	4.2 Methods	24
18	4.2.1 Data	24
19	4.2.2 Evaluation framework	24
20	4.2.3 Models	25
21	4.2.4 Hyperparameter optimisation using random search	25
22	4.3 Results	26
23	4.3.1 Hyperparameter selection	26
24	4.3.2 Forecast performance	26
25	4.3.3 Number of model to aggregate	30
26	4.3.4 Input feature importance	30
27	4.4 Discussion	34
28	5 Discussion and conclusion	34
29	Acknowledgements	36
30	References	36
31	Session information	39

32 **1 Introduction**

33 Reservoir Computing (RC) is a prominent machine learning method, proposed by Jaeger (2001), Maass,
34 Natschläger, and Markram (2002) and Lukoševičius and Jaeger (2009) that has gained significant
35 attention in recent years for its ability to efficiently process information generated by dynamical
36 systems. This innovative approach leverages the dynamics of a high-dimensional “reservoir” (we
37 define it below) to perform complex computations and solve various tasks based on the response
38 of this dynamical system to input signals. RC has demonstrated its efficacy in tackling various
39 challenges, encompassing pattern classification and time series forecasting in applications ranging
40 from electrocardiogram analysis to bird calls (Trouvain and Hinaut 2021), language processing
41 (Hinaut and Dominey 2013), power plants, internet traffic, stock prices, and beyond (Lukoševičius

42 and Jaeger 2009; Tanaka et al. 2019).

43 Originally, the RC paradigm was implemented in artificial firing-rate neurons (“Echo State Networks”,
44 Jaeger (2001)) and spiking neurons (“Liquid State Machine”, Maass, Natschläger, and Markram (2002))
45 as a recurrent neural network (RNN) where the internal recurrent connections, denoted as the
46 reservoir, are randomly generated and only the output layer (named “read-out”) is trained. The
47 reservoir projects temporal input signals onto a high-dimensional feature space, facilitating the
48 learning of non-linear and temporal interactions. Thus, this recurrent layer contains high-dimensional
49 non-linear recombination of the inputs and past states: it is a “reservoir of computations” from
50 which useful information can be linearly extracted (or “read-out”) to provide the desired outputs.
51 This offers the advantage of decreasing the computing time compared to conventional RNNs while
52 consistently maintaining performance (Vlachas et al. 2020). Besides, this RC paradigm fostered
53 increasing interest thanks to its ability to be implemented on classical computers, as the hidden
54 recurrent layer can be kept untrained. A wide range of physical media can be also used to replace
55 it and Tanaka et al. (2019) recently reviewed this prolific field: from FPGA hardware (Penkovsky,
56 Larger, and Brunner 2018), to spin waves using magnetic properties (Nakane, Tanaka, and Hirose
57 2018), skrymions (Prychynenko et al. 2018) or optical implementations (Rafayelyan et al. 2020).
58 This provides interesting and potentially more efficient alternative to traditional machine learning
59 computing.

60 RC leverages various hyperparameters to introduce prior knowledge about the relationship between
61 input variables and output targets. But because the connections within the reservoir are randomly
62 initialized, the same set of hyperparameters may exhibit diverse behaviors across different instances
63 of the reservoir connections. This unpredictability makes it challenging to anticipate the performance
64 of a particular hyperparameter setting, as identical settings may produce varying outcomes when
65 applied to distinct instances of the reservoir. Moreover, selecting the most suitable hyperparameters
66 often requires researchers to experiment with multiple combinations on a training dataset and
67 evaluate their performance on a separate test set². Although this approach can be resource-intensive
68 and time-consuming, it is a compromise that is acceptable considering the rapid simulation capabilities
69 offered by RC. Furthermore, there is a current absence of implementation in R, rendering the method
70 challenging for users unfamiliar with Python (Trouvain and Hinaut 2022) or Julia (Martinuzzi et al.
71 2022).

72 Here, we offer comprehensive guidance to assist new users in maximizing the benefits of RC. Initially,
73 a broad introduction to reservoir computing is presented in Section 2, followed in Section 3 by a
74 tutorial on its application using `reservoirnet`, an R package built upon the `ReservoirPy` Python
75 module (Trouvain, Rougier, and Hinaut 2022; Trouvain and Hinaut 2022; Trouvain et al. 2020).
76 Section 3 then introduces the workflow usage on `reservoirnet` for RC with two basic use-cases,
77 and finally, in Section 4 we investigate the various challenges associated with an advanced case-
78 study leveraging RC for forecasting COVID-19 hospitalizations. This case-study exploration includes
79 detailed guidance on the modeling strategy, the selection of hyperparameters, and the implementation
80 process.

81 2 RC presentation

82 RC is a machine learning paradigm which is most often implemented as Echo State Networks (ESNs),
83 i.e. the firing-rate neuron version (Jaeger 2001). An ESN is described by three matrices of connectivity:
84 an input layer W_{in} , a recurrent layer W and an output layer W_{out} . At each time step, the input vector
85 u_t is projected into the reservoir which is also combined with reservoir past state $x(t - 1)$ through

²In this article, we employ the term “train set” to refer to the combined dataset consisting of both the training and validation sets, which are cycled through in a cross-validation manner.

86 the recurrent connections. The output $y(t)$ is linearly read-out from the reservoir. Input W_{in} and
 87 recurrent W matrices are kept random; only the output matrix W_{out} is trained in an offline or online
 88 method. Often a ridge regression (i.e. a regularized linear regression) is used to obtain the desired
 89 outputs $y(t)$ from the reservoir states $x(t)$. Figure 1 depicts the architecture. For simplicity, we will
 90 use the term “reservoir computing” for “Echo State Network” in the remainder of the paper.

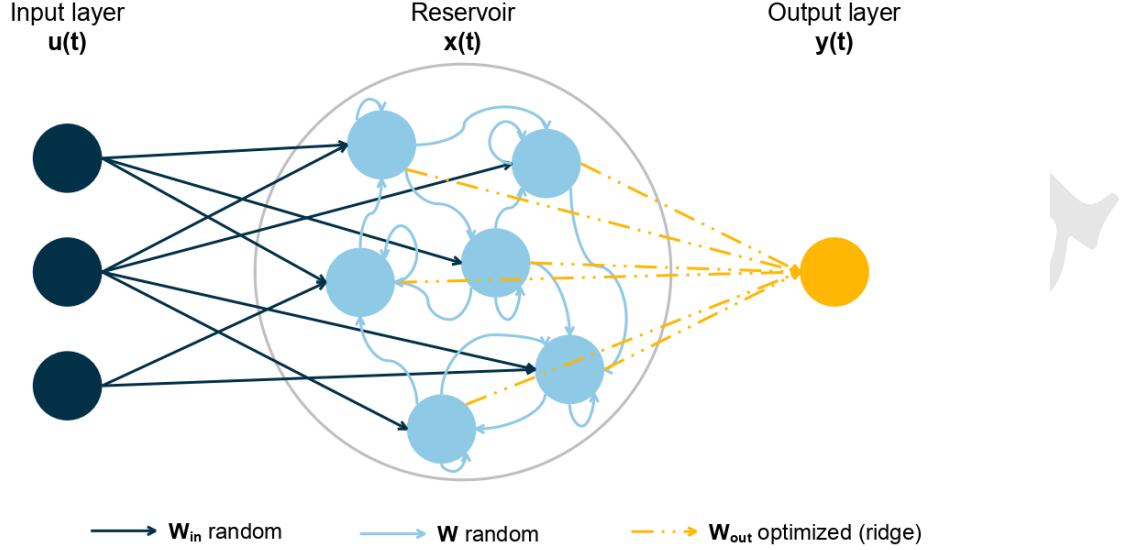


Figure 1: Reservoir computing is composed of an input layer, a reservoir and an output layer. Connection between input layer and reservoir and inside reservoir are random. Only the output layer is optimized based on a ridge regression. Adapted from Trouvain et al. (2020)

91 The input layer $u(t)$ is an M -dimension vector, where M is the number of input time series, which
 92 corresponds to the values of the input time series at time t where $t = 1, \dots, T$. The reservoir layer $x(t)$
 93 is an N_{res} -dimensional vector where N_{res} is the number of nodes in the reservoir. The value $x(t)$ is
 94 defined as follow:

$$x(t+1) = (1 - \alpha)x(t) + \alpha \tanh(Wx(t) + W_{in}u(t+1)) \quad (1)$$

95 The leaking rate $\alpha \in [0, 1]$ defines the update rate of the nodes. The closer α is to 1, the more the
 96 reservoir is sensitive to new inputs $u(t)$. Therefore, the reservoir state at time $t+1$ denoted $x(t+1)$
 97 depends on the reservoir state at the previous time $x(t)$ and the new inputs $u(t+1)$. Both W_{in} and W
 98 are random matrices of size $N_{res} \times M$ and $N_{res} \times N_{res}$ respectively.

99 W_{in} is a matrix (usually sparse) generated using a Bernoulli (bimodal) distribution where each value
 100 can be either $-I_{scale}(m)$ or $I_{scale}(m)$ with an equal probability where $m = 1, \dots, M$ corresponds to a
 101 given feature in the input layer. The input scaling, denoted I_{scale} , is a hyperparameter coefficient
 102 common to all features from the input layer or specific to each feature m . In that case, the more
 103 important the feature is, the greater should be its input scaling. W is a matrix (usually sparse) where
 104 values are generated from a Gaussian distribution $\mathcal{N}(0, 1)$. Then, the W matrix is scaled according to
 105 the defined spectral radius, a hyperparameter defining the highest eigen value of W .

106 The final layer is a linear regression with ridge penalization where the explanatory features are the
 107 reservoir state and the variable to be explained is the outcome to predict such that:

$$W_{out} = YX^T(XX^T + \lambda I)^{-1}$$

108 Where $x(t)$ and $y(t)$ are accumulated in X and Y respectively such that:

$$X = \begin{bmatrix} x(1) \\ x(2) \\ \dots \\ x(T) \end{bmatrix} \text{ and } Y = \begin{bmatrix} y(1) \\ y(2) \\ \dots \\ y(T) \end{bmatrix}$$

109 The parameter λ is the ridge penalization which aims to prevent overfitting. Additionally, one can also
110 connect the input layer to the output layer to the reservoir nodes. In that case, X is the accumulation
111 of both such that :

$$X = \begin{bmatrix} x(1), u(1) \\ x(2), u(2) \\ \dots \\ x(T), u(T) \end{bmatrix} \text{ and } Y = \begin{bmatrix} y(1) \\ y(2) \\ \dots \\ y(T) \end{bmatrix}$$

112 Overall, there are four main hyperparameters to be chosen by the user: i) the leaking rate which
113 defines the memory of the RC, ii) the input scaling which defines the relative importance of the
114 features, iii) the spectral radius which defines the connections of the neurons inside the reservoir
115 which in turn defines the degree of non-linear combination of features, and iv) the ridge penalization
116 which controls the degree of overfitting. The choice of hyperparameters often requires the user to
117 evaluate the performance of different combinations of hyperparameters on a validation set before
118 selecting the optimal combination to forecast on the test set.

119 3 Usage workflow

120 In this section, we will cover the basics of `reservoirnet` use including installation, classification and
121 regression. A more in depth description is provided in Section 4 with the covid-19 forecast use case.

122 3.1 Installation

123 `reservoirnet` is an R package making the Python module `ReservoirPy` easily callable from R using
124 `reticulate` R package Ushey, Allaire, and Tang (2024). It is available on CRAN (see <https://cran.r-project.org/package=reservoirnet>) and can be installed using:

```
# Install reservoirnet package from CRAN
install.packages("reservoirnet")
```

126 Reservoir Computing (RC) is well suited to both regression and classification tasks. We will introduce
127 a simple example for both task.

128 3.2 Package workflow overview

129 The workflow of `reservoirnet` is described in Figure 2. A reservoir model is created by the association
130 of an input layer (a matrix), a reservoir, and an output layer. Both the reservoir and the output layer
131 are created using the function `reservoirnet::createNode()` by specifying the node type (i.e., either
132 `Reservoir` or `Ridge`).

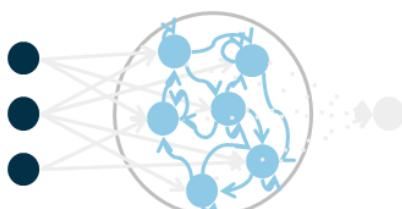
133 This function accepts several arguments to specify the hyperparameters of the reservoir and will be
134 detailed in future sections. After the reservoir and output layer are created, they can be connected
135 using the `%>>%` operator, a specific pipe operator dedicated to `reservoirnet`. The model can then be
136 fitted using `reservoirR_fit()` and used to make predictions on a new dataset using `predict_seq()`.

Input layer :X



Instantiate reservoir :

```
reservoir <- createNode(nodeType = "Reservoir")
```



Instantiate output layer :

```
readout <- createNode(nodeType = "Ridge")
```



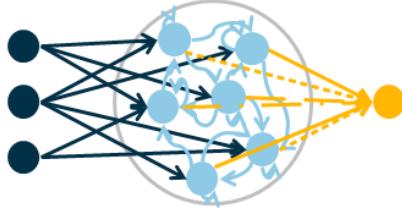
Build model :

```
model <- reservoir %>>% readout
```



Fit model :

```
fit <- reservoirR_fit(node = model,
  X = X,
  Y = Y)
```



Forecast :

```
predict_seq(node = fit$fit, X = X)
```

Figure 2: Workflow of reservoirnet.

137 **3.3 Basic regression use-case**

138 **3.3.1 Covid-19 data**

139 In this first use-case, we will introduce the fundamental usage of the `reservoirnet` package. This
140 demonstration will be conducted using the COVID-19 dataset that is included within the package.
141 These data encompass hospitalization, positive RT-PCR (Reverse Transcription Polymerase Chain
142 Reaction) results, and overall RT-PCR data sourced from Santé Publique France, which are publicly
143 available on `data.gouv.fr` (for further details, refer to `help(dfCovid)`). Our primary objective is to
144 predict the number of hospitalized patients 14 days into the future. To accomplish this, we will
145 initially train our model on data preceding the date of January 1, 2022, and subsequently apply it to
146 forecast values using the subsequent dataset.

147 We can proceed by loading useful packages - namely `ggplot2` Wickham, Navarro, and Pedersen
148 (2018) and `dplyr` Wickham et al. (2023), data and define the task:

```
# Load usefull packages
library(dplyr)
library(ggplot2)
library(reservoirnet)

# load dfCovid data from the reservoirnet package which contains Covid data
data("dfCovid")

# Set the forecast horizon to 14 days
dist_forecast = 14

# Set the train-test split to 2022-01-01
traintest_date = as.Date("2022-01-01")

Due to the substantial fluctuations observed in both RT-PCR metrics, our initial step involves applying
a moving average computation over the most recent 7-day periods for these features. Additionally,
we augment the dataset by introducing an outcome column and an outcomeDate column, which
will serve as valuable inputs for model training. Moreover, we calculate the outcome_deriv as the
difference between the outcome and the number of hospitalized patients (hosp), representing the
variation in hospitalization in relation to the current count of hospitalized individuals. The resulting
smoothed data is visualized in Figure 3.

dfOutcome <- dfCovid %>%
  # outcome at 14 days
  mutate(outcome = lead(x = hosp, n = dist_forecast),
         # Create a new column 'outcome' which contains the number of
         # hospitalizations ('hosp') shifted forward by 'dist_forecast' days
         # (14 days). This represents the outcome we want to predict.

         outcomeDate = date + dist_forecast,
         # Create a new column 'outcomeDate' which is the current date plus the
         # forecast period (14 days).

         outcome_deriv = outcome - hosp) %>%
  # Create a new column 'outcome_deriv' which is the difference between
  # the predicted outcome and current hospitalizations.
  # This represents the change in hospitalizations over the forecast
  # period.

  # rolling average for tested and positive_pcr
```

```

mutate_at(.vars = c("Positive", "Tested"),
  .funs = function(x) slider::slide_dbl(.x = x,
    .before = 6,
    .f = mean))
# Apply a rolling mean (7-day average) to the 'Positive' and
# 'Tested' columns.
# The 'slider::slide_dbl' function is used to calculate the mean
# over a window of 7 days (current day + 6 days before). This
# smooths out daily fluctuations and provides a better trend
# indicator.

```




Figure 3: Hospitalizations, number of positive RT-PCR and number of RT-PCR of Bordeaux University Hospital.

156 3.3.2 First reservoir

157 Setting a reservoir is done with the `createNode()` function. The important hyperparameters are the
158 following :

- 159 • Number of nodes (`units`) : it corresponds to the number of nodes inside the reservoir. Usually,
160 the more the better, but more nodes increases the computation time.
- 161 • Leaking rate (`lr`) : the leaking rate corresponds to the balance between the new inputs and the
162 previous state. A leaking rate of 1 only consider information from new inputs.
- 163 • Spectral radius (`sr`): the spectral radius is the maximum absolute eigenvalue of the reservoir
164 connectivity matrix. A small spectral radius induces stable dynamics inside the reservoir, a
165 high spectral radius induces a chaotic regime inside the reservoir.
- 166 • Input scaling (`input_scaling`): the input scaling is a gain applied to the input features of the
167 reservoir.
- 168 • Warmup (`warmup`) : it corresponds to the number of time step during which the data are
169 propagating into the reservoir but not used to fit the output layer. This hyperparameter is set
170 in the `reservoirR_fit()` function.

171 In addition, we can set the seed (`seed`). Because the reservoir connections are set at random, setting
172 the seed is a good approach to ensure reproducibility.

173 For this part of the tutorial, we will set the hyperparameter at a given value. Hyperparameter
174 optimization will be detailed at Section 4.

```
# Create a reservoir computing node using the 'createNode' function from the
# reservoirnet package.

# Arguments:
# - nodeType = "Reservoir": Specify the type of node to be a reservoir.
# - seed = 1: Set the seed for reproducibility, ensuring consistent results
#             when the model is run multiple times.
# - units = 500: Set the number of reservoir units (neurons) to 500.
# - lr = 0.7: Set the leakage rate (lr) of the reservoir, which controls how
#             quickly the reservoir state decays over time.
# - sr = 1: Set the spectral radius (sr) of the reservoir, which influences the
#             stability and memory capacity of the reservoir.
# - input_scaling = 1: Set the input scaling factor, which scales the input
#                     signal before it is fed into the reservoir.

reservoir <- reservoirnet::createNode(nodeType = "Reservoir",
                                       seed = 1,
                                       units = 500,
                                       lr = 0.7,
                                       sr = 1,
                                       input_scaling = 1)
```

175 Then we can feed the data to the reservoir and see the activation state of the reservoir $x(t)$. To do so,
176 we first prepare the data and transform it to a matrix.

```
## select explanatory features of the train set and transform it to an array
X <- dfOutcome %>%
  filter(outcomeDate < traintest_date) %>%
  select(hosp, Positive, Tested) %>%
  as.matrix()
```

177 Then we run the `predict_seq()` function. It takes as input a node (i.e a reservoir or a reservoir
178 associated with an output layer) and the feature matrix.

```
# Generate the state of the reservoir using the 'predict_seq' function from the
# reservoirnet package.

# Arguments:
# - node = reservoir: The reservoir computing node created earlier.
# - X = X: The input data matrix containing the features 'hosp', 'Positive',
#           and 'Tested'.
# The function computes the state of the reservoir for each time step in the
# input sequence, effectively transforming the input data into the reservoir's
# high-dimensional state space.

reservoir_state <- predict_seq(node = reservoir, X = X)
```

179 Now we can visualize node activation using the `plot()` function presented at Figure 4 .

```
# Plot the reservoir state activation over time
```

```
plot(reservoir_state)
```




Figure 4: 20 random nodes activation over time.

180 Numerous nodes within the system exhibit a consistent equilibrium state. The challenge arises when
 181 the output layer attempts to extract knowledge from these nodes, as they do not convey meaningful
 182 information. This issue can be attributed to the disparate scales of the features. To address this
 183 concern, a practical approach involves normalizing the features by dividing each of them by their
 184 respective maximum values, thereby scaling them within the range of -1 to 1 by dividing by the
 185 maximum of the absolute value. Of note, here the features will be scaled between 0 and 1 because all
 186 features are positive.

```
# Standardise features by dividing by the maximum value can improve performance
# After standardisation, all features are on a similar scale which helps RC
stand_max <- function(x) return(x/max(abs(x)))
# scaled features
Xstand <- df0$Outcome %>%
  filter(date < traintest_date) %>%
  select(hosp, Positive, Tested) %>%
  mutate_all(.funs = stand_max) %>%
  as.matrix() %>%
  as.array()
```

187 We then feed them to the reservoir and plot the node activation again. Compared to Figure 4, the
 188 obtained node activation at Figure 5 shows interesting trend outputs as no node seems saturated.

```
# feed the scaled features to the reservoir
reservoir_state_stand <- predict_seq(node = reservoir,
                                       X = Xstand,
                                       reset = TRUE)

# plot the output
plot(reservoir_state_stand)
```

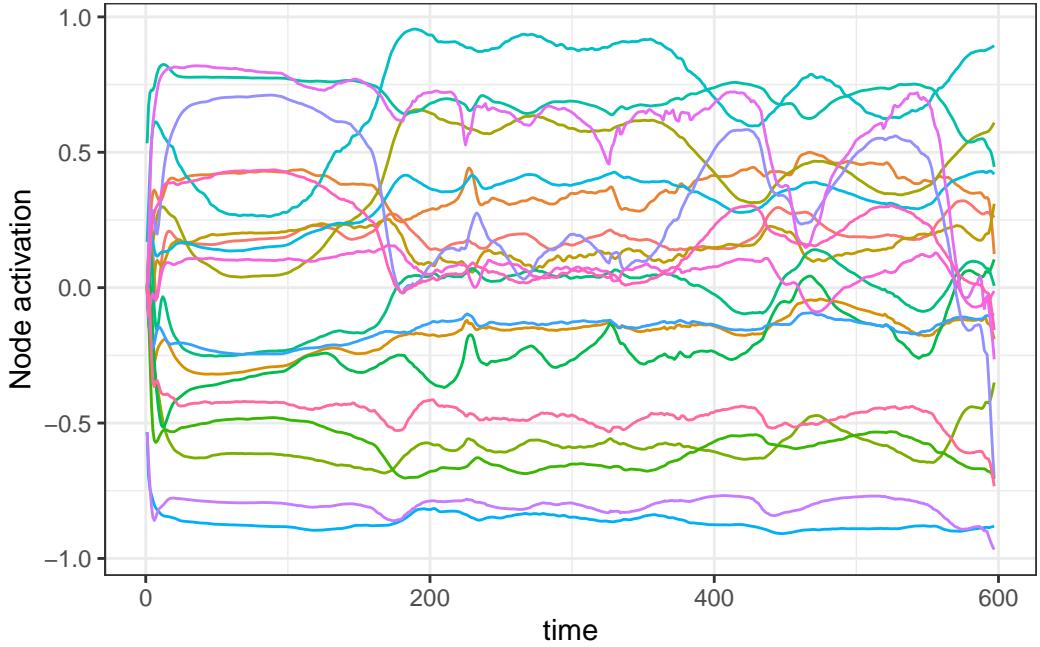


Figure 5: 20 random node activation over time. Scaled features.

189 3.3.3 Forecast

190 In order to train the reservoir, we should train the last layer which linearly combines the neuron's
 191 output.

192 3.3.3.1 Set the ESN

193 Initially, we establish the output layer with the `createNode()` function, incorporating a ridge penalty
 194 set at `1e3`. It's important to note that this hyperparameter can be subject to optimization, a topic
 195 that will be explored in Section 4. This parameter plays a pivotal role in fine-tuning the model's
 196 conformity to the data. When set excessively high, the risk of underfitting arises, whereas setting
 197 it too low can lead to overfitting. We connect the output layer to the reservoir, with the `link()`
 198 function, making the model ready to be trained.

```
readout <- reservoirnet::createNode(nodeType = "Ridge",
                                     ridge = 1e3)
# Create a readout node using ridge regression with the 'createNode' function
# from the reservoirnet package.
# Arguments:
# - nodeType = "Ridge": Specify the type of node to be a ridge regression
#                     readout.
# - ridge = 1e3: Set the regularization parameter (ridge) for the ridge
#                 regression to 1000.
# Ridge regression is used to prevent overfitting by adding a penalty on the
# size of the coefficients.

model <- reservoirnet::link(reservoir, readout)
# Link the reservoir and readout nodes to form a complete reservoir computing
# model. The 'link' function connects the high-dimensional state generated by
```

```
# the reservoir to the readout layer, allowing the model to learn the mapping  
# from the reservoir states to the target outputs.
```

199 3.3.3.2 Set the data

200 First we separate the train set on which we will learn the ridge coefficients and the test set on which
201 we will make the forecast. We define the train set to be all the data before 2022-01-01 and the test
202 data to be all the data to have forecast both on train and test sets.

```
# Perform some data management to isolate train and test sets  
# train set  
dftrain <- dfOutcome %>% filter(outcomeDate <= traintest_date)  
yTrain <- dftrain %>% select(outcome)  
yTrain_variation <- dftrain %>% select(outcome_deriv)  
xTrain <- dftrain %>% select(hosp, Positive, Tested)  
# test set  
xTest <- dfOutcome %>% select(hosp, Positive, Tested)
```

203 We standardize with the same formula as seen before. We learn the standardization on the training
204 set and apply it on the test set. Then we convert the dataframe to matrix.

```
# copy train and test sets  
xTrainstand <- xTrain  
xTeststand <- xTest  
# standardise based on training set values  
ls_fct_stand <- apply(xTrain,  
                        MARGIN = 2,  
                        FUN = function(x) feature/(max(x)))  
lapply(X = names(ls_fct_stand),  
       FUN = function(x){  
         xTrainstand[,x] <- ls_fct_stand[[x]](feature = xTrain[,x])  
         xTeststand[,x] <- ls_fct_stand[[x]](feature = xTest[,x])  
         return()  
       })  
# convert to array  
lsdf <- lapply(list(yTrain = yTrain,  
                     yTrain_variation = yTrain_variation,  
                     xTrain = xTrainstand,  
                     xTest = xTeststand),  
               function(x) as.matrix(x))
```

205 3.3.3.3 Train the model and predict

206 We then feed the reservoir with the train set using the `reservoirR_fit()` function. To do so, we set
207 a `warmup` of 30 days during which the data are propagating into the reservoir but not used to fit the
208 output layer.

```
### train the reservoir ridge output  
fit <- reservoirnet::reservoirR_fit(node = model,  
                                      X = lsdf$xTrain,  
                                      Y = lsdf$yTrain,  
                                      warmup = 30,  
                                      reset = TRUE)
```

209 Now that the ridge layer is trained, we can forecast using the `predict_seq()` function. We set the
210 parameter `reset` to TRUE in order to clean the reservoir from the data used by the training set.

```
# Forecast with the trained reservoir on the test data
vec_pred <- reservoirnet::predict_seq(node = fit$fit,
                                         X = lsdf$xTest,
                                         reset = TRUE)

# Make figure to represent forecast on the train and test sets.

dfOutcome %>%
  mutate(pred = vec_pred) %>%
  na.omit() %>%
  ggplot(mapping = aes(x = outcomeDate)) +
  geom_line(mapping = aes(y = outcome,
                           color = "observed")) +
  geom_line(mapping = aes(y = pred,
                           color = "forecast")) +
  annotate("rect",
           xmin = traintest_date,
           xmax = max(dfOutcome$outcomeDate, na.rm = T),
           ymin = 0,
           ymax = max(dfOutcome$outcome, na.rm = T)*1.1,
           alpha = .2) +
  annotate("text", label = "Test set",
           x = as.Date("2022-08-01"), y = 2200, size = 7) +
  annotate("text", label = "Train set",
           x = as.Date("2021-03-01"), y = 2200, size = 7) +
  scale_color_manual(values = c("#3772ff", "#080708")) +
  theme_minimal() +
  labs(color = "", x = "Date", y = "Hospitalizations")
```

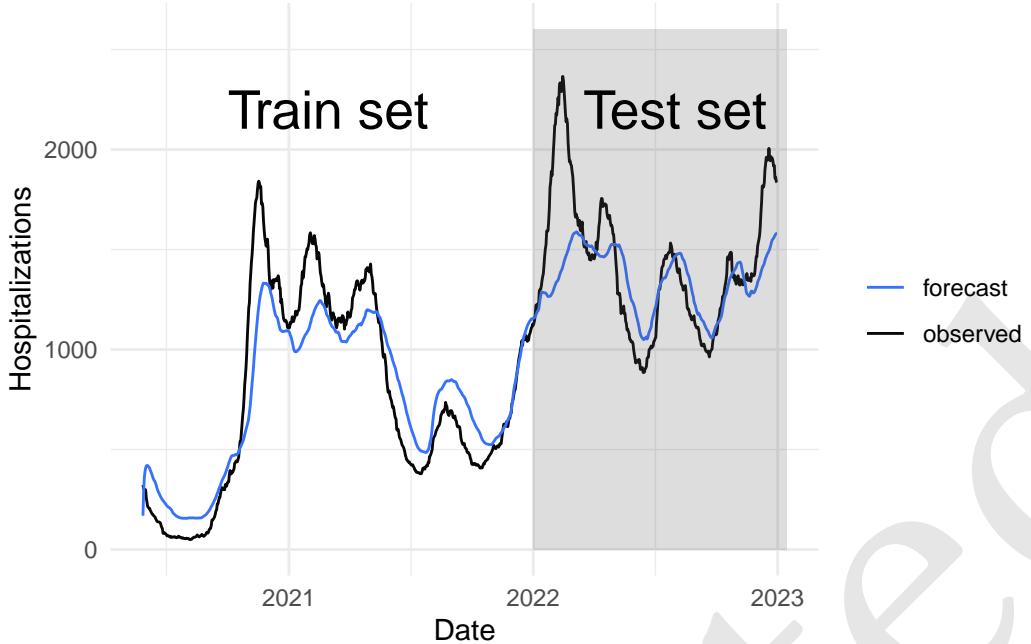


Figure 6: Forecast

211 We observe that the model forecast at Figure 6 is not fully accurate, both on the test set and the
 212 train set. In that case, one option could be to reduce ridge penalization to fit more closely the data,
 213 the optimization of ridge hyperparameter will be discussed at Section 4. Another possibility is to
 214 ease the learning of the algorithm by forecasting the variation of the hospitalization instead of
 215 the number of hospitalized patients. For that step, we will learn on the `outcome_deriv` contained
 216 in `yTrain_variation` data which is defined outcome as `outcome_deriv = outcome - hosp`. As
 217 depicted at Figure 7, this strategy improved the model forecast.

```
## Fit reservoir on outcome variation instead of raw outcome
fit2 <- reservoirnet::reservoirR_fit(node = model,
                                       X = lsdf$xTrain,
                                       Y = lsdf$yTrain_variation,
                                       warmup = 30,
                                       reset = TRUE)

## Get the forecast on the test set
vec_pred2_variation <- reservoirnet::predict_seq(node = fit2$fit,
                                                   X = lsdf$xTest,
                                                   reset = TRUE)

## Transform the outcome variation forecast into hospitalization forecast
vec_pred2 <- vec_pred2_variation + xTest$hosp

## Plot the results
dfOutcome %>%
  mutate(Raw = vec_pred2,
        Variation = vec_pred2) %>%
  tidyr::pivot_longer(cols = c(Raw, Variation),
                      names_to = "Outcome_type",
                      values_to = "Forecast") %>%
  na.omit() %>%
```

```

ggplot(mapping = aes(x = outcomeDate)) +
  geom_line(mapping = aes(y = outcome,
                          color = "observed")) +
  geom_line(mapping = aes(y = Forecast,
                          color = "Forecast")) +
  annotate("rect",
    xmin = traintest_date,
    xmax = max(dfOutcome$outcomeDate, na.rm = T),
    ymin = 0,
    ymax = max(dfOutcome$outcome, na.rm = T)*1.1,
    alpha = .2) +
  annotate("text", label = "Test set",
    x = as.Date("2022-08-01"), y = 2200, size = 5) +
  annotate("text", label = "Train set",
    x = as.Date("2021-03-01"), y = 2200, size = 5) +
  facet_wrap(Outcome_type ~ .,
    labeller = label_bquote(cols = "Outcome" : .(Outcome_type))) +
  scale_color_manual(values = c("#3772ff", "#080708")) +
  theme_minimal() +
  theme(legend.position = "bottom") +
  labs(color = "", x = "Date", y = "Hospitalizations")

```




Figure 7: Covid-19 hospitalizations forecast. The model is either trained to forecast the number of hospitalizations (denoted Raw) or the variation of the hospitalizations compared to current level of hospitalisation (denoted Variation)

218 **3.4 Classification**

219 **3.4.1 The Japanese vowel dataset**

220 This example is largely inspired from the [classification tutorial of reservoirpy](#). To illustrate the
221 classification task, we will use the Japanese vowel dataset (Kudo, Toyama, and Shimbo (1999)). The
222 data can be loaded from `reservoirnet` as follow :

```
# Get the Japanese vowels dataset using the 'generate_data' function from the
# reservoirnet package.
# The dataset contains preprocessed features and labels for classification.
# Then we isolate train and test sets
japanese_vowels <- reservoirnet::generate_data(dataset = "japanese_vowels")[[1]]
X_train <- japanese_vowels$X_train
Y_train <- japanese_vowels$Y_train
X_test <- japanese_vowels$X_test
Y_test <- japanese_vowels$Y_test
```

223 The dataset comprises 640 vocalizations of the Japanese vowel æ, contributed by nine distinct
224 speakers. Each vocalization represents a time series spanning between 7 and 29 time steps, encoded
225 as a 12-dimensional vector denoting the Linear Prediction Coefficients (LPC). A visual representation
226 of six distinct utterances from the test set, originating from three different speakers, is depicted in
227 Figure 8.



Figure 8: Vowel dataset, sample with 3 speakers and 2 utterance each.

228 The primary objective involves the attribution of each utterance to its respective speaker, this is
229 denoted as classification or sequence-to-vector encoding. The secondary objective involves the
230 attribution of each time step of each utterance to its speaker, this is denoted as transduction or
231 sequence-to-sequence encoding.

232 **3.4.2 Classification (sequence-to-vector model)**

233 The first approach is the sequence-to-vector encoding. For this task we aim to predict the speaker of
234 the whole utterance (i.e the label is assigned to the whole sequence). We first start by creating the
235 reservoir and the output layer using `createNode()` function.

```
reservoir <- reservoirnet::createNode("Reservoir", units = 500,  
                                      lr=0.1, sr=0.9,  
                                      seed = 1)  
# Create a reservoir computing node with 500 units using the 'createNode'  
# function from the reservoirnet package.  
# Arguments:  
# - units = 500: Set the number of reservoir units (neurons) to 500.  
# - lr = 0.1: Set the leakage rate (lr) of the reservoir to 0.1, controlling  
#             how quickly the reservoir state decays over time.  
# - sr = 0.9: Set the spectral radius (sr) of the reservoir to 0.9, influencing  
#             the stability and memory capacity of the reservoir.  
# - seed = 1: Set the seed for reproducibility, ensuring consistent results  
#             when the model is run multiple times.  
readout <- reservoirnet::createNode("Ridge", ridge=1e-6)  
# Create a readout node using ridge regression with the 'createNode' function  
# from the reservoirnet package.  
# Arguments:  
# - ridge = 1e-6: Set the regularization parameter (ridge) for the ridge  
#                 regression to 1e-6.  
# Ridge regression is used to prevent overfitting by adding a penalty on the  
# size of the coefficients.
```

236 To perform this task, we need to modify the training and testing processes. Leveraging the inherent
237 inertia of the reservoir, information from preceding time steps is preserved, effectively endowing the
238 RC with a form of memory. Consequently, the final state vector encapsulates insights gathered from
239 all antecedent states. In the context of the sequence-to-vector encoding task, only this ultimate state
240 is employed. This process is executed as follows:

```
states_train = list()  
k <- 1  
for (x in X_train) {  
  # Loop over each training sample in X_train.  
  states <- reservoirnet::predict_seq(node = reservoir, X = x,  
                                       reset=TRUE)  
  # Generate the reservoir states for the current training sample using the  
  # 'predict_seq' function.  
  states_train[[k]] <- t(as.matrix(states[nrow(states),]))  
  # Extract the final reservoir state for the current sample and store it in  
  # 'states_train'.  
  k <- k+1  
}  
}
```

241 Then we can train the readout based on this last state vector. In that case, `Y_train` contains a single
242 label for each utterance.

```
# Fit the reservoir using the last state vector (each observation is the whole  
# vowel sequence)
```

```

res <- reservoirnet::reservoirR_fit(readout, X = states_train, Y = Y_train)

243 The prediction is also modified using only the final state :

```

```

# The operation is repeated for the test set :
# - the last state vector is extracted for each utterance
# - the model is used to forecast the utterance from this last state vector
Y_pred <- list()
k <- 1
for (x in X_test) {
  states <- reservoirnet::predict_seq(node = reservoir, X = x,
                                         reset=TRUE)
  y <- reservoirnet::predict_seq(node = readout,
                                 X = as.array(states[nrow(states),]))
  Y_pred[[k]] <- y
  k <- k+1
}

```

```

244 Figure 9 shows the prediction for the 6 utterances depicted at Figure 8 where the model correctly
245 identifies the speaker.

```

```

# A figure represents the performance on the test set
dfplotseqtovec <- lapply(vec_sample,
  FUN = function(i){
    speaker <- which(Y_test[[i]][1,] == 1)
    Y_pred[[i]] %>%
      as.data.frame() %>%
      tidyr::pivot_longer(cols = everything(),
                          names_to = "pred_speaker",
                          values_to = "prediction") %>%
      mutate(pred_speaker = gsub(x = pred_speaker,
                                  pattern = "V", ""))
    mutate(speaker = speaker, .before = 1,
           utterance = i,
           target = speaker == pred_speaker) %>%
    return()
  }) %>%
bind_rows()

ggplot(dfplotseqtovec,
       mapping = aes(x = pred_speaker,
                      y = prediction,
                      fill = target)) +
  geom_bar(stat = "identity") +
  facet_wrap(utterance ~ speaker,
             labeller = label_bquote(cols = "speaker" : .(speaker)),
             ncol = 2) +
  scale_fill_manual(values = c("#BDBDBD", "#A3CEF1")) +
  theme_minimal() +
  theme(legend.position = "none") +
  labs(y = 'Score',
       x = "Speaker")

```

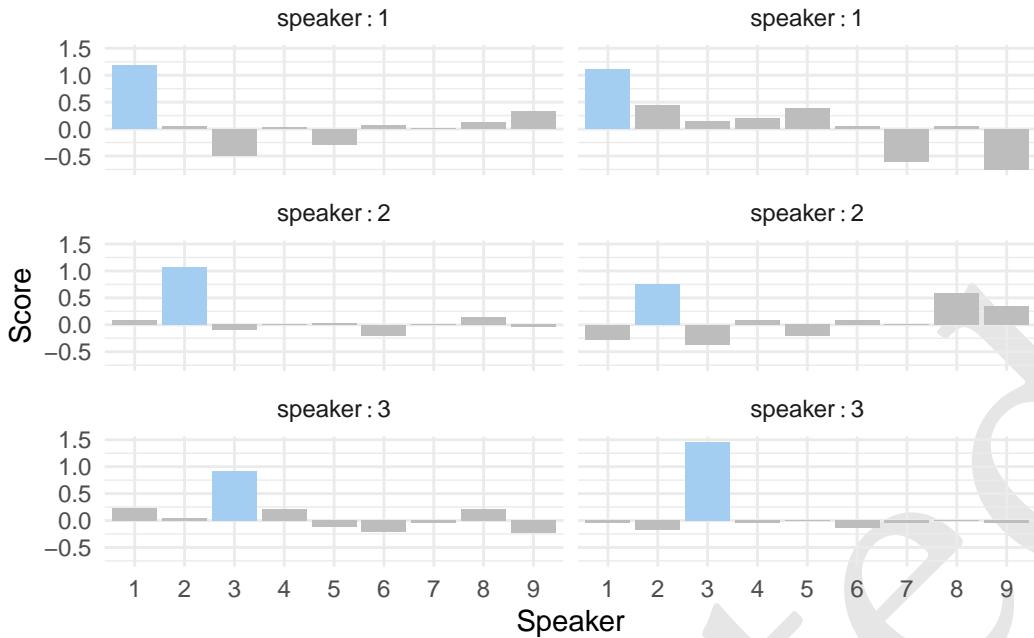


Figure 9: Prediction in a sequence-to-sequence approach 6 samples with 3 speakers and 2 utterance each. The speaker to predict is depicted in blue. For each of the 6 utterance, the model correctly identifies the speaker.

246 Then, we can also compute the overall accuracy :

```
# The overall accuracy is evaluated
accuracy <- function(pred, truth) mean(pred == truth)

Y_pred_class <- sapply(Y_pred,
                        FUN = function(x) apply(as.matrix(x), 1, which.max))
Y_test_class <- sapply(Y_test,
                        FUN = function(x) apply(as.matrix(x), 1, which.max))

score <- accuracy(pred = Y_test_class, truth = Y_pred_class)

print(paste0("Accuracy: ", round(score * 100, 3), "%"))

247 [1] "Accuracy: 92.703%"
```

248 3.4.3 Transduction (sequence-to-sequence model)

249 For this task, the goal is to predict the speaker for each time step of each utterance. The first
 250 step is to get the data where the label is repeated for each time step. This is easily done with the
 251 `repeat_targets` argument as follow :

```
# For this new task where we want to forecast for each time step (instead of each utterance)
# we start by getting the data in the appropriate format
# Then we split the train and test data
japanese_vowels <- reservoirnet::generate_data(
  dataset = "japanese_vowels",
  repeat_targets=TRUE)$japanese_vowels
```

```

X_train <- japanese_vowels$X_train
Y_train <- japanese_vowels$Y_train
X_test <- japanese_vowels$X_test
Y_test <- japanese_vowels$Y_test

```

252 Then we can train a simple Echo State Network to solve this task. For this example we will connect
 253 both the input layer and the reservoir layer to the readout layer which is performed by the `%>>%`
 254 operator :

```

# Create an input, a reservoir and an output layers
source <- createNode("Input")
readout <- createNode("Ridge", ridge=1e-6)
reservoir <- createNode("Reservoir", units = 500, lr=0.1, sr=0.9, seed = 1)
# Connect the input layer to the reservoir and connect both the input layer and
# the reservoir to the output layer
model <- list(source %>>% reservoir, source) %>>% readout

```

255 We can then fit the model and predict the labels for the test data. The `reset` parameter is set to TRUE
 256 to remove information from the reservoir from the training process.

```

# Fit the RC model
model_fit <- reservoirnet::reservoirR_fit(node = model,
                                             X = X_train,
                                             Y = Y_train,
                                             warmup = 2)

# Predict with the fitted model
Y_pred <- reservoirnet::predict_seq(node = model_fit$fit,
                                      X = X_test,
                                      reset = TRUE)

```

257 From the `Y_pred` and `Y_test` we represent at Figure 10 the predictions for the same patients as in
 258 Figure 8.

```

# Make a graph with a label for each time of each utterance
dfplotseqtoseq <- lapply(vec_sample,
                          FUN = function(i){
                            speaker <- which(Y_test[[i]][1,] == 1)
                            Y_pred[[i]] %>%
                              as.data.frame() %>%
                              tibble::rowid_to_column(var = "Time") %>%
                              tidyverse::pivot_longer(cols = -Time,
                                                     names_to = "pred_speaker",
                                                     values_to = "prediction") %>%
                              mutate(pred_speaker = gsub(x = pred_speaker,
                                                         pattern = "V", "")),
                            speaker = speaker,
                            utterance = i,
                            .before = 1) %>%
                            return()
                          }) %>%
bind_rows()

ggplot(dfplotseqtoseq, mapping = aes(x = Time,

```

```
    y = pred_speaker,  
    fill = prediction) +  
geom_tile() +  
facet_wrap(utterance ~ speaker,  
          labeller = label_bquote(cols = "speaker" : .(speaker)),  
          ncol = 2) +  
scale_fill_gradient2(low = "#8ECAE6", high = "#FB8500", mid = "#023047",  
                     midpoint = 0) +  
theme_minimal() +  
labs(y = 'Predicted speaker',  
     fill = "Prediction score")
```

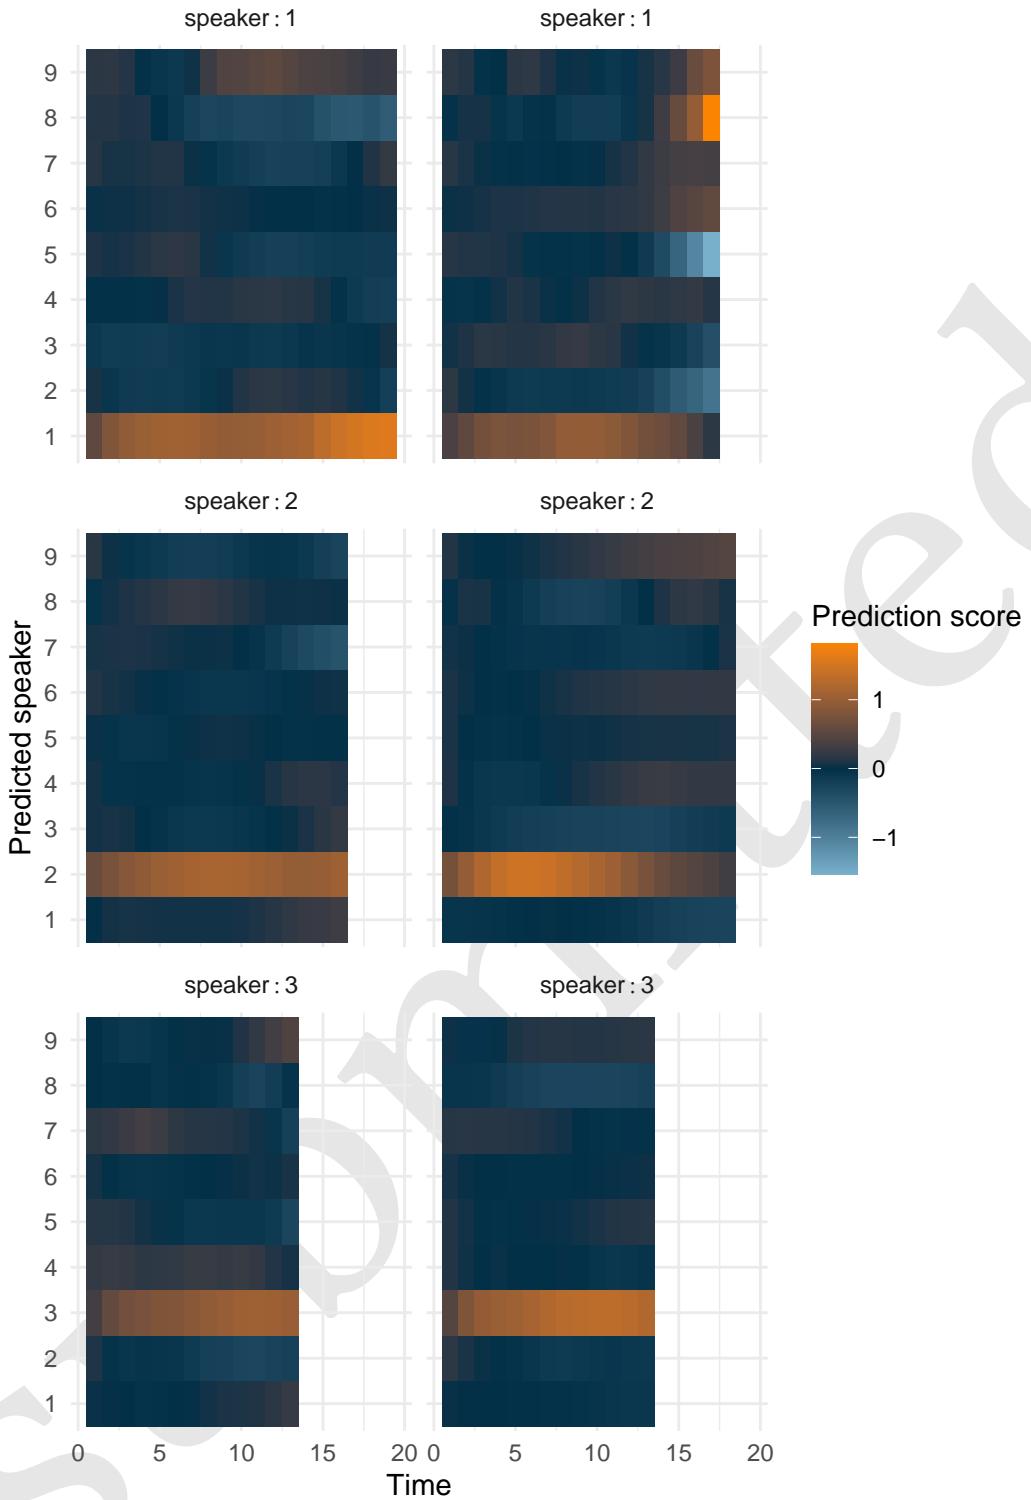


Figure 10: Prediction in a sequence-to-sequence approach 6 samples with 3 speakers and 2 utterance each. The higher the score of the speaker, the lighter the color.

259 For those 6 utterances, the model correctly identify the speaker for most of the time steps. We can
 260 then evaluate the overall accuracy of the model :

```
# Compute the accuracy
```

```

Y_pred_class <- sapply(Y_pred, FUN = function(x) apply(as.matrix(x),
1,
which.max))
Y_test_class <- sapply(Y_test, FUN = function(x) apply(as.matrix(x),
1,
which.max))
score <- accuracy(array(unlist(Y_pred_class)), array(unlist(Y_test_class)))

print(paste0("Accuracy: ", round(score * 100,3) ,"%"))

261 [1] "Accuracy: 92.456%"

```

262 4 Avanced case-study: Covid-19 hospitalizations forecast

263 4.1 Introduction

264 Since late 2020, millions of cases of SARS-CoV-2 infection have been documented across the globe
 265 (World Health Organisation 2020; COVID-19 Cumulative Infection Collaborators 2022; Carrat et al.
 266 2021). This ongoing pandemic has exerted significant strain on healthcare systems, resulting in a surge
 267 in hospitalizations. This surge, in turn, necessitated modifications to the healthcare infrastructure and
 268 gave rise to population-wide lockdown measures aimed at preventing the saturation of healthcare
 269 facilities (Simões et al. 2021; Hübner et al. 2020; Kim et al. 2020). The capacity to predict the
 270 trajectory of the epidemic on a regional scale is of paramount importance for effective healthcare
 271 system management.

272 Numerous COVID-19 forecasting algorithms have been proposed using different methods (e.g en-
 273 semble, deep learning, compartmental), yet none has proven entirely satisfactory (Cramer et al. 2022;
 274 Rahimi, Chen, and Gandomi 2021). In France, short-term forecasts with different methods have
 275 been evaluated with similar results (Paireau et al. 2022; Carvalho et al. 2021; Mohimont et al. 2021;
 276 Pottier 2021). In this context a machine learning algorithm based on linear regression with elastic-net
 277 penalization, leveraging both Electronic Health Records (EHRs) and public data, was implemented at
 278 Bordeaux University Hospital (Ferté et al. 2022). This model, which aimed at forecasting the number
 279 of hospitalized patients at 14 days, showed good performance but struggled to accurately anticipate
 280 dynamic shifts of the epidemic.

281 RC has been used in the context of covid-19 epidemic forecast (Kmet and Kmetova 2019; Liu et
 282 al. 2023; Ray, Chakraborty, and Ghosh 2021; Zhang et al. 2023; Ghosh et al. 2021). Among them,
 283 Ghosh et al. (2021), Liu et al. (2023) and Ray, Chakraborty, and Ghosh (2021) used it to forecast
 284 epidemic, Zhang et al. (2023) performed sentiment analysis and Kmet and Kmetova (2019) used
 285 it to solve optimal control related to vaccine. The evaluation of RC for epidemic forecast showed
 286 promising results in all approaches, being competitive with Long-Short Term Memory (LSTM) and
 287 Feed-Forward Neural Network (FFNN) in Ray, Chakraborty, and Ghosh (2021). However, the test
 288 period was short for Ghosh et al. (2021} (21 and 14 days) and Ray, Chakraborty, and Ghosh (2021)
 289 (86 days) making it difficult to evaluate the behavior of the methods during epidemic dynamic shift.
 290 This was not the case for Liu et al. (2023) (6 months) but they implemented daily ahead forecast
 291 which would be difficult to use to manage a hospital. Finally, all three implementations used only
 292 one time series as input whereas it has been shown that using different data sources could improve
 293 forecast Ferté et al. (2022). Therefore, it is still difficult to assess the usefulness of RC over a large
 294 period and using many time series as inputs.

295 RC can be viewed as an extension of penalized linear regression, where inputs undergo processing by a
 296 reservoir, introducing the capacity for memory and non-linear combinations. Given the effectiveness

297 of penalized linear regression in COVID-19 forecasting, as highlighted in Ferté et al. (2022), and the
298 promising results exhibited by RC in epidemic forecasting, as demonstrated in studies such as Ghosh
299 et al. (2021), Liu et al. (2023), and Ray, Chakraborty, and Ghosh (2021), we have opted to employ RC
300 for the prediction of hospitalizations at 14 days at the University Hospital of Bordeaux.

301 The primary aim of this study is to assess the performance of RC in this forecasting task. Secondary
302 objectives include (i) comparing the performance of RC with that of elastic-net penalized regression
303 (identified as the optimal model in Ferté et al. (2022)) and (ii) evaluating variations in RC performance
304 based on different architectural choices, such as the connection between the input layer and the
305 output layer, and the use of one input scaling per feature versus a common input scaling.

306 4.2 Methods

307 4.2.1 Data

308 The study utilized aggregated data spanning from May 16, 2020, to January 17, 2022, regarding
309 the COVID-19 epidemic in France, drawing from various sources to enhance forecasting accuracy.
310 These sources encompassed epidemiological statistics from Santé Publique France, weather data
311 from the National Oceanic and Atmospheric Administration (NOAA), both providing department-
312 level data (Smith, Lott, and Vose 2011; Etalab 2020) and Electronic Health Record (EHR) data from
313 the Bordeaux Hospital providing hospital-level data. All data were daily updated. Santé Publique
314 France data included information on hospitalizations, RT-PCR tests, positive RT-PCR results, variant
315 prevalence, and vaccination data, categorized by age groups. NOAA data contributed temperature,
316 wind speed, humidity, and dew point data, allowing for the computation of the COVID-19 Climate
317 Transmissibility Predict Index (Roumagnac et al. 2021). EHRs data included hospitalizations, ICU
318 admissions, ambulance service records, and emergency unit notes, with relevant COVID-19-related
319 concepts extracted from the notes. Data are discussed more in depth in Ferté et al. (2022).

320 First derivative over the last 7 days were computed to enrich model information. To take into account
321 measurement error and daily noise variation, data were smoothed using a local polynomial regression
322 with a span of 21 days. As previously described, input features were scaled between -1 and 1 by
323 dividing the observed value by the maximum of the absolute value of the given input feature.

324 All data are publicly available. Weather data can be obtained from Smith, Lott, and Vose (2011) using
325 R package `worldmet` (Carslaw 2023). [Vaccine data](#) can be downloaded from Etalab (2020). EHRs data
326 can be downloaded on dryad (Ferté et al. 2023). For privacy issues, publicly available EHRs data
327 below 10 patients were obfuscated to 0. For convenience, all data were downloaded, merged and
328 provided as replication material.

329 4.2.2 Evaluation framework

330 The task was to forecast 14 days ahead the number of hospitalized patients. As seen at Section 3.3,
331 we will train the model to predict the variation of hospitalization, denoted as \widehat{hosp} , defined as
332 $outcome_{t+14} = \widehat{hosp}_{t+14} - \widehat{hosp}_t$. Metrics computation and visualizations will be performed on the
333 predicted number of hospitalizations denoted as $\widehat{hosp}_{t+14} = outcome_{t+14} + \widehat{hosp}_t$.

334 The dataset was separated into two periods. First period from May 16, 2020 to March 1, 2021 served
335 to identify relevant hyperparameters. Remaining data was used to evaluate the model performance.

336 The performance of the model was evaluated according to several metrics:

- 337 • the mean absolute error : $MAE = mean(|\widehat{hosp}_{t+14} - hosp_{t+14}|)$.
- 338 • the median relative error : $MRE = median(|\frac{\widehat{hosp}_{t+14} - hosp_{t+14}}{hosp_{t+14}}|)$.

- the mean absolute error to baseline : $MAEB = \text{mean}(|\widehat{\text{hospt}_{t+14}} - \text{hospt}_{t+14}| - |\text{hospt}_t - \text{hospt}_{t+14}|)$.
- the median relative error to baseline : $MREB = \text{median}\left(\left|\frac{\widehat{\text{hospt}_{t+14}} - \text{hospt}_{t+14}}{\text{hospt}_t - \text{hospt}_{t+14}}\right|\right)$

Median was chosen over mean for *MRE* and *MREB* because those metrics tend to have extremely high values when the denominator is close to 0 (i.e when the number of hospitalized patients is close to 0 or the number of patients hospitalized at 14 days is close to the current number of hospitalized patients respectively). *MAEB* and *MREB* compare model performance to a baseline model which predicts the current number of hospitalized patients at 14 days. Those metrics help to determine the information added by the model and is a good baseline as covid-19 forecast model do not always outperform this basic forecast (Cramer et al. (2022)).

Because the outcome is obfuscated below 10 hospitalizations for privacy reason, we set both the outcome and the forecast to 10 when the observed value was 0 or the forecasted value was below 10 when evaluating the model performance.

4.2.3 Models

We compared RC to elastic-net penalized regression (denoted as *Enet*). Furthermore we evaluated RC based on several architectures. First we compared RC with a single input scaling common to all features and a RC with on specific input scaling per feature. Second we compared RC where the input layer is connected to the output layer in addition to the connection between reservoir and output layer. Therefore, five models were evaluated :

- Elastic-net penalized regression denoted *Enet*
- RC with a single input scaling and no connection between input and ouput layers denoted *Common IS R %»% O*
- RC with a single input scaling and connection between input and ouput layers denoted *Common IS I+R %»% O*
- RC with multiple input scaling and no connection between input and ouput layers denoted *Multiple IS R %»% O*
- RC with multiple input scaling and connection between input and ouput layers denoted *Multiple IS I+R %»% O*

Because of the randomness of the reservoir, we took the median forecast of 10 reservoir on the train set to evaluate the performance of a given hyperparameter set. On the test set we aggregated the forecast of 40 reservoirs, each of them having one of the 40 best hyperparameter sets found on the train set. In addition, because covid-19 hospitalization is a non-stationary process, models were re-trained everyday using all previous days. To ease computation burden, only one day over two was used to find hyperparameters on the training set.

4.2.4 Hyperparameter optimisation using random search

RC relies mainly on 4 hyperparameters including the leaking rate (i.e “memory” parameter), spectral radius (i.e “chaoticity” parameter), input scaling (i.e “feature gain” parameter) and ridge (i.e penalization parameter). Input scaling can be either, common to all features or specific to each feature which increases the number of hyperparameter by the number of features.

Following the notation from *glmnet* package (Friedman, Hastie, and Tibshirani 2010), elastic-net penalized linear regression relies on two hyperparameters, lambda (i.e the penalization parameter) and alpha (i.e the compromise between lasso and ridge penalty)

Hyperparameter were selected in the training set (i.e before March 1, 2021) using a wrapper approach and a random search sampler using 2000 samples for each model. The sampling distribution were

382 defined as follow :

- 383 • (RC) ridge and (Enet) lambda : log-uniform law defined between 1e-10 and 1e5
- 384 • (RC) input scaling and spectral radius : log-uniform law defined between 1e-5 and 1e5
- 385 • (RC) leaking rate : log-uniform law defined between 1e-3 and 1
- 386 • (Enet) alpha : uniform defined between 0 and 1

387 We provided large search space for all hyper-parameters. Search space was slightly reduced for
388 leaking rate based on previous results and because a leaking rate of 1e-3 already imply that new
389 inputs make the reservoir change really slowly which is not inline with the dynamic of covid-19 but
390 would be appropriate for an application where the phenomena to forecast has a slow dynamic.

391 Finally, we provided an additional Enet model similar to the one in Ferté et al. (2022) where alpha
392 was set to 0.5 and lambda was re-evaluated everyday in the test set based on previous data using the
393 cross-validation procedure provided by `glmnet`.

394 4.3 Results

395 The goal of this task is to predict 14 days ahead the hospitalization. Figure 11 shows both the training
396 set (i.e before 2021-03-01) and the test set where the blue curve correspond to the input features (first
397 derivatives are not shown) and the orange curves correspond to the outcome the model is trained
398 on (i.e the hospitalization variation) and the hospitalizations at 14 days on which the performance
399 metrics are computed. The figures outline that the relation between the input features and the
400 outcome evolve over time and that the time series is not stationary. For instance IPTCC (*Index*
401 *PREDICT de Transmissivité Climatique de la COVID-19*) seems correlated to the outcome except that
402 it completely miss the summer 2021 increase.

403 4.3.1 Hyperparameter selection

404 Figure 12 shows the hyperparameter optimisation using random search for the different RC architec-
405 tures. We observe that model with multiple input scaling achieved better performance on the train
406 set compared to model with single input scaling which is expected as they can adapt more closely to
407 the data thanks to specific input scaling for each feature.

408 As expected, we observe that the optimal leaking rate is above 1e-2 for all RC which is coherent with
409 the short term dynamic of covid-19 epidemic. Trend for other hyperparameter are less clear even
410 though best hyperparameter sets were close for RC with common input scaling and for RC with
411 multiple input scaling.

412 Figure 13 shows the hyperparameter search for RC with multiple input scaling and connected input
413 layer. We observe that the random search tend to favor high importance given to derivative of
414 positive RT-PCR (including the elderly) and the derivative of IPTCC. The rest of the feature do not
415 show clear pattern.

416 4.3.2 Forecast performance

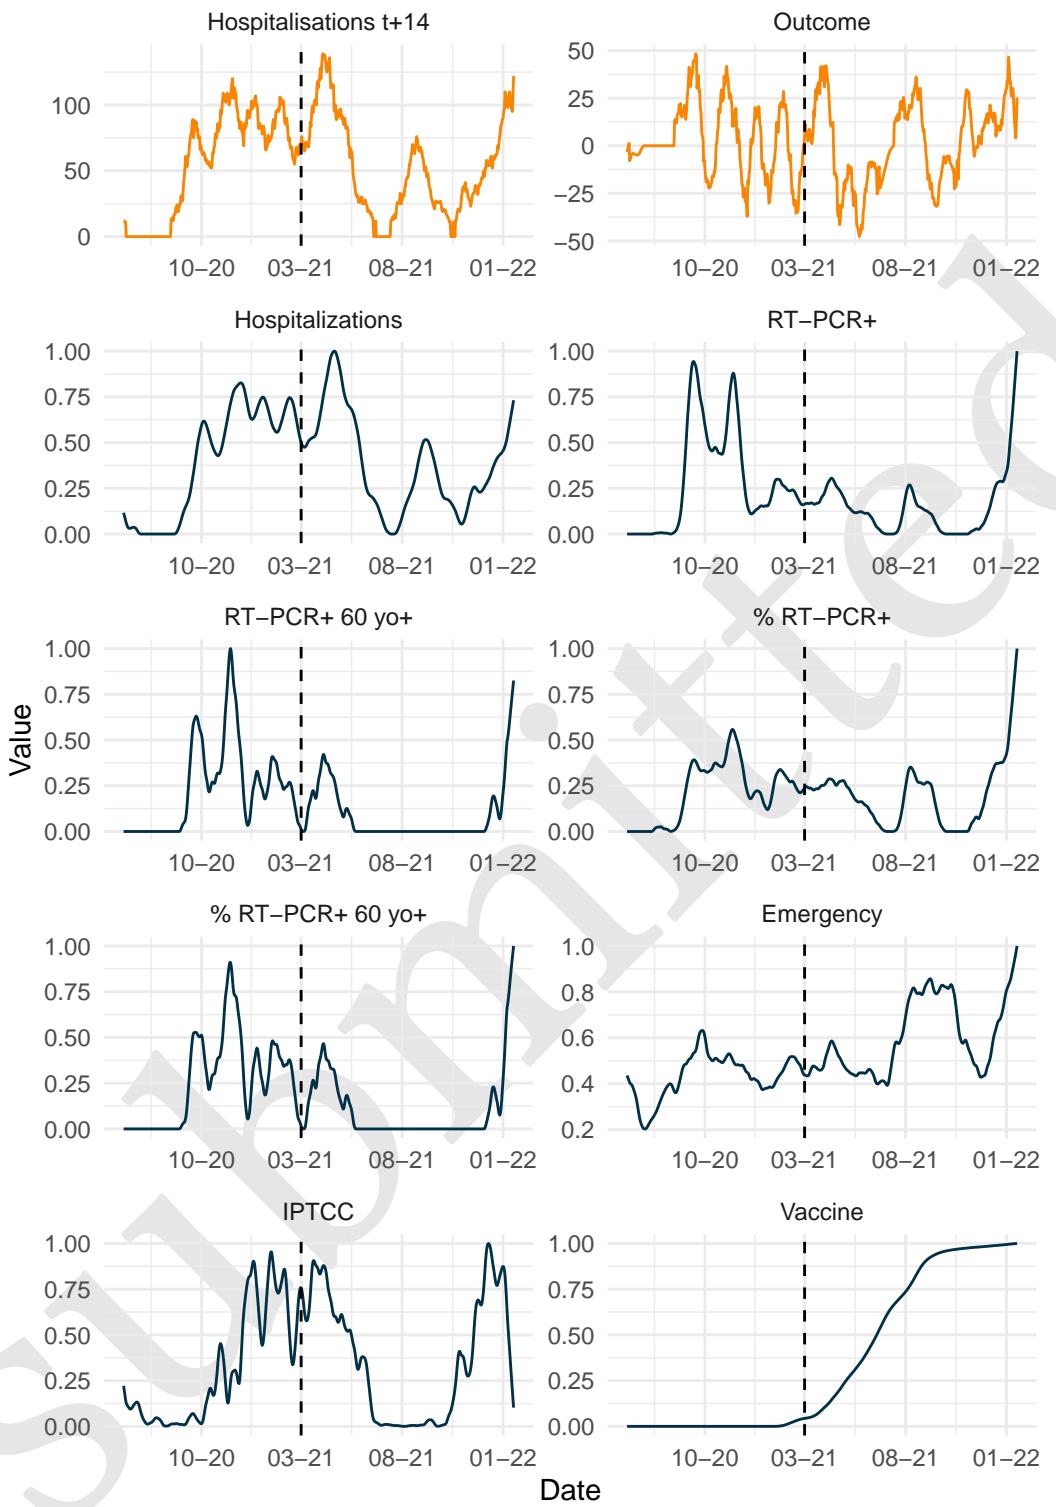


Figure 11: Covid-19 epidemic at BUH. Outcome of interest is presented in orange. Model is trained to forecast Outcome curve which corresponds to the difference between Hospitalisations at 14 days and current hospitalisations. Other features are scaled (divide by the maximum of the feature) represented in darkblue.

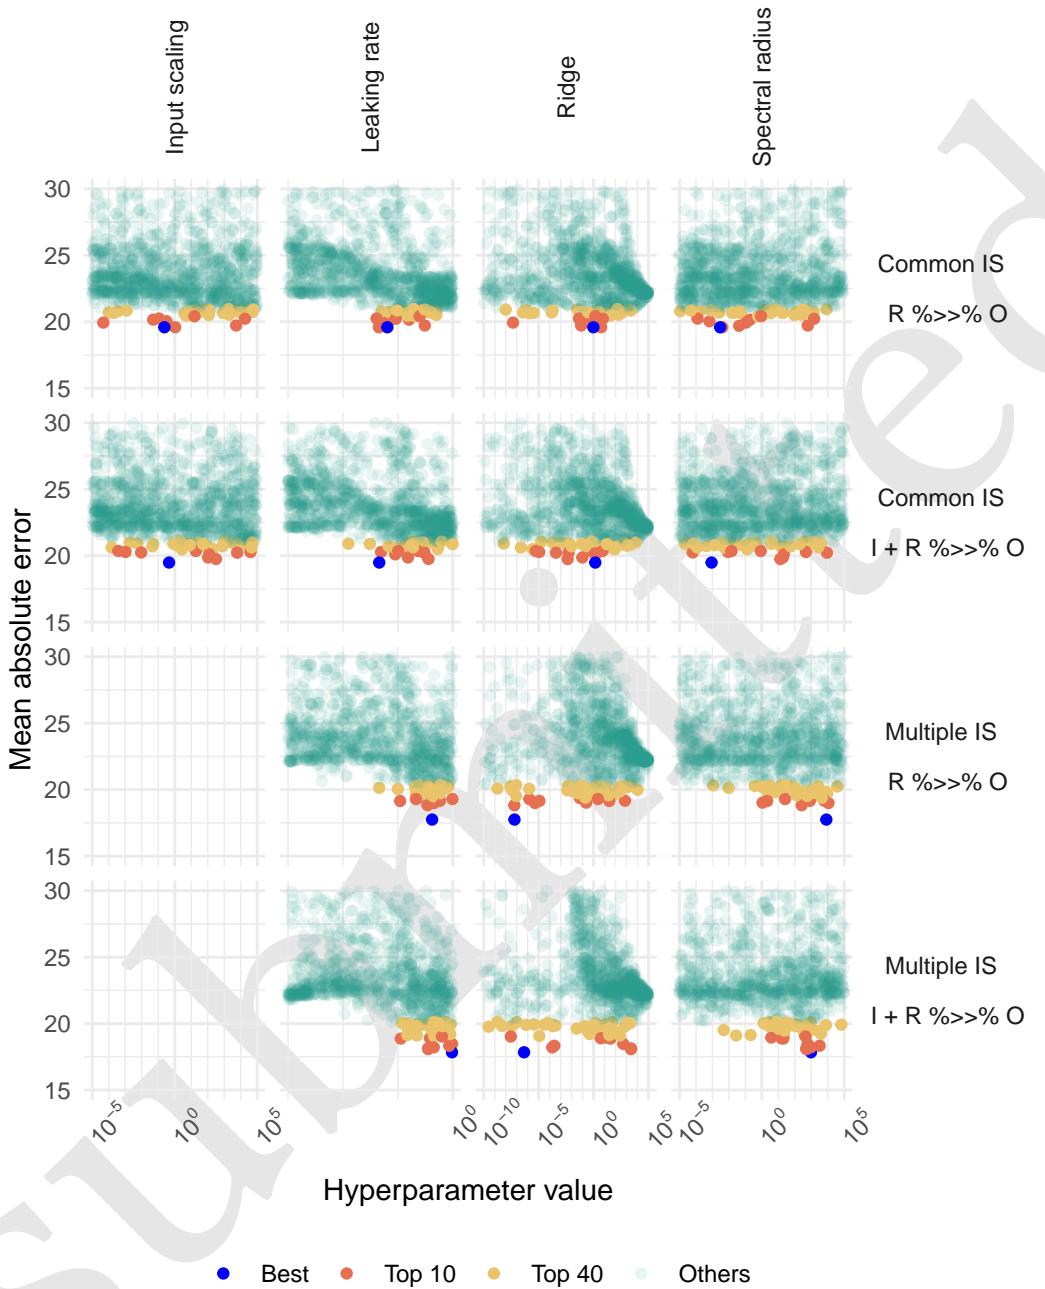


Figure 12: Hyperparameter evaluation on training set by random search. Hp sets with MAE above 30 were removed for clarity of visualisation.

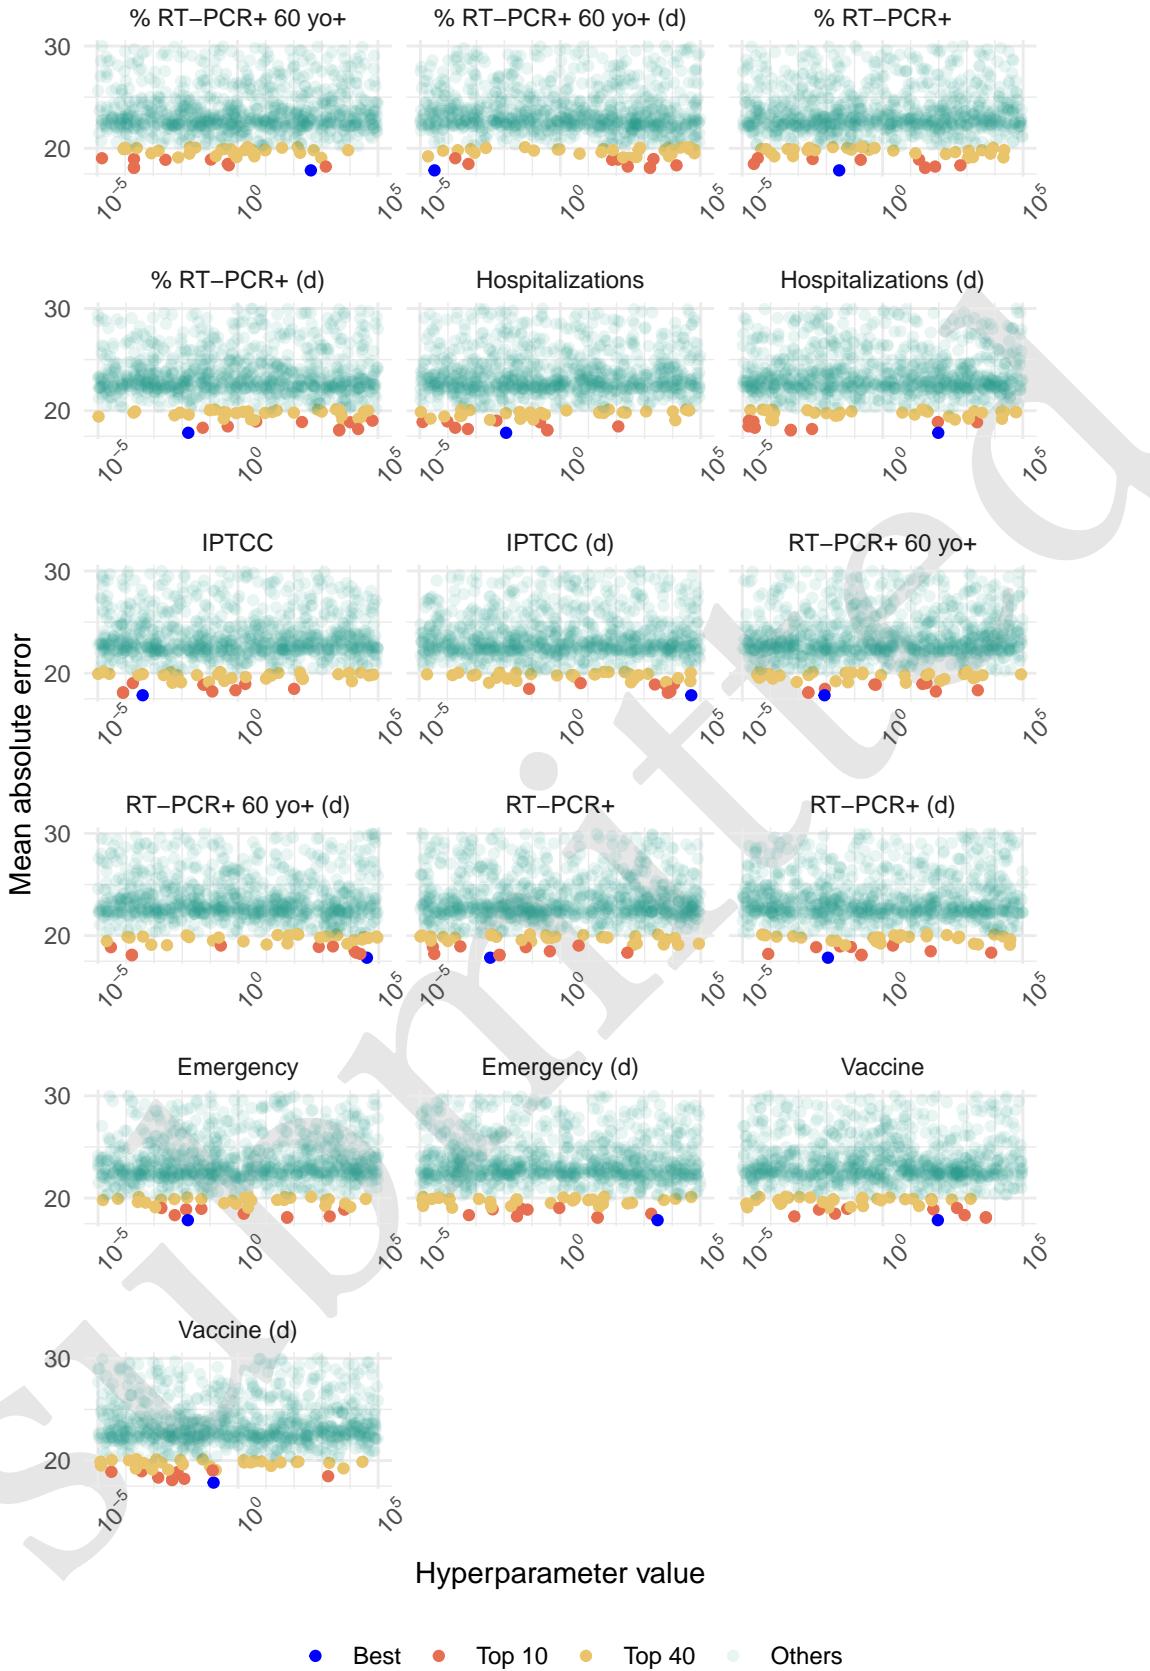


Figure 13: Hyperparameter evaluation on training set by random search of the model with multiple input scaling and no connection between input layer and output layer. Hp sets with MAE above 30 were removed for clarity of visualisation.

Table 1: Model performance with several reservoir configuration. For each setting, 40 reservoirs are computed and the forecast is the median of the 40 forecasts. Results show the performance metrics : MAE = Mean Absolute Error, MRE = Median Relative Error, MAEB = Mean Absolute Error to Baseline, MREB = Median Relative Error to Baseline.

Table 1: Model Performance

Model	MAE	MRE	MAEB	MREB
Common IS: R %»% O	15.23	0.26	-3.50	0.85
Common IS: I + R %»% O	14.84	0.26	-3.89	0.83
Multiple IS: R %»% O	15.38	0.28	-3.35	0.82
Multiple IS: I + R %»% O	15.25	0.28	-3.49	0.83
Elastic-net	16.40	0.29	-2.34	0.93

417 Table 1 shows the performance on the test set. Best model according to all metrics was RC with
 418 common input scaling and connection between input and output layers. Having one input scaling per
 419 feature did not improve the model which might be due to low generalisability of the hyperparameter
 420 of the training set to the test set due to non-stationarity. Additionally, connecting input layer to
 421 output layer improved the model forecast. All RC models performed better than the elastic-net
 422 model.

423 Figure 14 shows the forecast of the different models. We note that models struggle to accurately
 424 forecast slope shifts. For instance, summer 2021 initial increase is partially predicted by all models
 425 but its decrease is not well predicted. Winter 2021 increase is anticipated by all models but they tend
 426 to overestimate it because of the rise of vaccine effect.

427 4.3.3 Number of model to aggregate

428 Figure 15 show the individual forecast for the 40 best sets of hyperparameters of each RC architecture.
 429 Due to the internal random connection of the reservoir, we observe forecast stochasticity and relying
 430 on only one forecast is unreliable. We explored the number of model needed at Figure 16 which
 431 shows that after 10 models, forecast is stable and even 5 models for the simpler model with common
 432 input scaling which rely on less hyperparamters.

433 4.3.4 Input feature importance

434 We compared the coefficients of the output layer estimated for the input layer and the reservoir
 435 nodes. Additionally, we compared the coefficient given to the input layer by the output layer in the
 436 reservoir and the coefficient estimated by the elastic-net model.

437 Figure 17 illustrates the ranking of input layer compared to all connections to the output layer,
 438 including the 500 reservoir nodes and the 16 features of the input layer (excluding bias). The figure
 439 shows that the model with common input scaling tends to assign less weight to input layer compared
 440 to the model with multiple input scaling. This suggests that the reservoir with common input scaling
 441 provides more information than the reservoir with multiple input scaling, which aligns with its better
 442 performance, as shown in Table Table 1.

443 Furthermore, Figure 18 compares the coefficients assigned to input features by the elastic-net model
 444 and the RC models. While the coefficients are generally consistent across RC models, there are
 445 some notable differences with elastic-net. Specifically, certain features deemed important by the
 446 elastic-net model, such as the derivative of RT-PCR, and the derivative of Vaccine, are less important

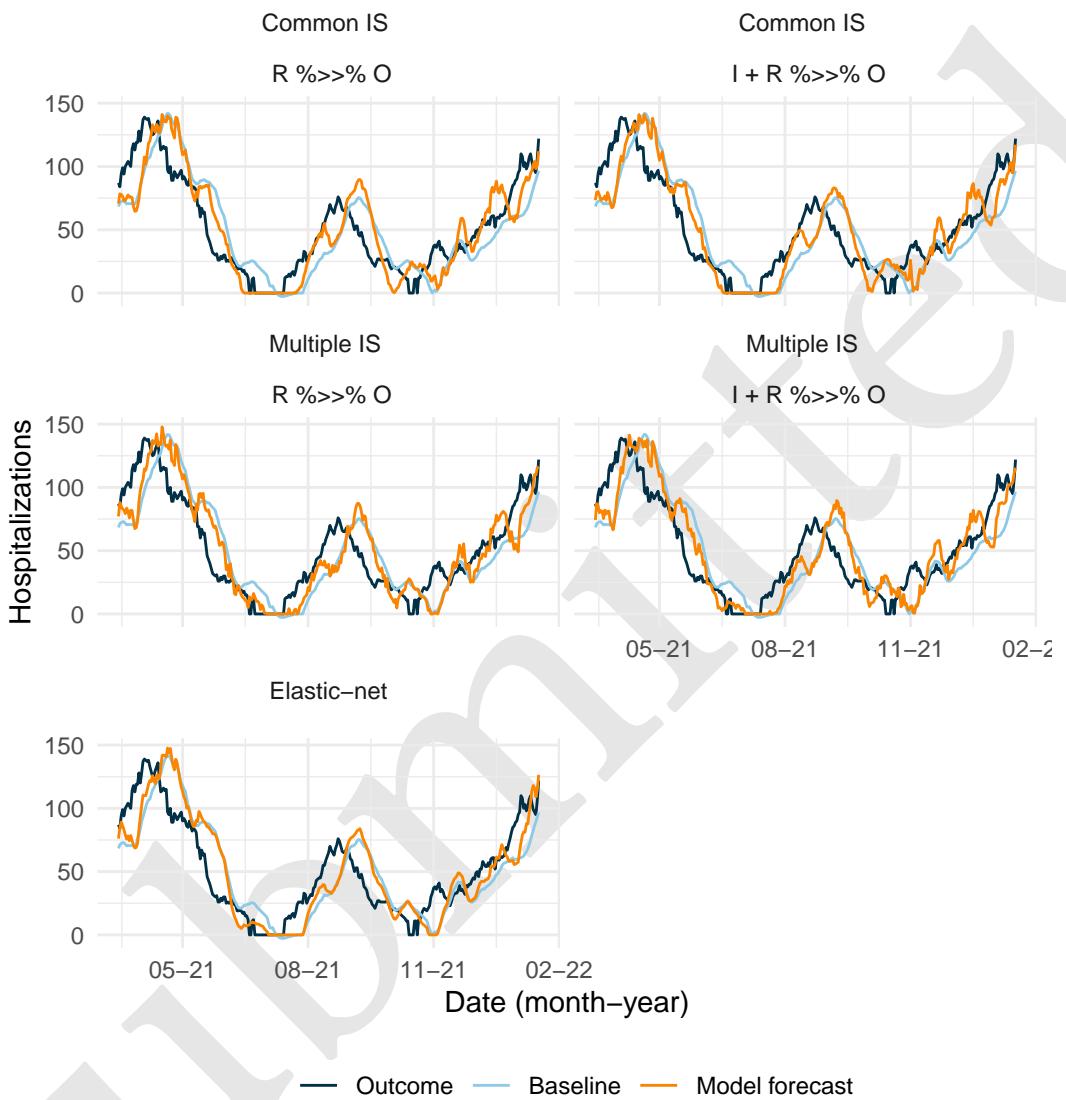


Figure 14: Reservoir computing forecast depending on the setting with and without monthly update. Red line is the median forecast of 40 reservoirs. Grey lines are individual forecast of each of the 40 reservoirs.

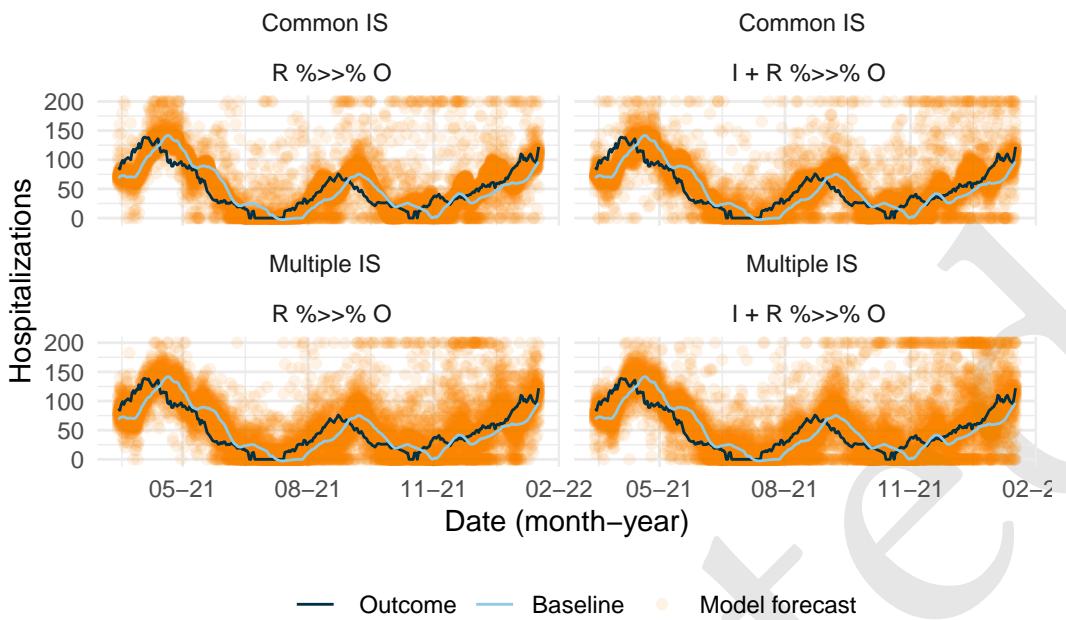


Figure 15: Individual forecast the 40 best hyperparameter sets for the different RC configuration. Forecast value above 200 were set to 200 for clarity.

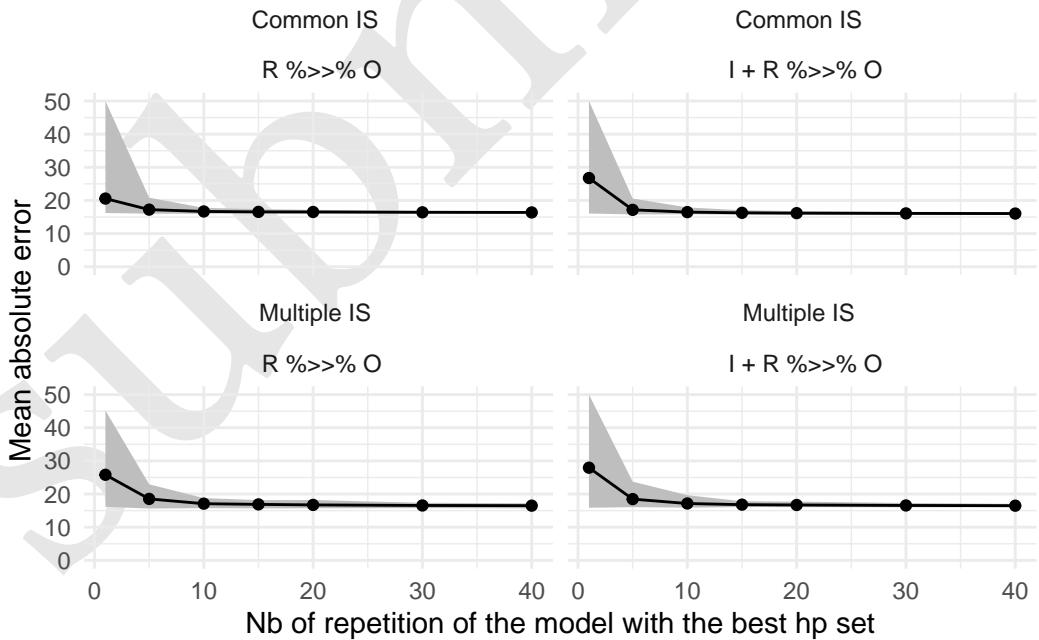


Figure 16: Mean absolute error depending on the number of aggregated reservoir.

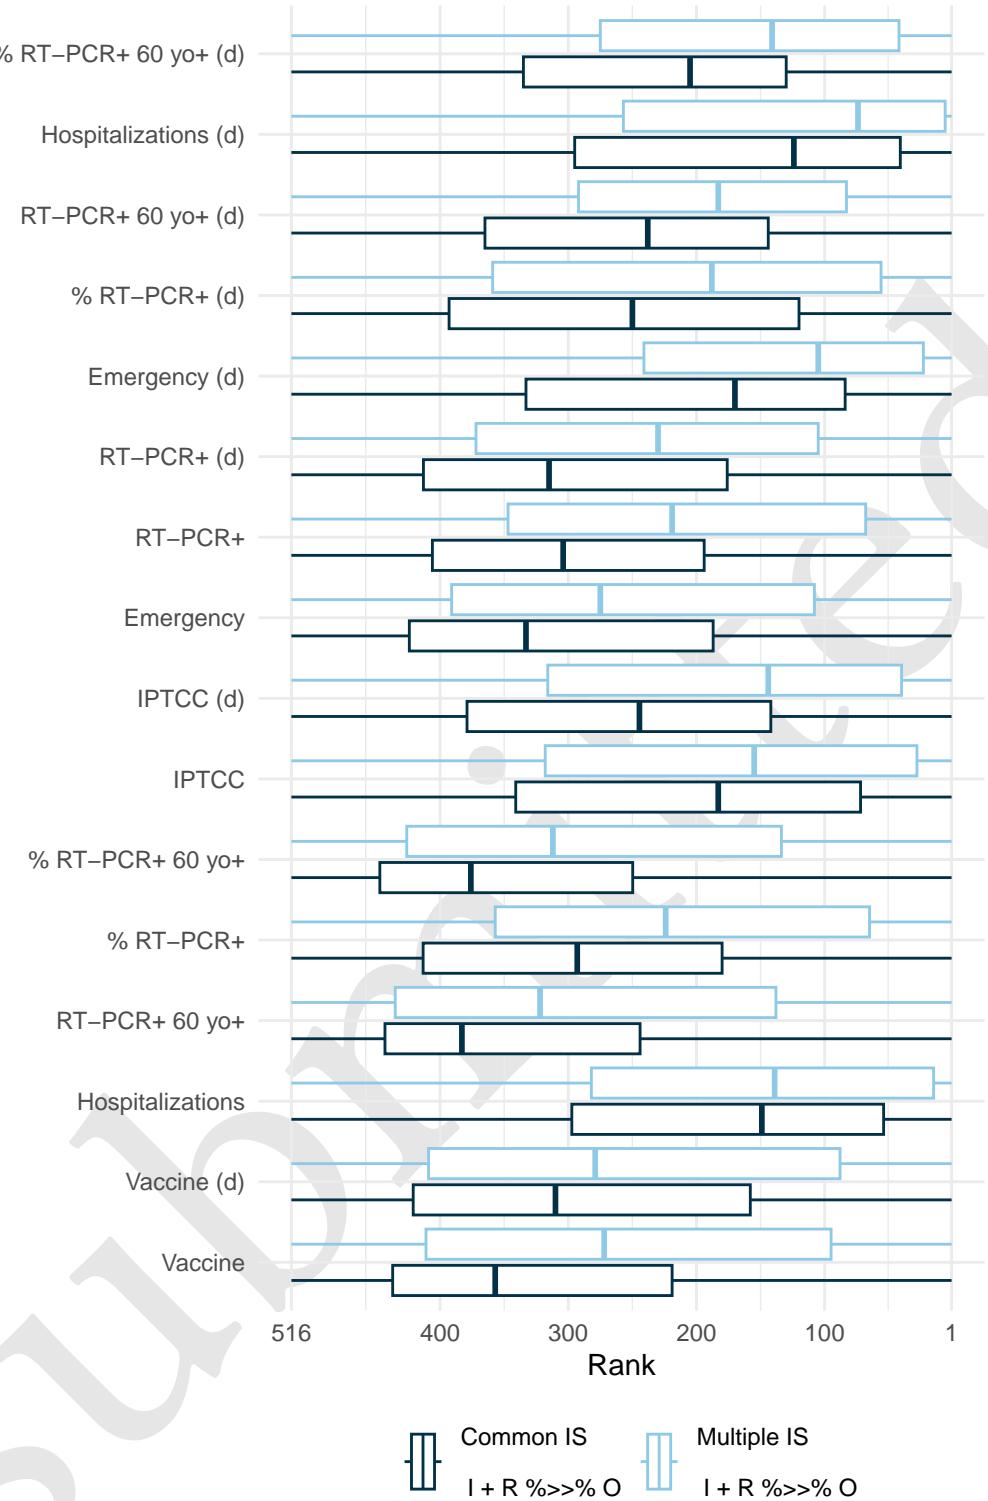


Figure 17: Mean feature importance of the 40 best hyperparameter sets by model, focus on the connection between the input and output layers. Models with direct connection between input and output layer are included. The rank is obtained by comparing the feature input layer and all other connection coefficients (both input and reservoir corresponding coefficients) attributed by the output layer at each date for each hyperparameter set. The higher the output layer's coefficient for the input layer, the closer its rank will be to 1 and the more important the feature is.

447 for the reservoir computing model. This may indicate that these features predictive ability is better
448 conveyed by their relationship with other features, which is captured by the reservoir computing
449 model but might not be by the elastic-net model. Conversely, emergency, IPTCC, proportion of
450 positive RT-PCR, and hospitalizations are more important for the reservoir computing model than
451 for the elastic-net model.

452 **4.4 Discussion**

453 In this specific application, we have demonstrated that RC exhibits commendable performance in
454 comparison to Elastic-net, which serves as the reference model. Furthermore, we highlight the
455 inherent challenges in forecasting within this context, primarily stemming from the non-stationarity
456 of the time series.

457 All computations in this study were conducted using the `reservoirnet` package, and the entire
458 codebase is accessible on Zenodo (Ferté et al. 2024). This R package demonstrates its efficacy in
459 implementing various reservoir architectures, including connection between the input layer and the
460 output layer, as well as the utilization of several input scaling, all within the context of a real-world
461 use case.

462 Given the substantial number of hyperparameters involved, we acknowledge that random search
463 may not be the most efficient optimization algorithm. We have retained this approach for the sake
464 of simplicity in this tutorial paper; however, meta-heuristic approaches, particularly those utilizing
465 evolutionary algorithms, may prove more efficient, especially when employing multiple input scaling
466 (Bala et al. 2018).

467 This study represents a novel contribution to epidemic forecasting utilizing RC. Notably, previous
468 literature predominantly focused on simpler problems characterized by fewer input features or
469 shorter evaluation periods (Liu et al. 2023; Ray, Chakraborty, and Ghosh 2021; Ghosh et al. 2021).
470 Our findings underscore the potential of this approach for future epidemics, suggesting its potential
471 to surpass more traditional epidemiological tools while maintaining a lightweight model structure
472 compared to other RNNs.

473 It is worth noting that all models, including the Ferté et al. (2022) models, encounter challenges
474 in accurately anticipating slope shifts, indicating the need for further investigation. Specifically,
475 additional work is warranted to extend the application of RC to high-dimensional settings, building
476 upon the insights gained from models based on a more extensive set of features.

477 **5 Discussion and conclusion**

478 In this paper, we introduce the R package `reservoirnet`, which serves as a versatile tool for imple-
479 menting reservoir computing based on ReservoirPy’s Python library. It offers flexibility in defining
480 the reservoir architecture, including options for specifying connections between the input layer and
481 the output layer, as well as variations in input scaling as demonstrated on a real-world use case.

482 We provided a comprehensive overview of the basic usage of the `reservoirnet` package through
483 illustrative examples in regression and classification tasks. This introductory section serves as
484 a foundation for R users, offering step-by-step guidance on constructing and training reservoir
485 computing models using the package. By demonstrating the application of RC in both regression
486 and classification scenarios, we aim to equip users with the essential knowledge and skills needed to
487 harness the capabilities of reservoir computing for diverse tasks.

488 Drawing on the robust foundation of the ReservoirPy structure, a well-maintained Python library,
489 this package inherits its reliability and longevity. We have focused on providing access to the




Figure 18: Mean feature coefficient of the 40 best hyperparameter sets by model and the elastic-net model. Only models with direct connection between input and output layer are included. The coefficients were calculated as the average value across all dates for each feature, model and hyperparameter set.

490 fundamental features, building upon the strong base provided by ReservoirPy. Therefore, this initial
491 version of reservoirnet must evolve in tandem with the growing understanding and adoption of
492 RC within the R community.

493 Acknowledgements

494 We thank Romain Griffier for his expertise on the Bordeaux University Hospital Data.

495 Experiments presented in this paper were conducted using the PlaFRIM experimental testbed, sup-
496 ported by Inria, CNRS (LABRI and IMB), Université de Bordeaux, Bordeaux INP and Conseil Régional
497 d'Aquitaine (see <https://www.plafrim.fr>), as well as by the MCIA (Mésocentre de Calcul Intensif
498 Aquitain).

499 This study was carried out in the framework of the University of Bordeaux's France 2030 program /
500 RRI PHDS.

501 References

- 502 Bala, Abubakar, Idris Ismail, Rosdiazli Ibrahim, and Sadiq M. Sait. 2018. "Applications of Meta-
503 heuristics in Reservoir Computing Techniques: A Review." *IEEE Access* 6: 58012–29. <https://doi.org/10.1109/ACCESS.2018.2873770>.
- 505 Carrat, Fabrice, Julie Figoni, Joseph Henny, Jean-Claude Desenclos, Sofiane Kab, Xavier de Lamballerie,
506 and Marie Zins. 2021. "Evidence of Early Circulation of SARS-CoV-2 in France: Findings from
507 the Population-Based 'CONSTANCES' Cohort." *European Journal of Epidemiology*, February, 1–4.
508 <https://doi.org/10.1007/s10654-020-00716-2>.
- 509 Carslaw, David. 2023. "Worldmet: Import Surface Meteorological Data from NOAA Integrated
510 Surface Database (ISD)." <https://cran.r-project.org/web/packages/worldmet/index.html>.
- 511 Carvalho, Kathleen, João Paulo Vicente, Mihajlo Jakovljevic, and João Paulo Ramos Teixeira. 2021.
512 "Analysis and Forecasting Incidence, Intensive Care Unit Admissions, and Projected Mortality
513 Attributable to COVID-19 in Portugal, the UK, Germany, Italy, and France: Predictions for 4
514 Weeks Ahead." *Bioengineering* 8 (6): 84. <https://doi.org/10.3390/bioengineering8060084>.
- 515 COVID-19 Cumulative Infection Collaborators. 2022. "Estimating Global, Regional, and National
516 Daily and Cumulative Infections with SARS-CoV-2 Through Nov 14, 2021: A Statistical Analysis."
517 *Lancet*. [https://doi.org/10.1016/S0140-6736\(22\)00484-6](https://doi.org/10.1016/S0140-6736(22)00484-6).
- 518 Cramer, Estee Y., Evan L. Ray, Velma K. Lopez, Johannes Bracher, Andrea Brennen, and et al.
519 2022. "Evaluation of Individual and Ensemble Probabilistic Forecasts of COVID-19 Mortality
520 in the United States." *Proceedings of the National Academy of Sciences* 119 (15): e2113561119.
521 <https://doi.org/10.1073/pnas.2113561119>.
- 522 Etalab. 2020. "Les Données Relatives Au COVID-19 En France - Data.gouv.fr." <https://www.data.gouv.fr/fr/pages/donnees-coronavirus/>.
- 524 Ferté, Thomas, Kalidou Ba, Dan Dutartre, Pierrick Legrand, Vianney Jouhet, Romain Griffier, Rodolphe
525 Thiébaut, Xavier Hinaut, and Boris P. Hejblum. 2024. "Thomasferte/Jss_reservoirnet: First
526 Release." Zenodo. <https://doi.org/10.5281/ZENODO.11281341>.
- 527 Ferté, Thomas, Vianney Jouhet, Romain Griffier, Boris P. Hejblum, Rodolphe Thiébaut, and Bordeaux
528 University Hospital Covid-19 Crisis Task Force. 2022. "The Benefit of Augmenting Open Data
529 with Clinical Data-Warehouse EHR for Forecasting SARS-CoV-2 Hospitalizations in Bordeaux
530 Area, France." *JAMIA Open* 5 (4): ooac086. <https://doi.org/10.1093/jamiaopen/ooac086>.
- 531 Ferté, Thomas, Vianney Jouhet, Romain Griffier, Boris Hejblum, Rodolphe Thiébaut, and Bordeaux
532 University Hospital Covid-19 Crisis Task Force. 2023. "The Benefit of Augmenting Open Data
533 with Clinical Data-Warehouse EHR for Forecasting SARS-CoV-2 Hospitalizations in Bordeaux

- 534 Area, France." Dryad. <https://doi.org/10.5061/DRYAD.HHMGQNKX>.
- 535 Friedman, Jerome, Trevor Hastie, and Rob Tibshirani. 2010. "Regularization Paths for Generalized
536 Linear Models via Coordinate Descent." *Journal of Statistical Software* 33 (1): 1–22.
- 537 Ghosh, Subrata, Abhishek Senapati, Arindam Mishra, Joydev Chattopadhyay, Syamal K. Dana,
538 Chittaranjan Hens, and Dibakar Ghosh. 2021. "Reservoir Computing on Epidemic Spreading:
539 A Case Study on COVID-19 Cases." *Physical Review E* 104 (1): 014308. <https://doi.org/10.1103/PhysRevE.104.014308>.
- 540 Hinaut, Xavier, and Peter Ford Dominey. 2013. "Real-Time Parallel Processing of Grammatical
541 Structure in the Fronto-Striatal System: A Recurrent Network Simulation Study Using Reservoir
542 Computing." *PLOS ONE* 8 (2): e52946. <https://doi.org/10.1371/journal.pone.0052946>.
- 543 Hübner, Martin, Tobias Zingg, David Martin, Philippe Eckert, and Nicolas Demartines. 2020. "Surgery
544 for Non-Covid-19 Patients During the Pandemic." *PLoS ONE* 15 (10): e0241331. <https://doi.org/10.1371/journal.pone.0241331>.
- 545 Jaeger, Herbert. 2001. "The" Echo State" Approach to Analysing and Training Recurrent Neu-
546 ral Networks-with an Erratum Note'." Bonn, Germany: German National Research Center for
547 Information Technology GMD Technical Report 148 (January).
- 548 Kim, Gina, Mengru Wang, Hanh Pan, Giana H. Davidson, Alison C. Roxby, Jen Neukirch, Danna Lei,
549 Elicia Hawken-Dennis, Louise Simpson, and Thuan D. Ong. 2020. "A Health System Response to
550 COVID-19 in Long-Term Care and Post-Acute Care: A Three-Phase Approach." *Journal of the
551 American Geriatrics Society* 68 (6): 1155–61. <https://doi.org/10.1111/jgs.16513>.
- 552 Kmet, Tibor, and Maria Kmetova. 2019. "Bézier Curve Parametrisation and Echo State Network
553 Methods for Solving Optimal Control Problems of SIR Model." *Biosystems* 186 (December): 104029.
554 <https://doi.org/10.1016/j.biosystems.2019.104029>.
- 555 Kudo, Mineichi, Jun Toyama, and Masaru Shimbo. 1999. "Multidimensional Curve Classification
556 Using Passing-Through Regions." *Pattern Recognition Letters* 20 (11): 1103–11. [https://doi.org/10.1016/S0167-8655\(99\)00077-X](https://doi.org/10.1016/S0167-8655(99)00077-X).
- 557 Liu, Bocheng, Yiyuan Xie, Weichen Liu, Xiao Jiang, Yichen Ye, Tingting Song, Junxiong Chai, Manying
558 Feng, and Haodong Yuan. 2023. "Nanophotonic Reservoir Computing for COVID-19 Pandemic
559 Forecasting." *Nonlinear Dynamics* 111 (7): 6895–6914. <https://doi.org/10.1007/s11071-022-08190-z>.
- 560 Lukoševičius, Mantas, and Herbert Jaeger. 2009. "Reservoir Computing Approaches to Recurrent
561 Neural Network Training." *Computer Science Review* 3 (3): 127–49. <https://doi.org/10.1016/j.cosrev.2009.03.005>.
- 562 Maass, Wolfgang, Thomas Natschläger, and Henry Markram. 2002. "Real-Time Computing Without
563 Stable States: A New Framework for Neural Computation Based on Perturbations." *Neural
564 Computation* 14 (11): 2531–60. <https://doi.org/10.1162/089976602760407955>.
- 565 Martinuzzi, Francesco, Chris Rackauckas, Anas Abdelrehim, Miguel D. Mahecha, and Karin Mora.
566 2022. "ReservoirComputing.jl: An Efficient and Modular Library for Reservoir Computing
567 Models." *Journal of Machine Learning Research* 23 (288): 1–8. <http://jmlr.org/papers/v23/22-0611.html>.
- 568 Mohimont, Lucas, Amine Chemchem, François Alin, Michaël Krajecki, and Luiz Angelo Steffenel.
569 2021. "Convolutional Neural Networks and Temporal CNNs for COVID-19 Forecasting in France."
570 *Applied Intelligence*, April. <https://doi.org/10.1007/s10489-021-02359-6>.
- 571 Nakane, Ryosho, Gouhei Tanaka, and Akira Hirose. 2018. "Reservoir Computing With Spin Waves
572 Excited in a Garnet Film." *IEEE Access* PP (January): 1–1. <https://doi.org/10.1109/ACCESS.2018.2794584>.
- 573 Paireau, Juliette, Alessio Andronico, Nathanaël Hozé, Maylis Layen, Pascal Crépey, Alix Roumagnac,
574 Marc Lavielle, Pierre-Yves Boëlle, and Simon Cauchemez. 2022. "An Ensemble Model Based
575 on Early Predictors to Forecast COVID-19 Health Care Demand in France." *Proceedings of the
576 National Academy of Sciences* 119 (18): e2103302119. <https://doi.org/10.1073/pnas.2103302119>.
- 577 Penkovsky, Bogdan, Laurent Larger, and Daniel Brunner. 2018. "Efficient Design of Hardware-
578

- 584 Enabled Reservoir Computing in FPGAs.” *Journal of Applied Physics* 124 (16): 162101. <https://doi.org/10.1063/1.5039826>.
- 585 Pottier, Loïc. 2021. “Forecast of the Covid19 Epidemic in France.” *medRxiv*. <https://doi.org/10.1101/2021.04.13.21255418>.
- 588 Prychynenko, Diana, Matthias Sitte, Kai Litzius, Benjamin Krüger, George Bourianoff, Mathias Kläui,
589 Jairo Sinova, and Karin Everschor-Sitte. 2018. “Magnetic Skyrmion as a Nonlinear Resistive
590 Element: A Potential Building Block for Reservoir Computing.” *Physical Review Applied* 9 (1):
591 014034. <https://doi.org/10.1103/PhysRevApplied.9.014034>.
- 592 Rafayelyan, Mushegh, Jonathan Dong, Yongqi Tan, Florent Krzakala, and Sylvain Gigan. 2020. “Large-
593 Scale Optical Reservoir Computing for Spatiotemporal Chaotic Systems Prediction.” *Physical
594 Review X* 10 (4): 041037. <https://doi.org/10.1103/PhysRevX.10.041037>.
- 595 Rahimi, Iman, Fang Chen, and Amir H. Gandomi. 2021. “A Review on COVID-19 Forecasting Models.”
596 *Neural Computing & Applications*, February, 1–11. <https://doi.org/10.1007/s00521-020-05626-8>.
- 597 Ray, Arnob, Tanujit Chakraborty, and Dibakar Ghosh. 2021. “Optimized Ensemble Deep Learning
598 Framework for Scalable Forecasting of Dynamics Containing Extreme Events.” *Chaos (Woodbury,
599 N.Y.)* 31 (11): 111105. <https://doi.org/10.1063/5.0074213>.
- 600 Roumagnac, Alix, Eurico de Carvalho Filho, Raphaël Bertrand, Anne-Kim Banchereau, and Guillaume
601 Lahache. 2021. “Étude de l’influence Potentielle de l’humidité Et de La Température Dans La
602 Propagation de La Pandémie COVID-19.” *Médecine de Catastrophe - Urgences Collectives*, Douleur
603 et situations d’exceptionPandémie COVID-19, 5 (1): 87–102. <https://doi.org/10.1016/j.pxur.2021.01.002>.
- 605 Simões, Jorge, João Paulo Moreira Magalhães, André Biscaia, António da Luz Pereira, Gonçalo
606 Figueiredo Augusto, and Inês Fronteira. 2021. “Organisation of the State, Model of Health System
607 and COVID-19 Health Outcomes in Six European Countries, During the First Months of the
608 COVID-19 Epidemic in 2020.” *The International Journal of Health Planning and Management*, June,
609 10.1002/hpm.3271. <https://doi.org/10.1002/hpm.3271>.
- 610 Smith, Adam, Neal Lott, and Russ Vose. 2011. “The Integrated Surface Database: Recent Developments
611 and Partnerships.” *Bulletin of the American Meteorological Society* 92 (6): 704–8. <https://doi.org/10.1175/2011BAMS3015.1>.
- 613 Tanaka, Gouhei, Toshiyuki Yamane, Jean Benoit Héroux, Ryosho Nakane, Naoki Kanazawa, Seiji
614 Takeda, Hidetoshi Numata, Daiju Nakano, and Akira Hirose. 2019. “Recent Advances in Physical
615 Reservoir Computing: A Review.” *Neural Networks* 115 (July): 100–123. <https://doi.org/10.1016/j.neunet.2019.03.005>.
- 617 Trouvain, Nathan, and Xavier Hinaut. 2021. “Canary Song Decoder: Transduction and Implicit
618 Segmentation with ESNs and LTSMs.” In *Artificial Neural Networks and Machine Learning – ICANN
619 2021*, edited by Igor Farkaš, Paolo Masulli, Sebastian Otte, and Stefan Wermter, 71–82. Lecture
620 Notes in Computer Science. Cham: Springer International Publishing. https://doi.org/10.1007/978-3-030-86383-8_6.
- 622 ——. 2022. “Reservoirpy: A Simple and Flexible Reservoir Computing Tool in Python.” <https://inria.hal.science/hal-03699931>.
- 624 Trouvain, Nathan, Luca Pedrelli, Thanh Trung Dinh, and Xavier Hinaut. 2020. “ReservoirPy:
625 An Efficient and User-Friendly Library to Design Echo State Networks.” In *Artificial Neural
626 Networks and Machine Learning – ICANN 2020*, 494–505. Springer International Publishing.
627 <https://inria.hal.science/hal-02595026>.
- 628 Trouvain, Nathan, Nicolas Rougier, and Xavier Hinaut. 2022. “Create Efficient and Complex Reservoir
629 Computing Architectures with ReservoirPy.” In *From Animals to Animats 16*, edited by Lola
630 Cañamero, Philippe Gaussier, Myra Wilson, Sofiane Boucenna, and Nicolas Cuperlier, 91–102.
631 Lecture Notes in Computer Science. Cham: Springer International Publishing. https://doi.org/10.1007/978-3-031-16770-6_8.
- 633 Ushey, Kevin, JJ Allaire, and Yuan Tang. 2024. *Reticulate: Interface to ‘Python’*. <https://rstudio.github.io/reticulate/>.

634 [io/reticulate/](#).
 635 Vlachas, P. R., J. Pathak, B. R. Hunt, T. P. Sapsis, M. Girvan, E. Ott, and P. Koumoutsakos. 2020.
 636 “Backpropagation Algorithms and Reservoir Computing in Recurrent Neural Networks for the
 637 Forecasting of Complex Spatiotemporal Dynamics.” *Neural Networks* 126 (June): 191–217. <https://doi.org/10.1016/j.neunet.2020.02.016>.
 638
 639 Wickham, Hadley, Romain François, Lionel Henry, Kirill Müller, and Davis Vaughan. 2023. “dplyr: A
 640 Grammar of Data Manipulation.” <https://CRAN.R-project.org/package=dplyr>.
 641 Wickham, Hadley, Danielle Navarro, and Thomas Lin Pedersen. 2018. *ggplot2: Elegant Graphics for*
 642 *Data Analysis (3e)*. 3rd ed. Springer-Verlag New York. <https://ggplot2-book.org/>.
 643 World Health Organisation. 2020. “WHO Coronavirus (COVID-19) Dashboard.” <https://covid19.who.int>.
 644
 645 Zhang, Qihuang, Grace Y. Yi, Li-Pang Chen, and Wenqing He. 2023. “Sentiment Analysis and
 646 Causal Learning of COVID-19 Tweets Prior to the Rollout of Vaccines.” *PloS One* 18 (2): e0277878.
 647 <https://doi.org/10.1371/journal.pone.0277878>.

648 Session information

```

sessionInfo()

649 R version 4.4.1 (2024-06-14)
650 Platform: x86_64-pc-linux-gnu
651 Running under: Ubuntu 24.04.1 LTS

652
653 Matrix products: default
654 BLAS:      /usr/lib/x86_64-linux-gnublas/libblas.so.3.12.0
655 LAPACK:   /usr/lib/x86_64-linux-gnulapack/liblapack.so.3.12.0

656
657 locale:
658 [1] LC_CTYPE=C.UTF-8          LC_NUMERIC=C           LC_TIME=C.UTF-8
659 [4] LC_COLLATE=C.UTF-8        LC_MONETARY=C.UTF-8    LC_MESSAGES=C.UTF-8
660 [7] LC_PAPER=C.UTF-8         LC_NAME=C             LC_ADDRESS=C
661 [10] LC_TELEPHONE=C        LC_MEASUREMENT=C.UTF-8 LC_IDENTIFICATION=C

662
663 time zone: Etc/UTC
664 tzcode source: system (glibc)

665
666 attached base packages:
667 [1] stats      graphics   grDevices datasets  utils      methods   base

668
669 other attached packages:
670 [1] reservoirnet_0.2.0 ggplot2_3.5.1       dplyr_1.1.4

671
672 loaded via a namespace (and not attached):
673 [1] utf8_1.2.4      generics_0.1.3    tidyverse_1.3.1   renv_1.0.11
674 [5] rstatix_0.7.2    lattice_0.20-45   stringi_1.8.4    digest_0.6.37
675 [9] magrittr_2.0.3    evaluate_1.0.1    grid_4.4.1       timechange_0.3.0
676 [13] fastmap_1.2.0    rprojroot_2.0.4   jsonlite_1.8.9   Matrix_1.7-1
677 [17] slider_0.3.2     backports_1.5.0   brio_1.1.5       Formula_1.2-5
678 [21] purrr_1.0.2      fansi_1.0.6       scales_1.3.0     abind_1.4-8
679 [25] cli_3.6.3       rlang_1.1.4       munspell_0.5.1   withr_3.0.2
  
```

```
680 [29] yaml_2.3.10      tools_4.4.1       ggsignif_0.6.4   colorspace_2.1-1
681 [33] ggpubr_0.6.0     here_1.0.1       broom_1.0.7      reticulate_1.40.0
682 [37] png_0.1-8        vctrs_0.6.5      R6_2.5.1         lifecycle_1.0.4
683 [41] lubridate_1.9.3   snakecase_0.11.1 stringr_1.5.1    car_3.1-3
684 [45] janitor_2.2.0    warp_0.2.1       pkgconfig_2.0.3  pillar_1.9.0
685 [49] gtable_0.3.6      Rcpp_1.0.13-1    glue_1.8.0       xfun_0.49
686 [53] tibble_3.2.1      tidyselect_1.2.1 knitr_1.49      farver_2.1.2
687 [57] htmltools_0.5.8.1 labeling_0.4.3    rmarkdown_2.29   carData_3.0-5
688 [61] testthat_3.2.1.1   compiler_4.4.1
```